

Deep Convolution Generative Adversarial Networks

This example implements the paper [Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks](#)

The implementation is very close to the Torch implementation [drgan.torch](#)

After every 100 training iterations, the files `real_samples.png` and `fake_samples.png` are written to disk with the samples from the generative model.

After every epoch, models are saved to: `netG_epoch_%d.pth` and `netD_epoch_%d.pth`

Downloading the dataset

You can download the LSUN dataset by cloning [this repo](#) and running

```
python download.py -c bedroom
```

Usage

```
usage: main.py [-h] --dataset DATASET --dataroot DATAROOT [--workers WORKERS]
               [--batchSize BATCHSIZE] [--imageSize IMAGESIZE] [--nz NZ]
               [--ngf NGF] [--ndf NDF] [--niter NITER] [--lr LR]
               [--beta1 BETA1] [--cuda] [--ngpu NGPU] [--netG NETG]
               [--netD NETD]
```

optional arguments:

<code>-h, --help</code>	show this help message and exit
<code>--dataset DATASET</code>	cifar10 lsun mnist imagenet folder lfw
<code>--dataroot DATAROOT</code>	path to dataset
<code>--workers WORKERS</code>	number of data loading workers
<code>--batchSize BATCHSIZE</code>	input batch size
<code>--imageSize IMAGESIZE</code>	the height / width of the input image to network
<code>--nz NZ</code>	size of the latent z vector
<code>--ngf NGF</code>	
<code>--ndf NDF</code>	
<code>--niter NITER</code>	number of epochs to train for
<code>--lr LR</code>	learning rate, default=0.0002
<code>--beta1 BETA1</code>	beta1 for adam. default=0.5
<code>--cuda</code>	enables cuda
<code>--ngpu NGPU</code>	number of GPUs to use
<code>--netG NETG</code>	path to netG (to continue training)
<code>--netD NETD</code>	path to netD (to continue training)
<code>--outf OUTF</code>	folder to output images and model checkpoints
<code>--manualSeed SEED</code>	manual seed

--classes CLASSES comma separated list of classes for the lsun dat

```
from __future__ import print_function
import argparse
import os
import random
import torch
import torch.nn as nn
import torch.nn.parallel
import torch.backends.cudnn as cudnn
import torch.optim as optim
import torch.utils.data
import torchvision.datasets as dset
import torchvision.transforms as transforms
import torchvision.utils as vutils

parser = argparse.ArgumentParser()
parser.add_argument('--dataset', required=True,
help='cifar10 | lsun | mnist | imagenet | folder | lfw | fake')
parser.add_argument('--dataroot', required=True,
help='path to dataset')
parser.add_argument('--workers', type=int,
help='number of data loading workers', default=2)
parser.add_argument('--batchSize', type=int,
default=64, help='input batch size')
parser.add_argument('--imageSize', type=int,
default=64, help='the height / width of the input image to network')
parser.add_argument('--nz', type=int, default=100,
help='size of the latent z vector')
parser.add_argument('--ngf', type=int, default=64)
parser.add_argument('--ndf', type=int, default=64)
parser.add_argument('--niter', type=int,
default=25, help='number of epochs to train for')
parser.add_argument('--lr', type=float,
default=0.0002, help='learning rate, default=0.0002')
```

```
parser.add_argument('--beta1', type=float,
                    default=0.5, help='beta1 for adam. default=0.5')
parser.add_argument('--cuda', action='store_true',
                    help='enables cuda')
parser.add_argument('--ngpu', type=int, default=1,
                    help='number of GPUs to use')
parser.add_argument('--netG', default='',
                    help="path to netG (to continue training)")
parser.add_argument('--netD', default='',
                    help="path to netD (to continue training)")
parser.add_argument('--outf', default='.',
                    help='folder to output images and model checkpoints')
parser.add_argument('--manualSeed', type=int,
                    help='manual seed')
parser.add_argument('--classes', default='bedroom',
                    help='comma separated list of classes for the lsun data set')
```

```
opt = parser.parse_args()
print(opt)
```

```
try:
    os.makedirs(opt.outf)
except OSError:
    pass
```

```
if opt.manualSeed is None:
    opt.manualSeed = random.randint(1, 10000)
print("Random Seed: ", opt.manualSeed)
random.seed(opt.manualSeed)
torch.manual_seed(opt.manualSeed)
```

```
cudnn.benchmark = True
```

```
if torch.cuda.is_available() and not opt.cuda:
    print("WARNING: You have a CUDA device, so you
```

```
should probably run with --cuda")

if opt.dataset in ['imagenet', 'folder', 'lfw']:
    # folder dataset
    dataset = dset.ImageFolder(root=opt.dataroot,

transform=transforms.Compose([

transforms.Resize(opt.imageSize),

transforms.CenterCrop(opt.imageSize),

transforms.ToTensor(),

transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5,
0.5))),
                        ]))
    nc=3
elif opt.dataset == 'lsun':
    classes = [ c + '_train' for c in
opt.classes.split(',') ]
    dataset = dset.LSUN(root=opt.dataroot,
classes=classes,

transform=transforms.Compose([

transforms.Resize(opt.imageSize),

transforms.CenterCrop(opt.imageSize),
                        transforms.ToTensor(),

transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5,
0.5))),
                        ]))
    nc=3
elif opt.dataset == 'cifar10':
    dataset = dset.CIFAR10(root=opt.dataroot,
```

```
download=True,

transform=transforms.Compose([
    transforms.Resize(opt.imageSize),
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5,
0.5))),
                        nc=3

elif opt.dataset == 'mnist':
    dataset = dset.MNIST(root=opt.dataroot,
download=True,

transform=transforms.Compose([
    transforms.Resize(opt.imageSize),
    transforms.ToTensor(),
    transforms.Normalize((0.5,), (0.5,)),
                        1))
    nc=1

elif opt.dataset == 'fake':
    dataset = dset.FakeData(image_size=(3,
opt.imageSize, opt.imageSize),

transform=transforms.ToTensor())
    nc=3

assert dataset
dataloader = torch.utils.data.DataLoader(dataset,
batch_size=opt.batchSize,
```

```
shuffle=True, num_workers=int(opt.workers))

device = torch.device("cuda:0" if opt.cuda else
"cpu")
ngpu = int(opt.ngpu)
nz = int(opt.nz)
ngf = int(opt.ngf)
ndf = int(opt.ndf)

# custom weights initialization called on netG and
netD
def weights_init(m):
    classname = m.__class__.__name__
    if classname.find('Conv') != -1:
        m.weight.data.normal_(0.0, 0.02)
    elif classname.find('BatchNorm') != -1:
        m.weight.data.normal_(1.0, 0.02)
        m.bias.data.fill_(0)

class Generator(nn.Module):
    def __init__(self, ngpu):
        super(Generator, self).__init__()
        self.ngpu = ngpu
        self.main = nn.Sequential(
            # input is Z, going into a convolution
            nn.ConvTranspose2d(      nz, ngf * 8, 4,
1, 0, bias=False),
            nn.BatchNorm2d(ngf * 8),
            nn.ReLU(True),
            # state size. (ngf*8) x 4 x 4
            nn.ConvTranspose2d(ngf * 8, ngf * 4, 4,
2, 1, bias=False),
            nn.BatchNorm2d(ngf * 4),
            nn.ReLU(True),
```

```

        # state size. (ngf*4) x 8 x 8
        nn.ConvTranspose2d(ngf * 4, ngf * 2, 4,
2, 1, bias=False),
        nn.BatchNorm2d(ngf * 2),
        nn.ReLU(True),
        # state size. (ngf*2) x 16 x 16
        nn.ConvTranspose2d(ngf * 2,      ngf, 4,
2, 1, bias=False),
        nn.BatchNorm2d(ngf),
        nn.ReLU(True),
        # state size. (ngf) x 32 x 32
        nn.ConvTranspose2d(      ngf,      nc, 4,
2, 1, bias=False),
        nn.Tanh()
        # state size. (nc) x 64 x 64
    )

```

```

    def forward(self, input):
        if input.is_cuda and self.ngpu > 1:
            output =
nn.parallel.data_parallel(self.main, input,
range(self.ngpu))
        else:
            output = self.main(input)
        return output

```

```

netG = Generator(ngpu).to(device)
netG.apply(weights_init)
if opt.netG != '':
    netG.load_state_dict(torch.load(opt.netG))
print(netG)

```

```

class Discriminator(nn.Module):
    def __init__(self, ngpu):
        super(Discriminator, self).__init__()

```



```
self.ngpu = ngpu
self.main = nn.Sequential(
    # input is (nc) x 64 x 64
    nn.Conv2d(nc, ndf, 4, 2, 1, bias=False),
    nn.LeakyReLU(0.2, inplace=True),
    # state size. (ndf) x 32 x 32
    nn.Conv2d(ndf, ndf * 2, 4, 2, 1,
bias=False),
    nn.BatchNorm2d(ndf * 2),
    nn.LeakyReLU(0.2, inplace=True),
    # state size. (ndf*2) x 16 x 16
    nn.Conv2d(ndf * 2, ndf * 4, 4, 2, 1,
bias=False),
    nn.BatchNorm2d(ndf * 4),
    nn.LeakyReLU(0.2, inplace=True),
    # state size. (ndf*4) x 8 x 8
    nn.Conv2d(ndf * 4, ndf * 8, 4, 2, 1,
bias=False),
    nn.BatchNorm2d(ndf * 8),
    nn.LeakyReLU(0.2, inplace=True),
    # state size. (ndf*8) x 4 x 4
    nn.Conv2d(ndf * 8, 1, 4, 1, 0,
bias=False),
    nn.Sigmoid()
)

def forward(self, input):
    if input.is_cuda and self.ngpu > 1:
        output =
nn.parallel.data_parallel(self.main, input,
range(self.ngpu))
    else:
        output = self.main(input)

    return output.view(-1, 1).squeeze(1)
```

```
netD = Discriminator(ngpu).to(device)
netD.apply(weights_init)
if opt.netD != '':
    netD.load_state_dict(torch.load(opt.netD))
print(netD)

criterion = nn.BCELoss()

fixed_noise = torch.randn(opt.batchSize, nz, 1, 1,
device=device)
real_label = 1
fake_label = 0

# setup optimizer
optimizerD = optim.Adam(netD.parameters(),
lr=opt.lr, betas=(opt.beta1, 0.999))
optimizerG = optim.Adam(netG.parameters(),
lr=opt.lr, betas=(opt.beta1, 0.999))

for epoch in range(opt.niter):
    for i, data in enumerate(dataloader, 0):
        #####
        # (1) Update D network: maximize log(D(x))
        + log(1 - D(G(z)))
        #####
        # train with real
        netD.zero_grad()
        real_cpu = data[0].to(device)
        batch_size = real_cpu.size(0)
        label = torch.full((batch_size,),
real_label, device=device)

        output = netD(real_cpu)
        errD_real = criterion(output, label)
        errD_real.backward()
        D_x = output.mean().item()
```

```

        # train with fake
        noise = torch.randn(batch_size, nz, 1, 1,
device=device)
        fake = netG(noise)
        label.fill_(fake_label)
        output = netD(fake.detach())
        errD_fake = criterion(output, label)
        errD_fake.backward()
        D_G_z1 = output.mean().item()
        errD = errD_real + errD_fake
        optimizerD.step()

#####
        # (2) Update G network: maximize
log(D(G(z)))
#####
        netG.zero_grad()
        label.fill_(real_label) # fake labels are
real for generator cost
        output = netD(fake)
        errG = criterion(output, label)
        errG.backward()
        D_G_z2 = output.mean().item()
        optimizerG.step()

        print('[%d/%d][%d/%d] Loss_D: %.4f Loss_G:
%.4f D(x): %.4f D(G(z)): %.4f / %.4f'
              % (epoch, opt.niter, i,
len(dataloader),
              errD.item(), errG.item(), D_x,
D_G_z1, D_G_z2))
        if i % 100 == 0:
            vutils.save_image(real_cpu,
                              '%s/real_samples.png' %
opt.outf,
                              normalize=True)
            fake = netG(fixed_noise)

```

```
        vutils.save_image(fake.detach(),
                           '%s/'
fake_samples_epoch_%03d.png' % (opt.outf, epoch),
                           normalize=True)

# do checkpointing
torch.save(netG.state_dict(), '%s/'
netG_epoch_%d.pth' % (opt.outf, epoch))
torch.save(netD.state_dict(), '%s/'
netD_epoch_%d.pth' % (opt.outf, epoch))
```