

Basic VAE Example

This **is** an improved implementation of the paper [Auto-Encoding Variational Bayes](<http://arxiv.org/abs/1312.6114>) by Kingma **and** Welling. It uses ReLUs **and** the adam optimizer, instead of sigmoids **and** adagrad. These changes make the network converge much faster.

```
```bash
pip install -r requirements.txt
python main.py

from __future__ import print_function
import argparse
import torch
import torch.utils.data
from torch import nn, optim
from torch.nn import functional as F
from torchvision import datasets, transforms
from torchvision.utils import save_image

parser = argparse.ArgumentParser(description='VAE
MNIST Example')
parser.add_argument('--batch-size', type=int,
default=128, metavar='N',
 help='input batch size for
training (default: 128)')
parser.add_argument('--epochs', type=int,
default=10, metavar='N',
 help='number of epochs to train
(default: 10)')
parser.add_argument('--no-cuda',
action='store_true', default=False,
 help='enables CUDA training')
parser.add_argument('--seed', type=int, default=1,
```

```
metavar='S',
 help='random seed (default: 1)')
parser.add_argument('--log-interval', type=int,
 default=10, metavar='N',
 help='how many batches to wait
before logging training status')
args = parser.parse_args()
args.cuda = not args.no_cuda and
torch.cuda.is_available()

torch.manual_seed(args.seed)

device = torch.device("cuda" if args.cuda else
"cpu")

kwargs = {'num_workers': 1, 'pin_memory': True} if
args.cuda else {}
train_loader = torch.utils.data.DataLoader(
 datasets.MNIST('../data', train=True,
download=True,
 transform=transforms.ToTensor()),
 batch_size=args.batch_size, shuffle=True,
**kwargs)
test_loader = torch.utils.data.DataLoader(
 datasets.MNIST('../data', train=False,
transform=transforms.ToTensor()),
 batch_size=args.batch_size, shuffle=True,
**kwargs)

class VAE(nn.Module):
 def __init__(self):
 super(VAE, self).__init__()

 self.fc1 = nn.Linear(784, 400)
 self.fc21 = nn.Linear(400, 20)
 self.fc22 = nn.Linear(400, 20)
```

```
self.fc3 = nn.Linear(20, 400)
self.fc4 = nn.Linear(400, 784)

def encode(self, x):
 h1 = F.relu(self.fc1(x))
 return self.fc21(h1), self.fc22(h1)

def reparameterize(self, mu, logvar):
 std = torch.exp(0.5*logvar)
 eps = torch.randn_like(std)
 return mu + eps*std

def decode(self, z):
 h3 = F.relu(self.fc3(z))
 return torch.sigmoid(self.fc4(h3))

def forward(self, x):
 mu, logvar = self.encode(x.view(-1, 784))
 z = self.reparameterize(mu, logvar)
 return self.decode(z), mu, logvar

model = VAE().to(device)
optimizer = optim.Adam(model.parameters(), lr=1e-3)

Reconstruction + KL divergence losses summed over
all elements and batch
def loss_function(recon_x, x, mu, logvar):
 BCE = F.binary_cross_entropy(recon_x,
x.view(-1, 784), reduction='sum')

 # see Appendix B from VAE paper:
 # Kingma and Welling. Auto-Encoding Variational
 Bayes. ICLR, 2014
 # https://arxiv.org/abs/1312.6114
 # 0.5 * sum(1 + log(sigma^2) - mu^2 - sigma^2)
```

```
KLD = -0.5 * torch.sum(1 + logvar - mu.pow(2) -
logvar.exp())

return BCE + KLD

def train(epoch):
 model.train()
 train_loss = 0
 for batch_idx, (data, _) in
enumerate(train_loader):
 data = data.to(device)
 optimizer.zero_grad()
 recon_batch, mu, logvar = model(data)
 loss = loss_function(recon_batch, data, mu,
logvar)
 loss.backward()
 train_loss += loss.item()
 optimizer.step()
 if batch_idx % args.log_interval == 0:
 print('Train Epoch: {} [{}/{} ({:.0f}%)]
\tLoss: {:.6f}'.format(
 epoch, batch_idx * len(data),
len(train_loader.dataset),
 100. * batch_idx /
len(train_loader),
 loss.item() / len(data)))

 print('====> Epoch: {} Average loss: {:.
4f}'.format(
 epoch, train_loss /
len(train_loader.dataset)))

def test(epoch):
 model.eval()
 test_loss = 0
```

```
 with torch.no_grad():
 for i, (data, _) in enumerate(test_loader):
 data = data.to(device)
 recon_batch, mu, logvar = model(data)
 test_loss += loss_function(recon_batch,
 data, mu, logvar).item()
 if i == 0:
 n = min(data.size(0), 8)
 comparison = torch.cat([data[:n],

recon_batch.view(args.batch_size, 1, 28, 28)[:n]])
 save_image(comparison.cpu(),
 'results/reconstruction_'
+ str(epoch) + '.png', nrow=n)

 test_loss /= len(test_loader.dataset)
 print('====> Test set loss: {:.
4f}'.format(test_loss))

if __name__ == "__main__":
 for epoch in range(1, args.epochs + 1):
 train(epoch)
 test(epoch)
 with torch.no_grad():
 sample = torch.randn(64, 20).to(device)
 sample = model.decode(sample).cpu()
 save_image(sample.view(64, 1, 28, 28),
 'results/sample_' +
str(epoch) + '.png')
```