

MULTITASK LEARNING AND SEMI-SUPERVISED LEARNING WITH NOISY DATA FOR AUDIO TAGGING

Technical Report

Osamu Akiyama

Osaka University
Faculty of Medicine
oakiyama1986@gmail.com

Junya Sato

Osaka University
Faculty of Medicine
junya.sto@gmail.com

ABSTRACT

This paper describes our submission to the DCASE 2019 challenge Task 2 "Audio tagging with noisy labels and minimal supervision" [1]. This task is a multi-label audio classification with 80 classes. The training data is composed of a small amount of reliably labeled data (curated data) and a larger amount of data with unreliable labels (noisy data). Additionally, there is a difference between data distribution between curated data and noisy data. To tackle this difficulty, we propose three strategies. The first is multitask learning using noisy data. The second is semi-supervised learning (SSL) using input data with a different distribution from labeled input data. The third is an ensemble method that averages models learned with different time windows. By using these methods, we achieved a score of 0.750 with label-weighted label-ranking average precision (lwrap), which is in the top 1% on the public leaderboard (LB).

Index Terms— Audio-Tagging, Noisy Labels, Multitask Learning, Semi-supervised Learning, Model Ensemble

1. MULTITASK LEARNING

In this task, the curated data and noisy data are labeled in a different manner, therefore treating them as the same one makes the model performance worse. To tackle this problem, we used a multitask learning approach [2]. The aim of multitasking learning is to get synergy between 2 tasks without reducing the performance of each task. Multitask learning learns features shared between two tasks and can be expected to achieve higher performance than learning independently. In our proposal, an encoder architecture learns the features shared between curated and noisy data, and the two separated FC layers learn the difference between the two data (Fig. 1). In this way, we can get the advantages of feature learning from noisy data and avoid the disadvantages of noisy label perturbation. The loss weight ratio of curated and noisy is set as 1:1. By this method, cross validation (CV) lwrap improved from 0.829 to 0.849 and score on the public LB increased + 0.021.

2. SEMI-SUPERVISED LEARNING

Because treating the noisy labels as same as the curated labels makes the model performance worse, it may be promising to do semi-supervised learning (SSL) using the noisy data without the noisy labels. However, this task is different from the data that SSL

is generally applied in two point. The first, there is a difference of data distribution between labeled data and unlabeled data. The second, this is a multi-label classification task. This makes it difficult to apply SSL. In particular, the method that generates labels online like Mean teacher [3] or MixMatch [4] tends to collapse. We tried pseudo label [5], Mean Teacher and MixMatch and all of them are not successful.




Figure 1: Over all architecture of our proposed model. (1) Basic classification (2) Soft pseudo labeling (3) Multitask learning with noisy labels

Therefore, we propose an SSL method that is robust to data distribution difference and can handle multi-label data (Fig. 1). For each noisy data sample, we guess the label using the trained model. The guessed label is sharpened by sharpening function proposed by MixMatch. We call this soft pseudo label. The basic pseudo label is a one-hot label with only one positive label so that can not apply to multi-label data. In contrast, the soft pseudo label is slightly sharpened label distribution and suits for multi-label data. Learning with soft pseudo labels is performed in parallel with multitask learning. As the temperature of sharpening function, we tried a value of 1, 1.5 or 2 and 2 was the best. Predictions of the trained model are obtained using snapshot ensemble [6] with all the folds and cycle snapshots of 5-fold CV. We used mean squared error (MSE) as a loss function. we set loss weight of semi-supervised learning as 20.

By soft pseudo labeling, The CV lwrap improved from 0.849 to 0.870 (Table 1, 5-fold soft pseudo label). On the other hand, on the public LB, improvement in score was slight (+0.001). We use predictions of all fold models to generate soft pseudo label so that high CV may be because of indirect label leak. However, even if we use labels generated by only the same fold model which has no label leak, CV was improved as compared to one without SSL (Table 1, 1-fold soft pseudo label). The model trained with the soft pseudo label is useful as a component of model averaging (+0.003 on the public LB).

3. PREPROCESSING

We used both waveform and log mel spectrogram as input data. These two data types are expected to compensate for each other.

3.1. Waveform

We tried a sampling rate of 44.1 kHz (original data) and 22.05 kHz and 44.1 kHz was better. Each input data was regularized into a range of from -1 to +1 by dividing by 32,768, the full range of 16-bit audio.

3.2. Log mel spectrogram

For the log mel spectrogram transformation, we used 128 mel frequency channels. we tried 64 and 256 but model performance decreased. The STFT hop size of 347 was used that makes log mel spectrogram 128 Hz time resolution. Log mel spectrogram was converted from power to dB after all augmentations applied. Thereafter, it was normalized by the mean and standard deviation of each data. Therefore, the mean and standard deviation values change every time and this works as a kind of augmentation. Normalization using the mean and standard deviation of the whole data decreased model performance.

4. AUGMENTATIONS

4.1. Augmentations for log mel spectrogram

4.1.1. MixUp/BC learning

MixUp/BC learning [7, 8] is an augmentation that mixes two pairs of inputs and labels with some ratio. The mixing ratio is selected from the beta distribution. We used alpha of 1.0 for Beta distribution. This makes Beta distribution equal to uniform distribution.

4.1.2. SpecAugment

SpecAugment [9] is an augmentation method for log mel spectrogram consists of three kinds of deformation. The first is time warping that deforms time-series in the time direction. The other two augmentations are time and frequency masking, modifications of Cutout [10], that masks a block of consecutive time steps or mel frequency channels. We applied frequency masking and masking width is chosen from 8 to 32 from a uniform distribution. Time warping and time masking are not effective in this task and we didn't apply them to our models.

4.1.3. Slicing

For training, audio samples which have various time lengths are converted to a certain length by random slicing. The sound samples which have short length than the slicing length are extended to the slicing length by zero paddings. We tried 2, 4 and 8 seconds (256, 521 and 1024 dimensions) as a slicing length and 4 seconds scores the best. Averaging models trained with 4-second slicing and 8-second slicing achieved a better score.

Expecting more strong augmentation effect, after basic slicing, we shorten data samples in a range of 25 - 100% of the basic slicing-length by additional slicing and extend to the basic slicing-length by zero paddings. For data samples with a time length shorter than the basic slicing-length, we shorten data samples in a range of 25 - 100% of original length by additional slicing and extend to the basic slicing-length by zero paddings. We call this additional slicing.

4.1.4. Other augmentations

We used gain augmentation with a factor randomly selected from a range of 0.80 - 1.20. We tried scaling augmentation and white noise augmentation but model performance decreased.

4.2. Augmentations for waveform

4.2.1. MixUp/BC learning

We applied MixUp to waveform input. We used alpha of 1.0 for Beta distribution as same as the case of log mel spectrogram.

4.2.2. Slicing

We applied slicing to waveform input. We tried 1.51, 3.02 and 4.54 seconds (66,650, 133,300 and 200,000 dimensions) as a slicing length and 4.54 seconds scores the best. Averaging models trained with 3.02-second slicing and 4.54-second slicing achieved a better score.

4.2.3. Other augmentations

We used scale augmentation with a factor randomly selected from a range of 0.8 - 1.25 and gain augmentation with a factor randomly selected from a range of 0.501 - 2.00.

5. MODEL ARCHITECTURE

5.1. ResNet

We selected ResNet [11] as a classification encoder model for log mel spectrogram. We compared ResNet18, ResNet34 and SE-ResNeXt50 [12] and ResNet34 performed the best. The number of trainable parameters including the multitask module is 44,210,576. We applied a global max pooling (GMP) after convolutional layers to allow variable input length.

5.2. EnvNet

We selected EnvNet-v2 [8] as a classification encoder model for waveform. The number of trainable parameters including the multitask module is 4,128,912. As same as ResNet, we applied a GMP after convolutional layers to allow variable input length.

5.3. Multitask module

For multitask learning, two separate full-connect (FC) layer sequences follow after encoder and GMP. The contents of both sequences are same and consist of FC (1024 cells) – ReLU - DropOut (drop rate = 0.2) - FC (1024 cells) – ReLU - DropOut (drop rate = 0.1) - FC (80 cells) - sigmoid. Sigmoid is replaced by softmax in model # 5 and 6 of EnvNet (Table 1).

6. TRAINING

6.1. ResNet

Adam [13] was used for optimization. Cyclic cosine annealing [14] was used for learning rate schedule. In each cycle, the learning rate is started with $1e-3$ and decrease to $1e-6$. There are 64 epochs per cycle. We used a batch size of 32 or 64. We used binary crossentropy as a loss function for basic classification and multitask learning with noisy data. We used mean squared error as a loss function for the soft pseudo label.

6.2. EnvNet

Stochastic gradient descent (SGD) was used for optimization. Cyclic cosine annealing was used for learning rate schedule. In each cycle, the learning rate is started with $1e-1$ and decrease to $1e-6$. There are 80 epochs per cycle. We used binary crossentropy as a loss function for the model using sigmoid and Kullback-Leibler (KL) divergence for the model using softmax. We used a batch size of 16 for the model using sigmoid and 64 for the model using softmax.

7. POSTPROCESSING AND ENSEMBLE

Prediction using the full length of audio input scores better than prediction using test time augmentation (TTA) with sliced audio input. This may be because important components for classification is concentrated on the beginning part of audio samples. Actually, prediction with slices of the beginning part scores better than prediction with slices of the latter part. In order to speed up the calculation, audio samples with similar lengths were grouped together, and the lengths of samples in the same group were adjusted to the same length by zero paddings and converted to minibatches. The patience for the difference of length within a group (patience rate) was adjusted based on the prediction speed.

We found padding augmentation is effective TTA. This is an augmentation method that applies zero paddings to both sides of audio samples with various length and averages prediction results. We think this method has an effect to emphasize the start and the end part of audio samples.

For model averaging, we prepared models trained with various conditions. Especially, we found that averaging with models trained with different time window is effective (Table 1 #7)

The model weights of each cycle were saved and used for snapshot ensemble. In order to reduce prediction time, the cycles and padding lengths used for the ensemble were chosen based on CV. The predictions selected are 5 fold \times (model #1 \times cycle = 2, 4, 7 and 8 \times padding = 8, 64 + model #2 \times cycle = 1, 2, 6 and 7 \times padding = 8, 64 + model #3 \times cycle = 2, 4 and 6 \times padding = 8, 64 + model #4 \times cycle = 2 and 3 \times padding = 0, 32k + model #5 \times cycle 3 and 5 \times padding = 8k, 32k + model #6 \times cycle = 7 and 9 \times padding = 8k, 32k) = 170 predictions (submission 1). The

weights of each model for averaging are model # 1: 2: 3: 4: 5: 6 = 3: 4: 3: 1: 1: 1, which is chosen based on CV. In the simplified version submission, the selected predictions are 5 fold \times (model #1 \times cycle = 2, 4, 7 and 8 \times padding = 8 + model #2 \times cycle = 1, 2, 4, 6 and 7 \times padding = 8 + model #3 \times cycle = 2, 4 and 6 \times padding = 8 + model #4 \times cycle = 2 and 3 \times padding = 8k + model #5 \times cycle 3 and 5 \times padding = 8k + model #6 \times cycle = 6, 7 and 9 \times padding = 8k) = 95 predictions (submission 2). The weights for averaging are the same as submission 1.

8. COMPARISON

Table 1 shows the results of each learning condition. The score is lwlap of 5-fold CV. Table 2 shows the results of each model averaging condition.

Table 1: Comparison of each learning condition. CV lwlap is calculated based on the best epoch of each fold in 5-fold CV except for #8, which is calculated based on the final epoch.

#	condition	CV lwlap
1	ResNet34, 512 epoch/cycle \times 1, slice length = 512, batch size = 64	0.724
2	#1 + MixUp, frequency masking and gain	0.829
3	#2 changed to 64 epoch/cycle \times 8	0.829
4	#3 + multitask (model #1)	0.849
5	#4 + 5-fold soft pseudo label, batch size = 32, + additional slicing, 64 epoch/cycle \times 7, use #1 weights as pretrained weights (model #2)	0.870
6	#5 changed to 1-fold soft pseudo label	0.858
7	#4 changed to slice length = 1,024 (model #3), 64 epoch/cycle \times 6	0.840
8	EnvNetV2, 400 epoch/cycle \times 1, slice length = 133,300, batch size = 16, augmentation = MixUp, gain and scaling, multitask, softmax	0.809
9	#8 changed to sigmoid, batch size = 64, 80 epoch/cycle \times 3, use #8 weights as pretrained weight (model #4)	0.814
10	#8 changed to 80 epoch/cycle \times 5, use #8 weights as pretrained weight (model #5)	0.818
11	#10 changed to 80 epoch/cycle \times 10, slice = 200,000 (model #6)	0.820

Table 2: Comparison of model averaging.

#	condition	CV lwlap
1	model #1 cycle = 1-8, pad = 8, 32	0.868
2	model #2 cycle = 1-7, pad = 8, 32	0.886
3	model #3 cycle = 1-6, pad = 8, 32	0.862
4	model #4 cycle = 1-3, pad = 8k, 32k	0.815
5	model #5 cycle = 1-5, pad = 8k, 32k	0.818
6	model #6 cycle = 5-10, pad = 8k, 32k	0.820
7	#1 + #3	0.876
8	#1 + #2 + #3	0.890
9	#4 + #5 + #6	0.836
10	submission 1	0.896
11	submission 2	0.895

9. REFERENCES

- [1] E. Fonseca, M. Plakal, F. Font, D. P. W. Ellis, and X. Serra, "Audio tagging with noisy labels and minimal supervision," *arXiv preprint arXiv:1906.02975*, 2019.
- [2] S. Ruder, "An Overview of Multi-Task Learning in Deep Neural Networks," *arXiv preprint arXiv:1706.05098*, 2017.
- [3] A. Tarvainen, and H. Valpola. "Mean teachers are better role models: Weight-averaged consistency targets improve semi-supervised deep learning results," *arXiv preprint arXiv:1703.01780*, 2017.
- [4] D. Berthelot, N. Carlini, I. Goodfellow, N. Papernot, A. Oliver, and C. Raffel, "MixMatch: A Holistic Approach to Semi-Supervised Learning," *arXiv preprint arXiv:1905.02249*, 2019.
- [5] D.-H. Lee, "Pseudo-label: The simple and efficient semi-supervised learning method for deep neural networks," in *ICML Workshop on Challenges in Representation Learning*, 2013.
- [6] G. Huang, Y. Li, G. Pleiss, Z. Liu, J. E. Hopcroft, and K. Q. Weinberger, "Snapshot Ensembles: Train 1, get M for free," *arXiv preprint arXiv:1704.00109*, 2017.
- [7] H. Zhang, M. Cisse, Y. N. Dauphin, and D. Lopez-Paz, "mixup: Beyond empirical risk minimization," *arXiv preprint arXiv:1710.09412*, 2017.
- [8] Y. Tokozume, Y. Ushiku, and T Harada, "Learning from Between-class Examples for Deep Sound Recognition," *arXiv preprint arXiv:1711.10282*, 2017.
- [9] D. S. Park, W. Chan, Y. Zhang, C.-C. Chiu, B. Zoph, E. D. Cubuk, and Q. V. Le, "SpecAugment: A Simple Data Augmentation Method for Automatic Speech Recognition," *arXiv preprint arXiv:1904.08779*, 2019.
- [10] T. DeVries, and G. W. Taylor, "Improved Regularization of Convolutional Neural Networks with Cutout," *arXiv preprint arXiv:1708.04552*, 2017.
- [11] K. He, X. Zhang, S. Ren, and J. Sun, "Deep residual learning for image recognition," in *Proceedings of the IEEE conference on computer vision and pattern recognition*, 2016, pp. 770-778.
- [12] J. Hu, L. Shen, and G. Sun, "Squeeze-and-Excitation Networks," *arXiv preprint arXiv:1709.01507*, 2017.
- [13] D. P. Kingma, and J. Ba, "Adam: A method for stochastic optimization," *arXiv preprint arXiv:1412.6980*, 2014.
- [14] I. Loshchilov, F. Hutter, "SGDR: Stochastic Gradient Descent with Warm Restarts," *arXiv preprint arXiv:1608.03983*, 2016.

Freesound-Audio-Tagging-2019

This is repository of the 4th place solution of [kaggleFreesound Audio Tagging 2019 competition](#).

The discription of this solution is available at

<http://dcase.community/challenge2019/task-audio-tagging-results#Akiyama2019>

<https://www.kaggle.com/c/freesound-audio-tagging-2019/discussion/96440>

Requirements

- Python 3.6.6
- CUDA 10.0
- numpy (1.16.4)
- pandas (0.23.4)
- matplotlib (3.1.0)
- Pytorch (1.1.0)
- librosa (0.6.3)
- sci-kit learn (0.21.2)
- scipy (1.2.1)
- pretrainedmodels (0.7.4)

Download the [dataset](#) and place them in input/.

Unzip zip files and place them to train_curated/, train_noisy/, test/.

In case you use pretrained weights, download the [weights](#), unzip zipped weights and place them to models/resnet_model1/, models/resnet_model2/ and so on.

Training

Run src/preprocess.py.

Run src/train_model1.py.

Run src/get_pseudo_label.py.

Run src/train_model2.py .

Run src/train_model3.py .

Run src/train_model4_0.py.

Run src/train_model4.py.

Run src/train_model5.py.

Run src/train_model6_0.py.

Run src/train_model6.py.

Prediction

Run src/make_final_submission1.py. The submission file output/ submission1.csv will be generted.

Run src/make_final_submission2.py. . The submission file output/

submission2.csv will be generated.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import time, random, sys
import sklearn.metrics
from sklearn.model_selection import KFold

import torch
import torch.nn as nn
import torch.backends.cudnn as cudnn
import torch.optim as optim
from torch.utils.data import DataLoader

sys.path.append('.')
from utils import AverageMeter, cycle, CosineLR, MelDataset, calculate_per_class_lwlap
from models import ResNet

# set parameters
NUM_FOLD = 5
NUM_CLASS = 80
SEED = 42
NUM_EPOCH = 64*8
NUM_CYCLE = 64
BATCH_SIZE = 64
LR = [1e-3, 1e-6]
FOLD_LIST = [1, 2, 3, 4, 5]
SLICE_LENGTH = 512
LOAD_DIR = "../models/resnet_model1"
OUTPUT_DIR = "../input/pseudo_label"

cudnn.benchmark = True
starttime = time.time()

def main():
    # load table data
    df_train = pd.read_csv("../input/train_curated.csv")
    df_noisy = pd.read_csv("../input/train_noisy.csv")
    df_test = pd.read_csv("../input/sample_submission.csv")
    labels = df_test.columns[1:].tolist()
    for label in labels:
        df_train[label] = df_train['labels'].apply(lambda x: label in x)
        df_noisy[label] = df_noisy['labels'].apply(lambda x: label in x)

    df_train['path'] = "../input/mel128/train/" + df_train['filename']
    df_test['path'] = "../input/mel128/test/" + df_train['filename']
    df_noisy['path'] = "../input/mel128/noisy/" + df_noisy['filename']

    # fold splitting
    folds = list(KFold(n_splits=NUM_FOLD, shuffle=True, random_s
```

```

tate=SEED).split(np.arange(len(df_train))))

    # build model
    model = ResNet(NUM_CLASS).cuda()

    # set generator
    dataset_noisy = MelDataset(df_noisy['path'], df_noisy[labels
].values)
    noisy_loader = DataLoader(dataset_noisy, batch_size=1,
                             shuffle=False, num_workers=1, pin_
memory=True,
                             )

    # predict
    preds_noisy = np.zeros([NUM_FOLD, NUM_EPOCH//NUM_CYCLE, len(
df_noisy), NUM_CLASS], np.float32)
    for fold, (ids_train_split, ids_valid_split) in enumerate(fo
lds):
        for cycle in range(NUM_EPOCH//NUM_CYCLE):
            print("fold: {} cycle: {}, sec: {:.1f}".format(fold+
1, cycle+1, time.time()-starttime))
            model.load_state_dict(torch.load("{}weight_fold{}_
epoch{}.pth".format(
                LOAD_DIR, fold+1, NUM_CYCLE*(cycle+1))))
            preds_noisy[fold, cycle] = predict(noisy_loader, mod
el)

    np.save("{}preds_noisy.npy".format(OUTPUT_DIR), preds_n
oisy)

def predict(test_loader, model):
    sigmoid = nn.Sigmoid().cuda()

    # switch to eval mode
    model.eval()

    preds = np.zeros([0, NUM_CLASS], np.float32)
    for i, (input, target) in enumerate(test_loader):
        input = torch.autograd.Variable(input.cuda())

        # compute output
        with torch.no_grad():
            output = model(input)
            pred = sigmoid(output)
            pred = pred.data.cpu().numpy()

        # measure accuracy and record loss
        preds = np.concatenate([preds, pred])
    return preds

if __name__ == '__main__':
    main()import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

```



```
import time, random, sys, os
import sklearn.metrics
from sklearn.model_selection import KFold
from multiprocessing import Pool
import gc
import shutil
from scipy.io import wavfile
import librosa
import concurrent.futures

import torch
import torch.nn as nn
import torch.backends.cudnn as cudnn
import torch.optim as optim
from torch.utils.data.dataset import Dataset
from torch.utils.data import DataLoader

sys.path.append('.')
from utils import AverageMeter, cycle, CosineLR, MelDataset, calculate_per_class_lwlap
from models import ResNet, EnvNetv2

# set parameters
NUM_FOLD = 5
NUM_CLASS = 80
SEED = 42
BATCH_DIR = "batch"
SIZE_LIMIT = 20000000
WIDTH_LIMIT = 80000
MAX_LEN = 1400000
MAX_PAD = 32000
MAX_BATCHSIZE = 512
LEN_DF_WAV_LIMIT = 500000000
LEN_DF_MEL_LIMIT = 2000000000
NUMBATCH_PER_NUMDATA = 1 / 55
MAX_PATIENCE = 0.2
wav_dir = "../input/test/"

RES_LIST = [
    {'dir': '../models/resnet_model1',
     'epoch': [2 * 64, 4 * 64, 7 * 64, 8 * 64],
     'pad': [8, 64],
    },
    {'dir': '../models/resnet_model2',
     'epoch': [1 * 64, 2 * 64, 6 * 64, 7 * 64],
     'pad': [8, 64],
    },
    {'dir': '../models/resnet_model3',
     'epoch': [2 * 64, 4 * 64, 6 * 64],
     'pad': [8, 64],
    },
]

ENV_LIST = [
    {'dir': '../models/envnet_model4',
     'epoch': [2 * 80, 3 * 80],
    },
]
```

```

        'pad': [0, 32000],
        'activation': 'sigmoid',
    },
    {'dir': '../models/envnet_model5',
     'epoch': [3 * 80, 5 * 80],
     'pad': [8000, 32000],
     'activation': 'softmax',
    },
    {'dir': '../models/envnet_model6',
     'epoch': [2 * 80, 4 * 80],
     'pad': [8000, 32000],
     'activation': 'softmax',
    },
]
LEN_RES_EPOCH = 0
for i in range(len(RES_LIST)):
    LEN_RES_EPOCH = max(LEN_RES_EPOCH, len(RES_LIST[i]['epoch']))
)
LEN_RES_PAD = 0
for i in range(len(RES_LIST)):
    LEN_RES_PAD = max(LEN_RES_PAD, len(RES_LIST[i]['pad']))
LEN_ENV_EPOCH = 0
for i in range(len(ENV_LIST)):
    LEN_ENV_EPOCH = max(LEN_ENV_EPOCH, len(ENV_LIST[i]['epoch']))
)
LEN_ENV_PAD = 0
for i in range(len(ENV_LIST)):
    LEN_ENV_PAD = max(LEN_ENV_EPOCH, len(ENV_LIST[i]['pad']))
starttime0 = time.time()

# cudnn speed up
cudnn.benchmark = True

def main():
    ### fix seed
    torch.manual_seed(SEED)
    random.seed(SEED)
    np.random.seed(SEED)
    torch.manual_seed(SEED)
    torch.cuda.manual_seed(SEED)

    # table data load
    starttime = time.time()
    df_test = pd.read_csv("../input/sample_submission.csv")
    labels = df_test.columns[1:].tolist()
    df_test['path'] = "{}/{}/".format(wav_dir) + df_test['fname']
    print("table data loading done. {:.1f}/{:.1f}".format(time.time() - starttime, time.time() - starttime0))

    # get data length
    starttime = time.time()
    p = Pool(2)
    len_list = p.map(get_len, df_test['path'].values)
    df_test['length'] = len_list
    print("getting data length done. {:.1f}/{:.1f}".format(time.time() - starttime, time.time() - starttime0))

```

```

time() - starttime, time.time() - starttime0))

    # data sort
    starttime = time.time()
    df_test_sort = df_test.copy()
    df_test_sort['index'] = np.arange(len(df_test_sort))
    df_test_sort = df_test_sort.sort_values(['length', 'index'])
    .reset_index(drop=True)
    print("data sort done. {:.1f}/{:.1f}".format(time.time() - s
tarttime, time.time() - starttime0))

    # batch splitting
    starttime = time.time()
    NUM_BATCH_LIMIT = 60 + int(len(df_test_sort) * NUMBATCH_PER_
NUMDATA)
    print("num batch limit: {}".format(NUM_BATCH_LIMIT))
    patience_rate = 0
    patience_rate_tmp = 0
    num_batch, count = get_num_batch(df_test_sort, patience_rate
)
    print("patience_rate_tmp: {:.2f}, patience_rate_tmp: {:.2f},
num_batch: {:3d}".format(
        patience_rate, patience_rate_tmp, num_batch))
    while num_batch > NUM_BATCH_LIMIT and patience_rate_tmp < MA
X_PATIENCE:
        patience_rate_tmp += 0.01
        num_batch_tmp, count_tmp = get_num_batch(df_test_sort, p
atience_rate_tmp)
        if num_batch_tmp < num_batch:
            num_batch = num_batch_tmp
            count = count_tmp
            patience_rate = patience_rate_tmp
        print("patience_rate_tmp: {:.2f}, patience_rate_tmp: {:.
2f}, num_batch_tmp: {:3d}".format(
            patience_rate, patience_rate_tmp, num_batch_tmp))
        num_batch, count = get_num_batch(df_test_sort, patience_rate
)
    print("num batch: {}, rate of padding patience: {:.2f}".form
at(num_batch, patience_rate))
    print("batch splitting done. {:.1f}/{:.1f}".format(time.time
() - starttime, time.time() - starttime0))

    # store batch id
    starttime = time.time()
    batch_list = []
    for i in range(num_batch):
        batch_list += [i] * count[i][1]
    df_test_sort['batch'] = batch_list
    print(df_test_sort[['path', 'length', 'batch']].head())
    print("save batch id done. {:.1f}/{:.1f}".format(time.time()
- starttime, time.time() - starttime0))

    # split dataframe if too big
    starttime = time.time()
    df_mel_split = get_df_split(df_test_sort, LEN_DF_MEL_LIMIT)
    print("df_mel_split")

```

```

    for i in range(len(df_mel_split)):
        print("{}: num data: {}, total length: {}".format(i + 1,
len(df_mel_split[i]), df_mel_split[i]['length'].sum()))
        print("dataframe splitting done. {:.1f}/{:.1f}".format(time.
time() - starttime, time.time() - starttime))

    # ### EnvNet part
    # build model
    model = EnvNetv2(NUM_CLASS).cuda()
    model.eval()

    # split df for EnvNet
    df_wav_split = get_df_split(df_test_sort, LEN_DF_WAV_LIMIT)
    print("df_wav_split")
    for i in range(len(df_wav_split)):
        print("{}: num data: {}, total length: {}".format(i + 1,
len(df_wav_split[i]), df_wav_split[i]['length'].sum()))

    print("predict wav...")

    # parallel threading
    executor = concurrent.futures.ThreadPoolExecutor(max_workers
=2)
    threadA = executor.submit(get_mel_batch, df_mel_split[0])
    threadB = executor.submit(predict_wav_split, model, df_wav_s
plit[0], ENV_LIST)
    preds_wav_split = []
    preds_wav_split.append(threadB.result())
    executor.shutdown()
    print("parallel threading done.", time.time() - starttime, t
ime.time() - starttime)

    # do remain EnvNet prediction
    if len(df_wav_split) > 1:
        for split in range(1, len(df_wav_split)):
            preds_wav_split.append(predict_wav_split(model, df_w
av_split[split], ENV_LIST))
            print("envnet prediction split {}/{}, done. {:.1f}/{
:.1f}".format(
                split + 1, len(df_wav_split), time.time() - star
ttime, time.time() - starttime))
            preds_test_wav = np.concatenate(preds_wav_split, axis=4)
            print("all envnet predict done.", time.time() - starttime, t
ime.time() - starttime)

    # build model
    starttime = time.time()
    model = ResNet(NUM_CLASS).cuda()
    model.eval()
    print("building ResNet model done. {:.1f}/{:.1f}".format(tim
e.time() - starttime, time.time() - starttime))

    # predict split #1
    preds_test_mel = []
    preds_test_mel.append(predict_mel_split(model, df_mel_split[
0], RES_LIST))

```

```

shutil.rmtree(BATCH_DIR)
print("mel prediction of split {} done. {:.1f}/{:.1f}".format(
    1, time.time() - starttime, time.time() - starttime0))

# process remain split
if len(df_mel_split) > 1:
    for split in range(1, len(df_mel_split)):
        # mel preprocessing
        starttime = time.time()
        df_test_sort_tmp = df_mel_split[split]
        get_mel_batch(df_test_sort_tmp)
        print("mel preprocessing of split {} done. {:.1f}/{:.1f}".format(
            split + 1, time.time() - starttime, time.time() - starttime0))
        preds_test_mel.append(predict_mel_split(model, df_test_sort_tmp, RES_LIST))
        shutil.rmtree(BATCH_DIR)
        print("mel prediction of split {} done. {:.1f}/{:.1f}".format(
            split + 1, time.time() - starttime, time.time() - starttime0))

    print("all prediction done. {:.1f}/{:.1f}".format(time.time() - starttime, time.time() - starttime0))

# concat
starttime = time.time()
preds_test_mel = np.concatenate(preds_test_mel, axis=4)
print("preds_test_mel.shape", preds_test_mel.shape)
print("concat done.", time.time() - starttime, time.time() - starttime0)

# make submission
preds_test_avr = (
    + preds_test_mel[:, 0, :len(RES_LIST[0]['epoch'])].mean(axis=(0, 1, 2)) * 4 / 13
    + preds_test_mel[:, 1, :len(RES_LIST[1]['epoch'])].mean(axis=(0, 1, 2)) * 3 / 13
    + preds_test_mel[:, 2, :len(RES_LIST[2]['epoch'])].mean(axis=(0, 1, 2)) * 3 / 13
    + preds_test_wav[:, 0, :len(ENV_LIST[0]['epoch'])].mean(axis=(0, 1, 2)) * 1 / 13
    + preds_test_wav[:, 1, :len(ENV_LIST[1]['epoch'])].mean(axis=(0, 1, 2)) * 1 / 13
    + preds_test_wav[:, 2, :len(ENV_LIST[2]['epoch'])].mean(axis=(0, 1, 2)) * 1 / 13)
print(preds_test_mel.shape, preds_test_wav.shape)
print(preds_test_avr.shape)
df_test_sort = df_test_sort.sort_values(['length', 'index']).reset_index(drop=True)
df_test_sort[labels] = preds_test_avr
df_test_sort = df_test_sort.sort_values('index').reset_index(drop=True)
df_test_sort[['fname'] + labels].to_csv("../output/submission1.csv", index=None)

```

```

    print("save submission done. {:.1f}/{:.1f}".format(time.time()
() - starttime, time.time() - starttime0))

def check_size_limit(width, num_batch):
    if (width + MAX_PAD * 2) * num_batch > SIZE_LIMIT:
        return True
    else:
        return False

def get_num_batch(df_test_sort, patience_rate):
    i = 0
    len_now = df_test_sort['length'][i]
    patience = int(len_now * patience_rate)
    count = []
    while (i < len(df_test_sort)):
        len_now = df_test_sort['length'][i]
        patience = int(len_now * patience_rate)
        if len(count) == 0 or count[-1][0] + patience < len_now
or count[-1][1] >= MAX_BATCHSIZE:
            count.append([len_now, 1])
        elif check_size_limit(len_now, count[-1][1] + 1):
            count.append([len_now, 1])
        else:
            count[-1][1] += 1
        i += 1
    return len(count), count

def predict_mel_split(model, df_split, RES_LIST):
    starttime = time.time()
    batch_idx = [df_split['batch'].min(), df_split['batch'].max(
) + 1]
    preds_test_mel_tmp = np.zeros([
        NUM_FOLD,
        len(RES_LIST),
        len(RES_LIST[0]['epoch']),
        len(RES_LIST[0]['pad']),
        len(df_split), NUM_CLASS], np.float32)

    dataset_valid = BatchDataset(df_split, 0)
    valid_loader = DataLoader(dataset_valid,
                               batch_size=1,
                               shuffle=False,
                               num_workers=1,
                               pin_memory=True,
                               collate_fn=my_collate
                               )
    for i in range(len(RES_LIST)):
        model_dir = RES_LIST[i]['dir']
        epoch_list = RES_LIST[i]['epoch']
        pad_list = RES_LIST[i]['pad']
        for fold in range(NUM_FOLD):
            for k, epoch in enumerate(epoch_list):
                model.load_state_dict(

```

```

        torch.load("{}weight_fold{}_epoch{}.pth".
format(model_dir, fold + 1, epoch))
        for j, pad in enumerate(pad_list):
            print("fold: {}, dir: {}, epoch: {}, pad: {}
, sec: {:.1f}".format(
                fold + 1, model_dir, epoch, pad, time.ti
me() - starttime))
                dataset_valid.pad = pad
                preds_test_mel_tmp[fold, i, k, j] = predict_
resnet(model, valid_loader)
            return preds_test_mel_tmp

def get_mel_batch(df_split):
    df_split['path'] = "{}{}".format(wav_dir) + df_split['fname']
    print(df_split[['path', 'batch']])
    os.makedirs(BATCH_DIR, exist_ok=True)
    p = Pool(2) # 2
    batch_idx = [df_split['batch'].min(), df_split['batch'].max(
) + 1]
    for i in range(batch_idx[0], batch_idx[1]):
        df_tmp = df_split[df_split['batch'] == i].reset_index(dr
op=True)
        args = []
        slice = int(np.ceil((df_tmp['length'].values[-1] + 1) /
347))
        slice = np.min([slice, WIDTH_LIMIT])
        for j in range(len(df_tmp)):
            args.append([df_tmp['path'][j], slice])
        batch = p.map(preprocess_mel, args)
        batch = np.array(batch)
        np.save("{}{}.np".format(BATCH_DIR, i), batch)

def predict_wav_split(model, df, ENV_LIST):
    starttime = time.time()
    batch_idx = [df['batch'].min(), df['batch'].max() + 1]
    p = Pool(2) # 2
    batch_list = []
    for i in range(batch_idx[0], batch_idx[1]):
        df_tmp = df[df['batch'] == i].reset_index(drop=True)
        args = []
        slice = df_tmp['length'].values[-1]
        for j in range(len(df_tmp)):
            args.append([df_tmp['path'][j], slice])
        batch = p.map(preprocess_wav, args)
        batch = np.array(batch)
        batch_list.append(batch)
    print("batch making done, sec: {:.1f}".format(time.time() -
starttime))

    # envnet predict
    starttime = time.time()
    print("predict valid...")
    preds_test_wav = np.zeros([
        NUM_FOLD,

```

```

        len(ENV_LIST),
        len(ENV_LIST[0]['epoch']),
        len(ENV_LIST[0]['pad']),
        len(df), NUM_CLASS], np.float32)

dataset_valid = BatchWavDataset(batch_list, 0)
valid_loader = DataLoader(dataset_valid,
                           batch_size=1,
                           shuffle=False,
                           num_workers=1,
                           pin_memory=True,
                           collate_fn=my_collate
                           )
for i in range(len(ENV_LIST)):
    model_dir = ENV_LIST[i]['dir']
    epoch_list = ENV_LIST[i]['epoch']
    pad_list = ENV_LIST[i]['pad']
    activation = ENV_LIST[i]['activation']

    for fold in range(NUM_FOLD):
        for k, epoch in enumerate(epoch_list):
            model.load_state_dict(
                torch.load("{}weight_fold{}_epoch{}.pth".
format(model_dir, fold + 1, epoch),
                           map_location='cuda:0'))
            for j, pad in enumerate(pad_list):
                print("fold: {}, dir: {}, epoch: {}, pad: {}
, sec: {:.1f}".format(
                    fold + 1, model_dir, epoch, pad, time.ti
me() - starttime))
                dataset_valid.pad = pad
                preds_test_wav[fold, i, k, j] = predict_envn
et(model, valid_loader, activation)
            return preds_test_wav

def get_df_split(df, size_limit):
    num_batch = df['batch'].max() + 1
    sum_len = df['length'].sum()
    df_split = []
    begin = 0
    sum_tmp = 0
    print("base df shape", df.shape)
    for i in range(num_batch):
        sum_tmp += df['length'][df['batch'] == i].sum()
        if sum_tmp > size_limit:
            df_split.append(
                df[(begin <= df['batch']) & (df['batch'] < i)].r
eset_index(drop=True))
            sum_tmp = df['length'][df['batch'] == i].sum()
            begin = i
        df_split.append(df[(begin <= df['batch'])].reset_index(drop=
True))
    return df_split

```



```

def my_collate(batch):
    return torch.Tensor(batch[0])

def get_len(path):
    _, data = wavfile.read(path)
    len_data = len(data)
    if len(data) > MAX_LEN:
        len_data = MAX_LEN
        print("File length {} is too long! This file is sliced to {}".format(len(data), MAX_LEN))

    return len_data

def get_wav(path):
    _, snd = wavfile.read(path)
    return snd

def get_mel(wave):
    wave = wave.astype(np.float32) / 32768.0
    data = librosa.feature.melspectrogram(
        wave,
        sr=44100,
        n_mels=128,
        hop_length=347 * 1,
        n_fft=128 * 20,
        fmin=20,
        fmax=44100 // 2,
    ).astype(np.float32)
    return data

def preprocess_mel(args):
    path, slice = args
    wav = get_wav(path)
    mel = get_mel(wav)
    mel_new = np.zeros([mel.shape[0], slice], np.float32)
    if mel.shape[1] > slice:
        print("wav length: {}, mel length: {}".format(wav.shape[0], mel.shape[1]))
        print("Mel file is sliced")
        mel_new[:] = mel[:, :slice]
    else:
        mel_new[:, :mel.shape[1]] = mel
    mel_new = librosa.power_to_db(mel_new)
    mel_new = mel_new.reshape([1, mel_new.shape[0], mel_new.shape[1]])
    return mel_new

def preprocess_wav(args):
    path, slice = args
    wav = get_wav(path)
    pad = (slice - len(wav)) // 2

```

```

wav_new = np.zeros([1, 1, slice], np.int16)

if wav.shape[0] > slice:
    print("wav length: {}".format(wav.shape[0]))
    print("Wav file is sliced")
    wav_new[0, 0, :] = wav[:slice]
else:
    wav_new[0, 0, pad:pad + len(wav)] = wav
return wav_new

class BatchDataset(Dataset):
    def __init__(self, df, pad=0):
        self.len_batch = df['batch'].max() - df['batch'].min() +
1
        self.X = np.arange(self.len_batch) + df['batch'].min()
        self.pad = pad

    def __getitem__(self, index):
        batch_base = np.load("{}/{}.npy".format(BATCH_DIR, self.
X[index]))
        batch_pad = np.zeros(batch_base.shape[:-1] + (batch_base
.shape[-1] + self.pad * 2,), np.float32)
        batch_max = batch_base.max(axis=(1, 2, 3)) - 80
        batch_max = np.maximum(batch_max, -100)
        batch_pad[:, :] = batch_max[:, np.newaxis, np.newaxis, np.n
ewaxis, ]
        if self.pad != 0:
            batch_pad[:, :, :, self.pad:-self.pad] = batch_base
        else:
            batch_pad[:, :] = batch_base
        batch_pad = (batch_pad - batch_pad.mean(axis=(1, 2, 3))[:
, np.newaxis, np.newaxis, np.newaxis, ]) / (
            batch_pad.std(axis=(1, 2, 3))[:, np.newaxis, np.
newaxis, np.newaxis, ] + 1e-7)
        return batch_pad

    def __len__(self):
        return self.len_batch

class BatchWavDataset(Dataset):
    def __init__(self, batch_list, pad=0):
        self.X = batch_list
        self.pad = pad

    def __getitem__(self, index):
        batch_base = self.X[index]
        if batch_base.shape[-1] + self.pad * 2 < 20580:
            pad = int(np.ceil((20580 - batch_base.shape[-1]) / 2
))
        else:
            pad = self.pad
        if pad != 0:
            batch = np.zeros([batch_base.shape[0], 1, 1, batch_b
ase.shape[-1] + pad * 2], np.float32)

```

```

        batch[:, :, :, pad:-pad] = batch_base.astype(np.float32) / 32768.0
    else:
        batch = batch_base.astype(np.float32) / 32768.0
    return batch

def __len__(self):
    return len(self.X)

def predict_resnet(model, dataloader):
    sigmoid = torch.nn.Sigmoid().cuda()
    preds = []
    model.eval()
    for i, input in enumerate(dataloader):
        input = input.cuda(async=True)
        with torch.no_grad():
            pred = sigmoid(model(input)).data.cpu().numpy()
        preds.append(pred)
    preds = np.concatenate(preds)
    return preds

def predict_envnet(model, dataloader, activation='sigmoid'):
    sigmoid = torch.nn.Sigmoid().cuda()
    softmax = torch.nn.Softmax(dim=1).cuda()
    if activation == 'sigmoid':
        f_act = sigmoid
    elif activation == 'softmax':
        f_act = softmax
    preds = []
    model.eval()
    for i, input in enumerate(dataloader):
        input = input.cuda(async=True)
        with torch.no_grad():
            pred = f_act(model(input)).data.cpu().numpy()
        preds.append(pred)
    preds = np.concatenate(preds)
    return preds

if __name__ == '__main__':
    main()
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import time, random, sys, os
import sklearn.metrics
from sklearn.model_selection import KFold
from multiprocessing import Pool
import gc
import shutil
from scipy.io import wavfile
import librosa
import concurrent.futures

import torch

```

```
import torch.nn as nn
import torch.backends.cudnn as cudnn
import torch.optim as optim
from torch.utils.data.dataset import Dataset
from torch.utils.data import DataLoader

sys.path.append('.')
from utils import AverageMeter, cycle, CosineLR, MelDataset, calculate_per_class_lwlrp
from models import ResNet, EnvNetv2

# set parameters
NUM_FOLD = 5
NUM_CLASS = 80
SEED = 42
BATCH_DIR = "batch"
SIZE_LIMIT = 20000000
WIDTH_LIMIT = 80000
MAX_LEN = 1400000
MAX_PAD = 32000
MAX_BATCHSIZE = 512
LEN_DF_WAV_LIMIT = 500000000
LEN_DF_MEL_LIMIT = 2000000000
NUMBATCH_PER_NUMDATA = 1 / 30
MAX_PATIENCE = 0.2
wav_dir = "../input/test/"

RES_LIST = [
    {'dir': '../models/resnet_model1',
     'epoch': [2*64, 4*64, 7*64, 8*64],
     'pad': [8],
    },
    {'dir': '../models/resnet_model2',
     'epoch': [1*64, 2*64, 4*64, 6*64, 7*64],
     'pad': [8],
    },
    {'dir': '../models/resnet_model3',
     'epoch': [2*64, 4*64, 6*64],
     'pad': [8],
    },
]

ENV_LIST = [
    {'dir': '../models/envnet_model4',
     'epoch': [2*80, 3*80],
     'pad': [8000],
     'activation': 'sigmoid',
    },
    {'dir': '../models/envnet_model5',
     'epoch': [3 * 80, 5 * 80],
     'pad': [8000],
     'activation': 'softmax',
    },
    {'dir': '../models/envnet_model6',
     'epoch': [1*80, 2*80, 4*80],
     'pad': [8000],
    },
]
```

```

        'activation': 'softmax',
    },
]
LEN_RES_EPOCH = 0
for i in range(len(RES_LIST)):
    LEN_RES_EPOCH = max(LEN_RES_EPOCH, len(RES_LIST[i]['epoch']))
)
LEN_RES_PAD = 0
for i in range(len(RES_LIST)):
    LEN_RES_PAD = max(LEN_RES_PAD, len(RES_LIST[i]['pad']))
LEN_ENV_EPOCH = 0
for i in range(len(ENV_LIST)):
    LEN_ENV_EPOCH = max(LEN_ENV_EPOCH, len(ENV_LIST[i]['epoch']))
)
LEN_ENV_PAD = 0
for i in range(len(ENV_LIST)):
    LEN_ENV_PAD = max(LEN_ENV_EPOCH, len(ENV_LIST[i]['pad']))
starttime0 = time.time()

# cudnn speed up
cudnn.benchmark = True

def main():
    ### fix seed
    torch.manual_seed(SEED)
    random.seed(SEED)
    np.random.seed(SEED)
    torch.manual_seed(SEED)
    torch.cuda.manual_seed(SEED)

    # table data load
    starttime = time.time()
    df_test = pd.read_csv("../input/sample_submission.csv")
    labels = df_test.columns[1:].tolist()
    df_test['path'] = "{}{}".format(wav_dir) + df_test['fname']
    print("table data loading done. {:.1f}/{:.1f}".format(time.time() - starttime, time.time() - starttime0))

    # get data length
    starttime = time.time()
    p = Pool(2)
    len_list = p.map(get_len, df_test['path'].values)
    df_test['length'] = len_list
    print("getting data length done. {:.1f}/{:.1f}".format(time.time() - starttime, time.time() - starttime0))

    # data sort
    starttime = time.time()
    df_test_sort = df_test.copy()
    df_test_sort['index'] = np.arange(len(df_test_sort))
    df_test_sort = df_test_sort.sort_values(['length', 'index']).reset_index(drop=True)
    print("data sort done. {:.1f}/{:.1f}".format(time.time() - starttime, time.time() - starttime0))

```

```

    # batch splitting
    starttime = time.time()
    NUM_BATCH_LIMIT = 50 + int(len(df_test_sort)*NUMBATCH_PER_NU
MDATA)
    print("num batch limit: {}".format(NUM_BATCH_LIMIT))
    patience_rate = 0
    patience_rate_tmp = 0
    num_batch, count = get_num_batch(df_test_sort, patience_rate
)
    print("patience_rate_tmp: {:.2f}, patience_rate_tmp: {:.2f},
num_batch: {:3d}".format(
        patience_rate, patience_rate_tmp, num_batch))
    while num_batch > NUM_BATCH_LIMIT and patience_rate_tmp < MA
X_PATIENCE:
        patience_rate_tmp += 0.01
        num_batch_tmp, count_tmp = get_num_batch(df_test_sort, p
atience_rate_tmp)
        if num_batch_tmp < num_batch:
            num_batch = num_batch_tmp
            count = count_tmp
            patience_rate = patience_rate_tmp
        print("patience_rate_tmp: {:.2f}, patience_rate_tmp: {:.
2f}, num_batch_tmp: {:3d}".format(
            patience_rate, patience_rate_tmp, num_batch_tmp))
        num_batch, count = get_num_batch(df_test_sort, patience_rate
)
    print("num batch: {}, rate of padding patience: {:.2f}".form
at(num_batch, patience_rate))
    print("batch splitting done. {:.1f}/{:.1f}".format(time.time
()) - starttime, time.time() - starttime0))

    # store batch id
    starttime = time.time()
    batch_list = []
    for i in range(num_batch):
        batch_list += [i] * count[i][1]
    df_test_sort['batch'] = batch_list
    print(df_test_sort[['path', 'length', 'batch']].head())
    print("save batch id done. {:.1f}/{:.1f}".format(time.time()
- starttime, time.time() - starttime0))

    # split dataframe if too big
    starttime = time.time()
    df_mel_split = get_df_split(df_test_sort, LEN_DF_MEL_LIMIT)
    print("df_mel_split")
    for i in range(len(df_mel_split)):
        print("{}: num data: {}, total length: {}".format(i + 1,
len(df_mel_split[i]), df_mel_split[i]['length'].sum()))
    print("dataframe splitting done. {:.1f}/{:.1f}".format(time.
time() - starttime, time.time() - starttime0))

    # ### EnvNet part
    # build model
    model = EnvNetv2(NUM_CLASS).cuda()
    model.eval()

```

```

# split df for EnvNet
df_wav_split = get_df_split(df_test_sort, LEN_DF_WAV_LIMIT)
print("df_wav_split")
for i in range(len(df_wav_split)):
    print("{}: num data: {}, total length: {}".format(i + 1,
len(df_wav_split[i]), df_wav_split[i]['length'].sum()))

print("predict wav...")

# parallel threading
executor = concurrent.futures.ThreadPoolExecutor(max_workers
=2)
threadA = executor.submit(get_mel_batch, df_mel_split[0])
threadB = executor.submit(predict_wav_split, model, df_wav_s
plit[0], ENV_LIST)
preds_wav_split = []
preds_wav_split.append(threadB.result())
executor.shutdown()
print("parallel threading done.", time.time() - starttime, t
ime.time() - starttime0)

# do remain EnvNet prediction
if len(df_wav_split) > 1:
    for split in range(1, len(df_wav_split)):
        preds_wav_split.append(predict_wav_split(model, df_w
av_split[split], ENV_LIST))
        print("envnet prediction split {}/{}, done. {:.1f}/{
:.1f}".format(
            split + 1, len(df_wav_split), time.time() - star
ttime, time.time() - starttime0))
        preds_test_wav = np.concatenate(preds_wav_split, axis=4)
        print("all envnet predict done.", time.time() - starttime, t
ime.time() - starttime0)

# build model
starttime = time.time()
model = ResNet(NUM_CLASS).cuda()
model.eval()
print("building ResNet model done. {:.1f}/{:.1f}".format(tim
e.time() - starttime, time.time() - starttime0))

# predict split #1
preds_test_mel = []
preds_test_mel.append(predict_mel_split(model, df_mel_split[
0], RES_LIST))
shutil.rmtree(BATCH_DIR)
print("mel prediction of split {} done. {:.1f}/{:.1f}".forma
t(1, time.time() - starttime, time.time() - starttime0))

# process remain split
if len(df_mel_split) > 1:
    for split in range(1, len(df_mel_split)):
        # mel preprocessing
        starttime = time.time()
        df_test_sort_tmp = df_mel_split[split]
        get_mel_batch(df_test_sort_tmp)

```

```

        print("mel preprocessing of split {} done. {:.1f}/{:.1f}").format(
            split + 1, time.time() - starttime, time.time()
            - starttime0))
        preds_test_mel.append(predict_mel_split(model, df_test_sort_tmp, RES_LIST))
        shutil.rmtree(BATCH_DIR)
        print("mel prediction of split {} done. {:.1f}/{:.1f}").format(
            split + 1, time.time() - starttime, time.time()
            - starttime0))

    print("all prediction done. {:.1f}/{:.1f}").format(time.time()
        - starttime, time.time() - starttime0))

    # concat
    starttime = time.time()
    preds_test_mel = np.concatenate(preds_test_mel, axis=4)
    print("preds_test_mel.shape", preds_test_mel.shape)
    print("concat done.", time.time() - starttime, time.time() -
        starttime0)

    # make submission
    preds_test_avr = (
        + preds_test_mel[:, 0, :len(RES_LIST[0]['epoch'])].mean(
            axis=(0, 1, 2)) * 4 / 13
        + preds_test_mel[:, 1, :len(RES_LIST[1]['epoch'])].mean(
            axis=(0, 1, 2)) * 3 / 13
        + preds_test_mel[:, 2, :len(RES_LIST[2]['epoch'])].mean(
            axis=(0, 1, 2)) * 3 / 13
        + preds_test_wav[:, 0, :len(ENV_LIST[0]['epoch'])].mean(
            axis=(0, 1, 2)) * 1 / 13
        + preds_test_wav[:, 1, :len(ENV_LIST[1]['epoch'])].mean(
            axis=(0, 1, 2)) * 1 / 13
        + preds_test_wav[:, 2, :len(ENV_LIST[2]['epoch'])].mean(
            axis=(0, 1, 2)) * 1 / 13)
    print(preds_test_mel.shape, preds_test_wav.shape)
    print(preds_test_avr.shape)
    df_test_sort = df_test_sort.sort_values(['length', 'index'])
    .reset_index(drop=True)
    df_test_sort[labels] = preds_test_avr
    df_test_sort = df_test_sort.sort_values('index').reset_index(
        drop=True)
    df_test_sort[['fname'] + labels].to_csv("../output/submission1.csv",
        index=None)
    print("save submission done. {:.1f}/{:.1f}").format(time.time()
        - starttime, time.time() - starttime0))

def check_size_limit(width, num_batch):
    if (width + MAX_PAD * 2) * num_batch > SIZE_LIMIT:
        return True
    else:
        return False

```



```

def get_num_batch(df_test_sort, patience_rate):
    i = 0
    len_now = df_test_sort['length'][i]
    patience = int(len_now * patience_rate)
    count = []
    while (i < len(df_test_sort)):
        len_now = df_test_sort['length'][i]
        patience = int(len_now * patience_rate)
        if len(count) == 0 or count[-1][0] + patience < len_now
or count[-1][1] >= MAX_BATCHSIZE:
            count.append([len_now, 1])
        elif check_size_limit(len_now, count[-1][1] + 1):
            count.append([len_now, 1])
        else:
            count[-1][1] += 1
        i += 1
    return len(count), count

def predict_mel_split(model, df_split, RES_LIST):
    starttime = time.time()
    batch_idx = [df_split['batch'].min(), df_split['batch'].max(
) + 1]
    preds_test_mel_tmp = np.zeros([
        NUM_FOLD,
        len(RES_LIST),
        len(RES_LIST[0]['epoch']),
        len(RES_LIST[0]['pad']),
        len(df_split), NUM_CLASS], np.float32)

    dataset_valid = BatchDataset(df_split, 0)
    valid_loader = DataLoader(dataset_valid,
                              batch_size=1,
                              shuffle=False,
                              num_workers=1,
                              pin_memory=True,
                              collate_fn=my_collate
                              )
    for i in range(len(RES_LIST)):
        model_dir = RES_LIST[i]['dir']
        epoch_list = RES_LIST[i]['epoch']
        pad_list = RES_LIST[i]['pad']
        for fold in range(NUM_FOLD):
            for k, epoch in enumerate(epoch_list):
                model.load_state_dict(
                    torch.load("{}weight_fold{}_epoch{}.pth".
format(model_dir, fold + 1, epoch)))
                for j, pad in enumerate(pad_list):
                    print("fold: {}, dir: {}, epoch: {}, pad: {}
, sec: {:.1f}".format(
                        fold + 1, model_dir, epoch, pad, time.ti
me() - starttime))
                    dataset_valid.pad = pad
                    preds_test_mel_tmp[fold, i, k, j] = predict_
resnet(model, valid_loader)
    return preds_test_mel_tmp

```

```

def get_mel_batch(df_split):
    print(1)
    df_split['path'] = "{}{}".format(wav_dir) + df_split['fname']
    print(df_split[['path', 'batch']])
    os.makedirs(BATCH_DIR, exist_ok=True)
    print(2)
    p = Pool(2) # 2=2
    batch_idx = [df_split['batch'].min(), df_split['batch'].max(
) + 1]
    for i in range(batch_idx[0], batch_idx[1]):
        df_tmp = df_split[df_split['batch'] == i].reset_index(drop=True)
        args = []
        slice = int(np.ceil((df_tmp['length'].values[-1] + 1) /
347))

        slice = np.min([slice, WIDTH_LIMIT])
        for j in range(len(df_tmp)):
            args.append([df_tmp['path'][j], slice])
        batch = p.map(preprocess_mel, args)
        batch = np.array(batch)
        np.save("{}{}.np".format(BATCH_DIR, i), batch)

def predict_wav_split(model, df, ENV_LIST):
    starttime = time.time()
    batch_idx = [df['batch'].min(), df['batch'].max() + 1]
    p = Pool(2) # 2=2
    batch_list = []
    for i in range(batch_idx[0], batch_idx[1]):
        df_tmp = df[df['batch'] == i].reset_index(drop=True)
        args = []
        slice = df_tmp['length'].values[-1]
        for j in range(len(df_tmp)):
            args.append([df_tmp['path'][j], slice])
        batch = p.map(preprocess_wav, args)
        batch = np.array(batch)
        batch_list.append(batch)
    print("batch making done, sec: {:.1f}".format(time.time() -
starttime))

    # envnet predict
    starttime = time.time()
    print("predict valid...")
    preds_test_wav = np.zeros([
        NUM_FOLD,
        len(ENV_LIST),
        len(ENV_LIST[0]['epoch']),
        len(ENV_LIST[0]['pad']),
        len(df), NUM_CLASS], np.float32)

    dataset_valid = BatchWavDataset(batch_list, 0)
    valid_loader = DataLoader(dataset_valid,
                                batch_size=1,
                                shuffle=False,

```

```

        num_workers=1,
        pin_memory=True,
        collate_fn=my_collate
    )
    for i in range(len(ENV_LIST)):
        model_dir = ENV_LIST[i]['dir']
        epoch_list = ENV_LIST[i]['epoch']
        pad_list = ENV_LIST[i]['pad']
        activation = ENV_LIST[i]['activation']

        for fold in range(NUM_FOLD):
            for k, epoch in enumerate(epoch_list):
                model.load_state_dict(
                    torch.load("{}weight_fold{}_epoch{}.pth".
format(model_dir, fold + 1, epoch),
                        map_location='cuda:0'))
                for j, pad in enumerate(pad_list):
                    print("fold: {}, dir: {}, epoch: {}, pad: {}
, sec: {:.1f}".format(
                        fold + 1, model_dir, epoch, pad, time.ti
me() - starttime))
                    dataset_valid.pad = pad
                    preds_test_wav[fold, i, k, j] = predict_envn
et(model, valid_loader, activation)
                return preds_test_wav

def get_df_split(df, size_limit):
    num_batch = df['batch'].max() + 1
    sum_len = df['length'].sum()
    df_split = []
    begin = 0
    sum_tmp = 0
    print("base df shape", df.shape)
    for i in range(num_batch):
        sum_tmp += df['length'][df['batch'] == i].sum()
        if sum_tmp > size_limit:
            df_split.append(
                df[(begin <= df['batch']) & (df['batch'] < i)].r
eset_index(drop=True))
            sum_tmp = df['length'][df['batch'] == i].sum()
            begin = i
        df_split.append(df[(begin <= df['batch'])].reset_index(drop=
True))
    return df_split

def my_collate(batch):
    return torch.Tensor(batch[0])

def get_len(path):
    _, data = wavfile.read(path)
    len_data = len(data)
    if len(data) > MAX_LEN:
        len_data = MAX_LEN

```

```

        print("File length {} is too long! This file is sliced t
o {}".format(len(data), MAX_LEN))

    return len_data

def get_wav(path):
    _, snd = wavfile.read(path)
    return snd

def get_mel(wave):
    wave = wave.astype(np.float32) / 32768.0
    data = librosa.feature.melspectrogram(
        wave,
        sr=44100,
        n_mels=128,
        hop_length=347 * 1,
        n_fft=128 * 20,
        fmin=20,
        fmax=44100 // 2,
    ).astype(np.float32)
    return data

def preprocess_mel(args):
    path, slice = args
    wav = get_wav(path)
    mel = get_mel(wav)
    mel_new = np.zeros([mel.shape[0], slice], np.float32)
    if mel.shape[1] > slice:
        print("wav length: {}, mel length: {}".format(wav.shape[
0], mel.shape[1]))
        print("Mel file is sliced")
        mel_new[:, :] = mel[:, :slice]
    else:
        mel_new[:, :mel.shape[1]] = mel
    mel_new = librosa.power_to_db(mel_new)
    mel_new = mel_new.reshape([1, mel_new.shape[0], mel_new.shap
e[1]])
    return mel_new

def preprocess_wav(args):
    path, slice = args
    wav = get_wav(path)
    pad = (slice - len(wav)) // 2
    wav_new = np.zeros([1, 1, slice], np.int16)

    if wav.shape[0] > slice:
        print("wav length: {}".format(wav.shape[0]))
        print("Wav file is sliced")
        wav_new[0, 0, :] = wav[:slice]
    else:
        wav_new[0, 0, pad:pad + len(wav)] = wav
    return wav_new

```

```

class BatchDataset(Dataset):
    def __init__(self, df, pad=0):
        self.len_batch = df['batch'].max() - df['batch'].min() +
1
        self.X = np.arange(self.len_batch) + df['batch'].min()
        self.pad = pad

    def __getitem__(self, index):
        batch_base = np.load("{}/{:.npy".format(BATCH_DIR, self.
X[index]))
        batch_pad = np.zeros(batch_base.shape[:-1] + (batch_base
.shape[-1] + self.pad * 2,), np.float32)
        batch_max = batch_base.max(axis=(1, 2, 3)) - 80
        batch_max = np.maximum(batch_max, -100)
        batch_pad[:, :, :, np.newaxis, np.newaxis, np.n
ewaxis, ]
        if self.pad != 0:
            batch_pad[:, :, :, self.pad:-self.pad] = batch_base
        else:
            batch_pad[:, :, :, ] = batch_base
        batch_pad = (batch_pad - batch_pad.mean(axis=(1, 2, 3)) [
:, np.newaxis, np.newaxis, np.newaxis, ]) / (
            batch_pad.std(axis=(1, 2, 3))[:, np.newaxis, np.
newaxis, np.newaxis, ] + 1e-7)
        return batch_pad

    def __len__(self):
        return self.len_batch

class BatchWavDataset(Dataset):
    def __init__(self, batch_list, pad=0):
        self.X = batch_list
        self.pad = pad

    def __getitem__(self, index):
        batch_base = self.X[index]
        if batch_base.shape[-1] + self.pad * 2 < 20580:
            pad = int(np.ceil((20580 - batch_base.shape[-1]) / 2
))
        else:
            pad = self.pad
        if pad != 0:
            batch = np.zeros([batch_base.shape[0], 1, 1, batch_b
ase.shape[-1] + pad * 2], np.float32)
            batch[:, :, :, pad:-pad] = batch_base.astype(np.floa
t32) / 32768.0
        else:
            batch = batch_base.astype(np.float32) / 32768.0
        return batch

    def __len__(self):
        return len(self.X)

```

```

def predict_resnet(model, dataloader):
    sigmoid = torch.nn.Sigmoid().cuda()
    preds = []
    model.eval()
    for i, input in enumerate(dataloader):
        input = input.cuda(async=True)
        with torch.no_grad():
            pred = sigmoid(model(input)).data.cpu().numpy()
        preds.append(pred)
    preds = np.concatenate(preds)
    return preds

def predict_envnet(model, dataloader, activation='sigmoid'):
    sigmoid = torch.nn.Sigmoid().cuda()
    softmax = torch.nn.Softmax(dim=1).cuda()
    if activation == 'sigmoid':
        f_act = sigmoid
    elif activation == 'softmax':
        f_act = softmax
    preds = []
    model.eval()
    for i, input in enumerate(dataloader):
        input = input.cuda(async=True)
        with torch.no_grad():
            pred = f_act(model(input)).data.cpu().numpy()
        preds.append(pred)
    preds = np.concatenate(preds)
    return preds

if __name__ == '__main__':
    main()
import torch
import torch.nn as nn
from torch.utils.data.dataset import Dataset
from torch.utils.data import DataLoader
import torch.optim as optim
import torch.nn.functional as F

import pretrainedmodels

class ResNet(nn.Module):
    def __init__(self, num_classes=2):
        super(ResNet, self).__init__()

        self.num_classes = num_classes
        self.mode = 'train'

        self.base_model = pretrainedmodels.__dict__['resnet34'](
num_classes=num_classes, pretrained=None)

        self.conv1 = nn.Conv2d(1, 64, kernel_size=7, stride=2, padding=3,
                                bias=False)

```

```

        self.bn1 = self.base_model.bn1
        self.relu = self.base_model.relu
        self.maxpool = self.base_model.maxpool
        self.layer1 = self.base_model.layer1
        self.layer2 = self.base_model.layer2
        self.layer3 = self.base_model.layer3
        self.layer4 = self.base_model.layer4
        self.gmp = nn.AdaptiveMaxPool2d((1, 1))
        self.last_linear = nn.Linear(self.base_model.layer4[1].conv1.in_channels, num_classes)
    def forward(self, input):
        bs, ch, h, w = input.size()
        x0 = self.conv1(input)
        x0 = self.bn1(x0)
        x0 = self.relu(x0)
        x1 = self.maxpool(x0)
        x1 = self.layer1(x1)
        x2 = self.layer2(x1)
        x3 = self.layer3(x2)
        x4 = self.layer4(x3)
        x = self.gmp(x4).view(bs, -1)
        x = self.last_linear(x)
        return x

    def noisy(self, input):
        bs, ch, h, w = input.size()
        x0 = self.conv1(input)
        x0 = self.bn1(x0)
        x0 = self.relu(x0)
        x1 = self.maxpool(x0)
        x1 = self.layer1(x1)
        x2 = self.layer2(x1)
        x3 = self.layer3(x2)
        x4 = self.layer4(x3)

```

```

        x = self.gmp(x4).view(bs, -1)
        x = self.last_linear2(x)

    return x

class ConvBnRelu(nn.Module):
    def __init__(self, in_channel, out_channel, kernel_size, stride=1, padding=0, dilation=1, groups=1):
        super(ConvBnRelu, self).__init__()
        self.conv_bn_relu = nn.Sequential(
            nn.Conv2d(in_channel, out_channel, kernel_size, stride, padding, dilation, groups,
                      False),
            nn.BatchNorm2d(out_channel),
            nn.ReLU(True))

    def forward(self, x):
        return self.conv_bn_relu(x)

class Flatten(nn.Module):
    def forward(self, x):
        return x.view(x.size()[0], -1)

class EnvNetv2(nn.Module):
    def __init__(self, num_classes=1):
        super(EnvNetv2, self).__init__()
        self.conv1 = ConvBnRelu(1, 32, (1, 64), stride=(1, 2))
        self.conv2 = ConvBnRelu(32, 64, (1, 16), stride=(1, 2))
        self.conv3 = ConvBnRelu(1, 32, (8, 8))
        self.conv4 = ConvBnRelu(32, 32, (8, 8))
        self.conv5 = ConvBnRelu(32, 64, (1, 4))
        self.conv6 = ConvBnRelu(64, 64, (1, 4))
        self.conv7 = ConvBnRelu(64, 128, (1, 2))
        self.conv8 = ConvBnRelu(128, 128, (1, 2))
        self.conv9 = ConvBnRelu(128, 256, (1, 2))
        self.conv10 = ConvBnRelu(256, 256, (1, 2))
        self.maxpool1 = nn.MaxPool2d((1, 64), stride=(1, 64))
        self.maxpool2 = nn.MaxPool2d((5, 3), stride=(5, 3))
        self.maxpool3 = nn.MaxPool2d((1, 2), stride=(1, 2))
        self.gmp = nn.AdaptiveMaxPool2d((10, 1))
        self.flatten = Flatten()
        self.last_linear1 = nn.Sequential(
            nn.Linear(256 * 10, 1024),
            nn.ReLU(),
            nn.Dropout(p=0.2),
            nn.Linear(1024, 1024),
            nn.ReLU(),
            nn.Dropout(p=0.1),
            nn.Linear(1024, num_classes),
        )
        self.last_linear2 = nn.Sequential(

```



```

        nn.Linear(256 * 10, 1024),
        nn.ReLU(),
        nn.Dropout(p=0.2),
        nn.Linear(1024, 1024),
        nn.ReLU(),
        nn.Dropout(p=0.1),
        nn.Linear(1024, num_classes),
    )

    def forward(self, input):
        h = self.conv1(input)
        h = self.conv2(h)
        h = self.maxpool1(h)
        h = h.transpose(1, 2)
        h = self.conv3(h)
        h = self.conv4(h)
        h = self.maxpool2(h)
        h = self.conv5(h)
        h = self.conv6(h)
        h = self.maxpool3(h)
        h = self.conv7(h)
        h = self.conv8(h)
        h = self.maxpool3(h)
        h = self.conv9(h)
        h = self.conv10(h)
        h = self.gmp(h)
        h = self.flatten(h)
        h = self.last_linear1(h)
        return h

    def noisy(self, input):
        h = self.conv1(input)
        h = self.conv2(h)
        h = self.maxpool1(h)
        h = h.transpose(1, 2)
        h = self.conv3(h)
        h = self.conv4(h)
        h = self.maxpool2(h)
        h = self.conv5(h)
        h = self.conv6(h)
        h = self.maxpool3(h)
        h = self.conv7(h)
        h = self.conv8(h)
        h = self.maxpool3(h)
        h = self.conv9(h)
        h = self.conv10(h)
        h = self.gmp(h)
        h = self.flatten(h)
        h = self.last_linear2(h)
        return h

import numpy as np
import pandas as pd
import time
import librosa

# parameters
SAMPLE_RATE = 44100

```

```

N_MELS = 128
HOP_LENGTH = 347
N_FFT = 128*20
FMIN = 20
FMAX = SAMPLE_RATE//2

starttime = time.time()

def convert(df, input_dir, output_dir):
    for i in range(len(df)):
        if (i+1)%100==0: print("{} / {}, sec: {:.1f}".format(i+1,
len(df), time.time()-starttime))
        file_path = "{} / {}".format(input_dir, df['fname'][i])
        data, _ = librosa.core.load(file_path, sr=SAMPLE_RATE, r
es_type="kaiser_fast")
        data = librosa.feature.melspectrogram(
            data,
            sr=SAMPLE_RATE,
            n_mels=N_MELS,
            hop_length=HOP_LENGTH, # 1sec -> 128
            n_fft=N_FFT,
            fmin=FMIN,
            fmax=FMAX,
        ).astype(np.float32)
        np.save("{} / {}.npy".format(output_dir, df['fname'][i][:
4]), data)

def main():
    # load table data
    df_train = pd.read_csv("../input/train_curated.csv")
    df_noisy = pd.read_csv("../input/train_noisy.csv")
    df_test = pd.read_csv("../input/sample_submission.csv")

    # convert to logmel
    print("converting train data...")
    convert(df_train, "../input/train_curated/", "../input/mel12
8/train")
    print("converting noisy data...")
    convert(df_noisy, "../input/train_noisy/", "../input/mel128/
noisy")
    print("converting test data...")
    convert(df_test, "../input/test/", "../input/mel128/test")

if __name__ == '__main__':
    main()
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import time, random, sys
import sklearn.metrics
from sklearn.model_selection import KFold

import torch
import torch.nn as nn

```

```

import torch.backends.cudnn as cudnn
import torch.optim as optim
from torch.utils.data import DataLoader

sys.path.append('.')
from utils import AverageMeter, cycle, CosineLR, MelDataset, calculate_per_class_lwlrwrap
from models import ResNet

# set parameters
NUM_FOLD = 5
NUM_CLASS = 80
SEED = 42
NUM_EPOCH = 64*8
NUM_CYCLE = 64
BATCH_SIZE = 64
LR = [1e-3, 1e-6]
FOLD_LIST = [1, 2, 3, 4, 5]
CROP_LENGTH = 512
OUTPUT_DIR = "../models/resnet_model1"

cudnn.benchmark = True
starttime = time.time()

def main():
    # load table data
    df_train = pd.read_csv("../input/train_curated.csv")
    df_noisy = pd.read_csv("../input/train_noisy.csv")
    df_test = pd.read_csv("../input/sample_submission.csv")
    labels = df_test.columns[1:].tolist()
    for label in labels:
        df_train[label] = df_train['labels'].apply(lambda x: label in x)
        df_noisy[label] = df_noisy['labels'].apply(lambda x: label in x)

    df_train['path'] = "../input/mel128/train/" + df_train['filename']
    df_test['path'] = "../input/mel128/test/" + df_train['filename']
    df_noisy['path'] = "../input/mel128/noisy/" + df_noisy['filename']

    # fold splitting
    folds = list(KFold(n_splits=NUM_FOLD, shuffle=True, random_state=SEED).split(np.arange(len(df_train))))

    # Training
    log_columns = ['epoch', 'bce', 'lwlrwrap', 'bce_noisy', 'lwlrwrap_noisy', 'val_bce', 'val_lwlrwrap', 'time']
    for fold, (ids_train_split, ids_valid_split) in enumerate(folds):
        if fold+1 not in FOLD_LIST: continue
        print("fold: {}".format(fold + 1))
        train_log = pd.DataFrame(columns=log_columns)

```

```

    # build model
    model = ResNet(NUM_CLASS).cuda()

    # prepare data loaders
    df_train_fold = df_train.iloc[ids_train_split].reset_index(drop=True)
    dataset_train = MelDataset(df_train_fold['path'], df_train_fold[labels].values,
                               crop=CROP_LENGTH, crop_mode='random',
                               mixup=True, freqmask=True, gain=True,
                               )
    train_loader = DataLoader(dataset_train, batch_size=BATCH_SIZE,
                              shuffle=True, num_workers=1, pin_memory=True,
                              )

    df_valid = df_train.iloc[ids_valid_split].reset_index(drop=True)
    dataset_valid = MelDataset(df_valid['path'], df_valid[labels].values,)
    valid_loader = DataLoader(dataset_valid, batch_size=1,
                              shuffle=False, num_workers=1, pin_memory=True,
                              )

    dataset_noisy = MelDataset(df_noisy['path'], df_noisy[labels].values,
                               crop=CROP_LENGTH, crop_mode='random',
                               mixup=True, freqmask=True, gain=True,
                               )
    noisy_loader = DataLoader(dataset_noisy, batch_size=BATCH_SIZE,
                              shuffle=True, num_workers=1, pin_memory=True,
                              )
    noisy_itr = cycle(noisy_loader)

    # set optimizer and loss
    optimizer = optim.Adam(filter(lambda p: p.requires_grad, model.parameters()), lr=LR[0])
    scheduler = CosineLR(optimizer, step_size_min=LR[1], t0=len(train_loader) * NUM_CYCLE, tmult=1)

    # training
    for epoch in range(NUM_EPOCH):
        # train for one epoch
        bce, lwlap, bce_noisy, lwlap_noisy = train((train_loader, noisy_itr), model, optimizer, scheduler, epoch)

        # evaluate on validation set

```

```

val_bce, val_lwlrp = validate(valid_loader, model)

# print log
endtime = time.time() - starttime
print ("Epoch: {}/{} ".format(epoch + 1, NUM_EPOCH)
      + "CE: {:.4f} ".format(bce)
      + "LwLRAP: {:.4f} ".format(lwlrp)
      + "Noisy CE: {:.4f} ".format(bce_noisy)
      + "Noisy LWLRAP: {:.4f} ".format(lwlrp_noisy)
      + "Valid CE: {:.4f} ".format(val_bce)
      + "Valid LWLRAP: {:.4f} ".format(val_lwlrp)
      + "sec: {:.1f} ".format(endtime)
      )

# save log and weights
train_log_epoch = pd.DataFrame(
    [[epoch+1, bce, lwlrp, bce_noisy, lwlrp_noisy,
val_bce, val_lwlrp, endtime]],
    columns=log_columns)
train_log = pd.concat([train_log, train_log_epoch])
train_log.to_csv("{}train_log_fold{}.csv".format(OUTPUT_DIR, fold+1), index=False)
if (epoch+1)%NUM_CYCLE==0:
    torch.save(model.state_dict(), "{}weight_fold_{}_epoch{}.pth".format(OUTPUT_DIR, fold+1, epoch+1))

def train(train_loaders, model, optimizer, scheduler, epoch):
    train_loader, noisy_itr = train_loaders
    bce_avr = AverageMeter()
    bce_noisy_avr = AverageMeter()
    criterion_bce = nn.BCEWithLogitsLoss().cuda()
    sigmoid = nn.Sigmoid().cuda()

    # switch to train mode
    model.train()

    # training
    preds = np.zeros([0, NUM_CLASS], np.float32)
    y_true = np.zeros([0, NUM_CLASS], np.float32)
    preds_noisy = np.zeros([0, NUM_CLASS], np.float32)
    y_true_noisy = np.zeros([0, NUM_CLASS], np.float32)
    for i, (input, target) in enumerate(train_loader):
        # get batches
        input = torch.autograd.Variable(input.cuda())
        target = torch.autograd.Variable(target.cuda())

        input_noisy, target_noisy = next(noisy_itr)
        input_noisy = torch.autograd.Variable(input_noisy.cuda())

    )
    target_noisy = torch.autograd.Variable(target_noisy.cuda())

    ()

    # compute output
    output = model(input)
    bce = criterion_bce(output, target)

```

```

        output_noisy = model.noisy(input_noisy)
        bce_noisy = criterion_bce(torch.sigmoid(output_noisy), target_
noisy)
        loss = bce + bce_noisy
        pred = torch.sigmoid(output)
        pred = pred.data.cpu().numpy()
        pred_noisy = torch.sigmoid(output_noisy)
        pred_noisy = pred_noisy.data.cpu().numpy()

        # backprop
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
        scheduler.step()

        # record log
        bce_avr.update(bce.data, input.size(0))
        bce_noisy_avr.update(bce_noisy.data, input.size(0))
        preds = np.concatenate([preds, pred])
        y_true = np.concatenate([y_true, target.data.cpu().numpy
    ()])

        preds_noisy = np.concatenate([preds_noisy, pred_noisy])
        y_true_noisy = np.concatenate([y_true_noisy, target_nois
y.data.cpu().numpy()])

        # calc metric
        per_class_lwlraps, weight_per_class = calculate_per_class_lwlraps(y_true, preds)
        lwlraps = np.sum(per_class_lwlraps * weight_per_class)
        per_class_lwlraps, weight_per_class = calculate_per_class_lwlraps(y_true_noisy, preds_noisy)
        lwlraps_noisy = np.sum(per_class_lwlraps * weight_per_class)

        return bce_avr.avg.item(), lwlraps, bce_noisy_avr.avg.item(), lwlraps_noisy

def validate(val_loader, model):
    bce_avr = AverageMeter()
    sigmoid = torch.nn.Sigmoid().cuda()
    criterion_bce = nn.BCEWithLogitsLoss().cuda()

    # switch to eval mode
    model.eval()

    # validate
    preds = np.zeros([0, NUM_CLASS], np.float32)
    y_true = np.zeros([0, NUM_CLASS], np.float32)
    for i, (input, target) in enumerate(val_loader):
        # get batches
        input = torch.autograd.Variable(input.cuda())
        target = torch.autograd.Variable(target.cuda())

        # compute output
        with torch.no_grad():
            output = model(input)

```

```

        bce = criterion_bce(output, target)
        pred = sigmoid(output)
        pred = pred.data.cpu().numpy()

        # record log
        bce_avr.update(bce.data, input.size(0))
        preds = np.concatenate([preds, pred])
        y_true = np.concatenate([y_true, target.data.cpu().numpy()])

    # calc metric
    per_class_lwlap, weight_per_class = calculate_per_class_lwlap(y_true, preds)
    lwlap = np.sum(per_class_lwlap * weight_per_class)

    return bce_avr.avg.item(), lwlap

if __name__ == '__main__':
    main()
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import time, random, sys
import sklearn.metrics
from sklearn.model_selection import KFold

import torch
import torch.nn as nn
import torch.backends.cudnn as cudnn
import torch.optim as optim
from torch.utils.data import DataLoader

sys.path.append('.')
from utils import AverageMeter, cycle, CosineLR, MelDataset, calculate_per_class_lwlap
from models import ResNet

# set parameters
NUM_FOLD = 5
NUM_CLASS = 80
SEED = 42
NUM_EPOCH = 64*7
NUM_CYCLE = 64
BATCH_SIZE = 64
LR = [1e-3, 1e-6]
FOLD_LIST = [1, 2, 3, 4, 5]
CROP_LENGTH = 512
C_SEMI = 20
TEMPERATURE = 2
CROP_RATE = 0.25
LOAD_DIR = "../models/resnet_model1"
OUTPUT_DIR = "../models/resnet_model2"

cudnn.benchmark = True
starttime = time.time()

```

```

def main():
    # load table data
    df_train = pd.read_csv("../input/train_curated.csv")
    df_noisy = pd.read_csv("../input/train_noisy.csv")
    df_test = pd.read_csv("../input/sample_submission.csv")
    labels = df_test.columns[1:].tolist()
    for label in labels:
        df_train[label] = df_train['labels'].apply(lambda x: label in x)
        df_noisy[label] = df_noisy['labels'].apply(lambda x: label in x)

    df_train['path'] = "../input/mel128/train/" + df_train['filename']
    df_test['path'] = "../input/mel128/test/" + df_train['filename']
    df_noisy['path'] = "../input/mel128/noisy/" + df_noisy['filename']

    # calc sampling weight
    df_train['weight'] = 1
    df_noisy['weight'] = len(df_train) / len(df_noisy)

    # generate pseudo label with sharpening
    tmp = np.load("../input/pseudo_label/preds_noisy.npy").mean(axis=0,1)
    tmp = tmp ** TEMPERATURE
    tmp = tmp / tmp.sum(axis=1)[:, np.newaxis]
    df_noisy_pseudo = df_noisy.copy()
    df_noisy_pseudo[labels] = tmp

    # fold splitting
    folds = list(KFold(n_splits=NUM_FOLD, shuffle=True, random_state=SEED).split(np.arange(len(df_train))))
    folds_noisy = list(KFold(n_splits=NUM_FOLD, shuffle=True, random_state=SEED).split(np.arange(len(df_noisy))))

    # Training
    log_columns = ['epoch', 'bce', 'lwlrap', 'bce_noisy', 'lwlrap_noisy', 'semi_mse', 'val_bce', 'val_lwlrap', 'time']
    for fold, (ids_train_split, ids_valid_split) in enumerate(folds):
        if fold+1 not in FOLD_LIST: continue
        print("fold: {}".format(fold + 1))
        train_log = pd.DataFrame(columns=log_columns)

        # build model
        model = ResNet(NUM_CLASS).cuda()
        model.load_state_dict(torch.load("{}weight_fold_{}_epoch_512.pth".format(LOAD_DIR, fold+1)))

        # prepare data loaders
        df_train_fold = df_train.iloc[ids_train_split].reset_index(drop=True)
        dataset_train = MelDataset(df_train_fold['path'], df_train_fold['weight'], df_train_fold[labels])

```



```

in_fold[labels].values,
                                crop=CROP_LENGTH, crop_mode=
'additional', crop_rate=CROP_RATE,
                                mixup=True, freqmask=True, g
ain=True,
                                )
    train_loader = DataLoader(dataset_train, batch_size=BATC
H_SIZE,
                                shuffle=True, num_workers=1, p
in_memory=True,
                                )

    df_valid = df_train.iloc[ids_valid_split].reset_index(dr
op=True)
    dataset_valid = MelDataset(df_valid['path'], df_valid[la
bels].values,)
    valid_loader = DataLoader(dataset_valid, batch_size=1,
                                shuffle=False, num_workers=1,
pin_memory=True,
                                )

    dataset_noisy = MelDataset(df_noisy['path'], df_noisy[la
bels].values,
                                crop=CROP_LENGTH, crop_mode=
'additional', crop_rate=CROP_RATE,
                                mixup=True, freqmask=True, g
ain=True,
                                )
    noisy_loader = DataLoader(dataset_noisy, batch_size=BATC
H_SIZE,
                                shuffle=True, num_workers=1, p
in_memory=True,
                                )
    noisy_itr = cycle(noisy_loader)

    df_semi = pd.concat([df_train.iloc[ids_train_split], df_
noisy_pseudo.iloc[folds_noisy[fold][0]]]).reset_index(drop=True)
    semi_sampler = torch.utils.data.sampler.WeightedRandomSa
mpler(df_semi['weight'].values, len(df_semi))
    dataset_semi = MelDataset(df_semi['path'], df_semi[label
s].values,
                                crop=CROP_LENGTH, crop_mode='
additional', crop_rate=CROP_RATE,
                                mixup=True, freqmask=True, ga
in=True,
                                )
    semi_loader = DataLoader(dataset_semi,
                                batch_size=BATCH_SIZE,
                                shuffle=False, num_workers=1, p
in_memory=True,
                                sampler=semi_sampler,
                                )
    semi_itr = cycle(semi_loader)

    # set optimizer and loss
    optimizer = optim.Adam(filter(lambda p: p.requires_grad,

```

```

model.parameters()), lr=LR[0])
    scheduler = CosineLR(optimizer, step_size_min=LR[1], t0=
len(train_loader) * NUM_CYCLE, tmult=1)

    # training
    for epoch in range(NUM_EPOCH):
        # train for one epoch
        bce, lwlap, bce_noisy, lwlap_noisy, mse_semi = tra
in((train_loader, noisy_itr, semi_itr), model, optimizer, schedu
ler, epoch)

        # evaluate on validation set
        val_bce, val_lwlap = validate(valid_loader, model)

        # print log
        endtime = time.time() - starttime
        print("Epoch: {}/{} ".format(epoch + 1, NUM_EPOCH)
              + "CE: {:.4f} ".format(bce)
              + "LwLRAP: {:.4f} ".format(lwlap)
              + "Noisy CE: {:.4f} ".format(bce_noisy)
              + "Noisy LWLRAP: {:.4f} ".format(lwlap_noisy)
              + "Semi MSE: {:.4f} ".format(mse_semi)
              + "Valid CE: {:.4f} ".format(val_bce)
              + "Valid LWLRAP: {:.4f} ".format(val_lwlap)
              + "sec: {:.1f} ".format(endtime)
              )

        # save log and weights
        train_log_epoch = pd.DataFrame(
            [[epoch+1, bce, lwlap, bce_noisy, lwlap_noisy,
mse_semi, val_bce, val_lwlap, endtime]],
            columns=log_columns)
        train_log = pd.concat([train_log, train_log_epoch])
        train_log.to_csv("{} /train_log_fold{}.csv".format(OU
TPUT_DIR, fold+1), index=False)
        if (epoch+1)%NUM_CYCLE==0:
            torch.save(model.state_dict(), "{} /weight_fold_{
}_epoch{}.pth".format(OUTPUT_DIR, fold+1, epoch+1))

def train(train_loaders, model, optimizer, scheduler, epoch):
    train_loader, noisy_itr, semi_itr = train_loaders
    bce_avr = AverageMeter()
    bce_noisy_avr = AverageMeter()
    mse_semi_avr = AverageMeter()
    criterion_bce = nn.BCEWithLogitsLoss().cuda()
    criterion_mse = nn.MSELoss().cuda()
    sigmoid = nn.Sigmoid().cuda()

    # switch to train mode
    model.train()

    # training
    preds = np.zeros([0, NUM_CLASS], np.float32)
    y_true = np.zeros([0, NUM_CLASS], np.float32)
    preds_noisy = np.zeros([0, NUM_CLASS], np.float32)

```

```

y_true_noisy = np.zeros([0, NUM_CLASS], np.float32)
for i, (input, target) in enumerate(train_loader):
    # get batches
    input = torch.autograd.Variable(input.cuda())
    target = torch.autograd.Variable(target.cuda())

    input_noisy, target_noisy = next(noisy_itr)
    input_noisy = torch.autograd.Variable(input_noisy.cuda())
)
    target_noisy = torch.autograd.Variable(target_noisy.cuda
())

    input_semi, target_semi = next(semi_itr)
    input_semi = torch.autograd.Variable(input_semi.cuda())
    target_semi = torch.autograd.Variable(target_semi.cuda())
)

    # compute output
    output = model(input)
    bce = criterion_bce(output, target)
    output_noisy = model.noisy(input_noisy)
    bce_noisy = criterion_bce(sigmoid(output_noisy), target_
noisy)

    output_semi = model(input_semi)
    mse_semi = criterion_mse(sigmoid(output_semi), target_se
mi)

    loss = bce + bce_noisy + C_SEMI * mse_semi
    pred = sigmoid(output)
    pred = pred.data.cpu().numpy()
    pred_noisy = sigmoid(output_noisy)
    pred_noisy = pred_noisy.data.cpu().numpy()

    # backprop
    optimizer.zero_grad()
    loss.backward()
    optimizer.step()
    scheduler.step()

    # record log
    bce_avr.update(bce.data, input.size(0))
    bce_noisy_avr.update(bce_noisy.data, input.size(0))
    mse_semi_avr.update(mse_semi.data, input.size(0))
    preds = np.concatenate([preds, pred])
    y_true = np.concatenate([y_true, target.data.cpu().numpy
])
    preds_noisy = np.concatenate([preds_noisy, pred_noisy])
    y_true_noisy = np.concatenate([y_true_noisy, target_nois
y.data.cpu().numpy()])

    # calc metric
    per_class_lwlraps, weight_per_class = calculate_per_class_lwlraps(y_true, preds)
    lwlraps = np.sum(per_class_lwlraps * weight_per_class)
    per_class_lwlraps, weight_per_class = calculate_per_class_lwlraps(y_true_noisy, preds_noisy)

```

```

    lwlapr_noisy = np.sum(per_class_lwlapr * weight_per_class)

    return bce_avr.avg.item(), lwlapr, bce_noisy_avr.avg.item(),
    lwlapr_noisy, mse_semi_avr.avg.item()

def validate(val_loader, model):
    bce_avr = AverageMeter()
    sigmoid = torch.nn.Sigmoid().cuda()
    criterion_bce = nn.BCEWithLogitsLoss().cuda()

    # switch to eval mode
    model.eval()

    # validate
    preds = np.zeros([0, NUM_CLASS], np.float32)
    y_true = np.zeros([0, NUM_CLASS], np.float32)
    for i, (input, target) in enumerate(val_loader):
        # get batches
        input = torch.autograd.Variable(input.cuda())
        target = torch.autograd.Variable(target.cuda())

        # compute output
        with torch.no_grad():
            output = model(input)
            bce = criterion_bce(output, target)
            pred = sigmoid(output)
            pred = pred.data.cpu().numpy()

        # record log
        bce_avr.update(bce.data, input.size(0))
        preds = np.concatenate([preds, pred])
        y_true = np.concatenate([y_true, target.data.cpu().numpy()])

    # calc metric
    per_class_lwlapr, weight_per_class = calculate_per_class_lwlapr(y_true, preds)
    lwlapr = np.sum(per_class_lwlapr * weight_per_class)

    return bce_avr.avg.item(), lwlapr

if __name__ == '__main__':
    main()
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import time, random, sys
import sklearn.metrics
from sklearn.model_selection import KFold

import torch
import torch.nn as nn
import torch.backends.cudnn as cudnn
import torch.optim as optim
from torch.utils.data import DataLoader

```

```

sys.path.append('.')
from utils import AverageMeter, cycle, CosineLR, MelDataset, calculate_per_class_lwlrp
from models import ResNet

# set parameters
NUM_FOLD = 5
NUM_CLASS = 80
SEED = 42
NUM_EPOCH = 64*8
NUM_CYCLE = 64
BATCH_SIZE = 64
LR = [1e-3, 1e-6]
FOLD_LIST = [1, 2, 3, 4, 5]
CROP_LENGTH = 1024
OUTPUT_DIR = "../models/resnet_model3"

cudnn.benchmark = True
starttime = time.time()

def main():
    # load table data
    df_train = pd.read_csv("../input/train_curated.csv")
    df_noisy = pd.read_csv("../input/train_noisy.csv")
    df_test = pd.read_csv("../input/sample_submission.csv")
    labels = df_test.columns[1:].tolist()
    for label in labels:
        df_train[label] = df_train['labels'].apply(lambda x: label in x)
        df_noisy[label] = df_noisy['labels'].apply(lambda x: label in x)

    df_train['path'] = "../input/mel128/train/" + df_train['filename']
    df_test['path'] = "../input/mel128/test/" + df_train['filename']
    df_noisy['path'] = "../input/mel128/noisy/" + df_noisy['filename']

    # fold splitting
    folds = list(KFold(n_splits=NUM_FOLD, shuffle=True, random_state=SEED).split(np.arange(len(df_train))))

    # Training
    log_columns = ['epoch', 'bce', 'lwlrp', 'bce_noisy', 'lwlrp_noisy', 'val_bce', 'val_lwlrp', 'time']
    for fold, (ids_train_split, ids_valid_split) in enumerate(folds):
        if fold+1 not in FOLD_LIST: continue
        print("fold: {}".format(fold + 1))
        train_log = pd.DataFrame(columns=log_columns)

        # build model
        model = ResNet(NUM_CLASS).cuda()

```

```

    # prepare data loaders
    df_train_fold = df_train.iloc[ids_train_split].reset_index(drop=True)
    dataset_train = MelDataset(df_train_fold['path'], df_train_fold[labels].values,
                               crop=CROP_LENGTH, crop_mode='random',
                               mixup=True, freqmask=True, gain=True,
                               )
    train_loader = DataLoader(dataset_train, batch_size=BATCH_SIZE,
                              shuffle=True, num_workers=1, pin_memory=True,
                              )

    df_valid = df_train.iloc[ids_valid_split].reset_index(drop=True)
    dataset_valid = MelDataset(df_valid['path'], df_valid[labels].values,)
    valid_loader = DataLoader(dataset_valid, batch_size=1,
                              shuffle=False, num_workers=1,
                              pin_memory=True,
                              )

    dataset_noisy = MelDataset(df_noisy['path'], df_noisy[labels].values,
                               crop=CROP_LENGTH, crop_mode='random',
                               mixup=True, freqmask=True, gain=True,
                               )
    noisy_loader = DataLoader(dataset_noisy, batch_size=BATCH_SIZE,
                              shuffle=True, num_workers=1, pin_memory=True,
                              )
    noisy_itr = cycle(noisy_loader)

    # set optimizer and loss
    optimizer = optim.Adam(filter(lambda p: p.requires_grad,
                                   model.parameters()), lr=LR[0])
    scheduler = CosineLR(optimizer, step_size_min=LR[1], t0=len(train_loader) * NUM_CYCLE, tmult=1)

    # training
    for epoch in range(NUM_EPOCH):
        # train for one epoch
        bce, lwlap, bce_noisy, lwlap_noisy = train((train_loader, noisy_itr), model, optimizer, scheduler, epoch)

        # evaluate on validation set
        val_bce, val_lwlap = validate(valid_loader, model)

        # print log

```

```

        endtime = time.time() - starttime
        print("Epoch: {}/{}".format(epoch + 1, NUM_EPOCH)
              + "CE: {:.4f}".format(bce)
              + "LwLRAP: {:.4f}".format(lwlrwrap)
              + "Noisy CE: {:.4f}".format(bce_noisy)
              + "Noisy LwLRAP: {:.4f}".format(lwlrwrap_noisy)
              + "Valid CE: {:.4f}".format(val_bce)
              + "Valid LwLRAP: {:.4f}".format(val_lwlrwrap)
              + "sec: {:.1f}".format(endtime)
              )

        # save log and weights
        train_log_epoch = pd.DataFrame(
            [[epoch+1, bce, lwlrwrap, bce_noisy, lwlrwrap_noisy,
              val_bce, val_lwlrwrap, endtime]],
            columns=log_columns)
        train_log = pd.concat([train_log, train_log_epoch])
        train_log.to_csv("{}train_log_fold{}.csv".format(OUTPUT_DIR,
            fold+1), index=False)
        if (epoch+1)%NUM_CYCLE==0:
            torch.save(model.state_dict(), "{}weight_fold_{}_epoch{}.pth".format(OUTPUT_DIR, fold+1, epoch+1))

def train(train_loaders, model, optimizer, scheduler, epoch):
    train_loader, noisy_itr = train_loaders
    bce_avr = AverageMeter()
    bce_noisy_avr = AverageMeter()
    criterion_bce = nn.BCEWithLogitsLoss().cuda()
    sigmoid = nn.Sigmoid().cuda()

    # switch to train mode
    model.train()

    # training
    preds = np.zeros([0, NUM_CLASS], np.float32)
    y_true = np.zeros([0, NUM_CLASS], np.float32)
    preds_noisy = np.zeros([0, NUM_CLASS], np.float32)
    y_true_noisy = np.zeros([0, NUM_CLASS], np.float32)
    for i, (input, target) in enumerate(train_loader):
        # get batches
        input = torch.autograd.Variable(input.cuda())
        target = torch.autograd.Variable(target.cuda())

        input_noisy, target_noisy = next(noisy_itr)
        input_noisy = torch.autograd.Variable(input_noisy.cuda())
        target_noisy = torch.autograd.Variable(target_noisy.cuda())

        # compute output
        output = model(input)
        bce = criterion_bce(output, target)
        output_noisy = model(input_noisy)
        bce_noisy = criterion_bce(sigmoid(output_noisy), target_noisy)

```

```

        loss = bce + bce_noisy
        pred = sigmoid(output)
        pred = pred.data.cpu().numpy()
        pred_noisy = sigmoid(output_noisy)
        pred_noisy = pred_noisy.data.cpu().numpy()

        # backprop
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
        scheduler.step()

        # record log
        bce_avr.update(bce.data, input.size(0))
        bce_noisy_avr.update(bce_noisy.data, input.size(0))
        preds = np.concatenate([preds, pred])
        y_true = np.concatenate([y_true, target.data.cpu().numpy()
    (]])
        preds_noisy = np.concatenate([preds_noisy, pred_noisy])
        y_true_noisy = np.concatenate([y_true_noisy, target_nois
y.data.cpu().numpy()])

        # calc metric
        per_class_lwlraps, weight_per_class = calculate_per_class_lwlraps(y_true, preds)
        lwlraps = np.sum(per_class_lwlraps * weight_per_class)
        per_class_lwlraps, weight_per_class = calculate_per_class_lwlraps(y_true_noisy, preds_noisy)
        lwlraps_noisy = np.sum(per_class_lwlraps * weight_per_class)

        return bce_avr.avg.item(), lwlraps, bce_noisy_avr.avg.item(), lwlraps_noisy

def validate(val_loader, model):
    bce_avr = AverageMeter()
    sigmoid = torch.nn.Sigmoid().cuda()
    criterion_bce = nn.BCEWithLogitsLoss().cuda()

    # switch to eval mode
    model.eval()

    # validate
    preds = np.zeros([0, NUM_CLASS], np.float32)
    y_true = np.zeros([0, NUM_CLASS], np.float32)
    for i, (input, target) in enumerate(val_loader):
        # get batches
        input = torch.autograd.Variable(input.cuda())
        target = torch.autograd.Variable(target.cuda())

        # compute output
        with torch.no_grad():
            output = model(input)
            bce = criterion_bce(output, target)
            pred = sigmoid(output)
            pred = pred.data.cpu().numpy()

```



```

        # record log
        bce_avr.update(bce.data, input.size(0))
        preds = np.concatenate([preds, pred])
        y_true = np.concatenate([y_true, target.data.cpu().numpy
    ()])

    # calc metric
    per_class_lwlrwrap, weight_per_class = calculate_per_class_lwlrwrap(y_true, preds)
    lwlrwrap = np.sum(per_class_lwlrwrap * weight_per_class)

    return bce_avr.avg.item(), lwlrwrap

if __name__ == '__main__':
    main()
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import time, random, sys
import sklearn.metrics
from sklearn.model_selection import KFold

import torch
import torch.nn as nn
import torch.backends.cudnn as cudnn
import torch.optim as optim
from torch.utils.data import DataLoader

sys.path.append('.')
from utils import AverageMeter, cycle, CosineLR, WaveDataset, calculate_per_class_lwlrwrap
from models import EnvNetv2

# set parameters
NUM_FOLD = 5
NUM_CLASS = 80
SEED = 42
NUM_EPOCH = 400 * 1
NUM_CYCLE = 400
BATCH_SIZE = 16
LR = [1e-1, 1e-6]
FOLD_LIST = [1, 2, 3, 4, 5]
FOLD_LIST = [1, 2,]
CROP_LENGTH = 133300
OUTPUT_DIR = "../models/envnet_model4_0"

cudnn.benchmark = True
starttime = time.time()

def main():
    # load table data
    df_train = pd.read_csv("../input/train_curated.csv")
    df_noisy = pd.read_csv("../input/train_noisy.csv")
    df_test = pd.read_csv("../input/sample_submission.csv")

```

```

    labels = df_test.columns[1:].tolist()
    for label in labels:
        df_train[label] = df_train['labels'].apply(lambda x: label in x)
        df_noisy[label] = df_noisy['labels'].apply(lambda x: label in x)

    df_train['path'] = "../input/train_curated/" + df_train['fname']
    df_test['path'] = "../input/test/" + df_train['fname']
    df_noisy['path'] = "../input/train_noisy/" + df_noisy['fname']

    # fold splitting
    folds = list(KFold(n_splits=NUM_FOLD, shuffle=True, random_state=SEED).split(np.arange(len(df_train))))

    # Training
    log_columns = ['epoch', 'kl', 'lwlrwrap', 'kl_noisy', 'lwlrwrap_noisy', 'val_kl', 'val_lwlrwrap', 'time']
    for fold, (ids_train_split, ids_valid_split) in enumerate(folds):
        if fold+1 not in FOLD_LIST: continue
        print("fold: {}".format(fold + 1))
        train_log = pd.DataFrame(columns=log_columns)

        # build model
        model = EnvNetv2(NUM_CLASS).cuda()

        # prepare data loaders
        df_train_fold = df_train.iloc[ids_train_split].reset_index(drop=True)
        dataset_train = WaveDataset(df_train_fold['path'], df_train_fold[labels].values,
                                     crop=CROP_LENGTH, crop_mode='random', padding=CROP_LENGTH//2,
                                     mixup=True, scaling=1.25, gaussian=6,
                                     )
        train_loader = DataLoader(dataset_train, batch_size=BATCH_SIZE,
                                   shuffle=True, num_workers=1, pin_memory=True,
                                   )

        df_valid = df_train.iloc[ids_valid_split].reset_index(drop=True)
        dataset_valid = WaveDataset(df_valid['path'], df_valid[labels].values, padding=CROP_LENGTH//2)
        valid_loader = DataLoader(dataset_valid, batch_size=1,
                                   shuffle=False, num_workers=1,
                                   pin_memory=True,
                                   )

        dataset_noisy = WaveDataset(df_noisy['path'], df_noisy[labels].values,

```

```

crop=CROP_LENGTH, crop_mode=
'random', padding=CROP_LENGTH//2,
mixup=True, scaling=1.25, ga
in=6,
)
noisy_loader = DataLoader(dataset_noisy, batch_size=BATC
H_SIZE,
shuffle=True, num_workers=1, p
in_memory=True,
)
noisy_itr = cycle(noisy_loader)

# set optimizer and loss
optimizer = optim.SGD(filter(lambda p: p.requires_grad,
model.parameters()), lr=LR[0], momentum = 0.9, nesterov = True)
scheduler = CosineLR(optimizer, step_size_min=LR[1], t0=
len(train_loader) * NUM_CYCLE, tmult=1)

# training
for epoch in range(NUM_EPOCH):
    # train for one epoch
    kl, lwlap, kl_noisy, lwlap_noisy = train((train_lo
ader, noisy_itr), model, optimizer, scheduler, epoch)

    # evaluate on validation set
    val_kl, val_lwlap = validate(valid_loader, model)

    # print log
    endtime = time.time() - starttime
    print("Epoch: {}/{}".format(epoch + 1, NUM_EPOCH)
        + "KL: {:.4f} ".format(kl)
        + "LwLRAP: {:.4f} ".format(lwlap)
        + "Noisy KL: {:.4f} ".format(kl_noisy)
        + "Noisy LWLRAP: {:.4f} ".format(lwlap_noisy)
        + "Valid KL: {:.4f} ".format(val_kl)
        + "Valid LWLRAP: {:.4f} ".format(val_lwlap)
        + "sec: {:.1f} ".format(endtime)
        )

    # save log and weights
    train_log_epoch = pd.DataFrame(
        [[epoch+1, kl, lwlap, kl_noisy, lwlap_noisy, v
al_kl, val_lwlap, endtime]],
        columns=log_columns)
    train_log = pd.concat([train_log, train_log_epoch])
    train_log.to_csv("{}train_log_fold{}.csv".format(OU
TPUT_DIR, fold+1), index=False)
    if (epoch+1)%NUM_CYCLE==0:
        torch.save(model.state_dict(), "{}weight_fold_
{}_epoch{}.pth".format(OUTPUT_DIR, fold+1, epoch+1))

def train(train_loaders, model, optimizer, scheduler, epoch):
    train_loader, noisy_itr = train_loaders
    kl_avr = AverageMeter()
    kl_noisy_avr = AverageMeter()

```

```

lsigmoid = nn.LogSigmoid().cuda()
lsoftmax = nn.LogSoftmax(dim=1).cuda()
softmax = nn.Softmax(dim=1).cuda()
criterion_kl = nn.KLDivLoss().cuda()

# switch to train mode
model.train()

# training
preds = np.zeros([0, NUM_CLASS], np.float32)
y_true = np.zeros([0, NUM_CLASS], np.float32)
preds_noisy = np.zeros([0, NUM_CLASS], np.float32)
y_true_noisy = np.zeros([0, NUM_CLASS], np.float32)
for i, (input, target) in enumerate(train_loader):
    # get batches
    input = torch.autograd.Variable(input.cuda())
    target = torch.autograd.Variable(target.cuda())

    input_noisy, target_noisy = next(noisy_itr)
    input_noisy = torch.autograd.Variable(input_noisy.cuda())
    target_noisy = torch.autograd.Variable(target_noisy.cuda())

    # compute output
    output = model(input)
    kl = criterion_kl(lsoftmax(output), target)
    output_noisy = model(input_noisy)
    kl_noisy = criterion_kl(lsoftmax(output_noisy), target_noisy)

    loss = kl + kl_noisy
    pred = softmax(output)
    pred = pred.data.cpu().numpy()
    pred_noisy = softmax(output_noisy)
    pred_noisy = pred_noisy.data.cpu().numpy()

    # backprop
    optimizer.zero_grad()
    loss.backward()
    optimizer.step()
    scheduler.step()

    # record log
    kl_avr.update(kl.data, input.size(0))
    kl_noisy_avr.update(kl_noisy.data, input.size(0))
    preds = np.concatenate([preds, pred])
    y_true = np.concatenate([y_true, target.data.cpu().numpy()])

    preds_noisy = np.concatenate([preds_noisy, pred_noisy])
    y_true_noisy = np.concatenate([y_true_noisy, target_noisy.data.cpu().numpy()])

# calc metric
per_class_lwlap, weight_per_class = calculate_per_class_lwlap(y_true, preds)
lwlap = np.sum(per_class_lwlap * weight_per_class)

```

```

    per_class_lwlap, weight_per_class = calculate_per_class_lwlap(y_true_noisy, preds_noisy)
    lwlap_noisy = np.sum(per_class_lwlap * weight_per_class)

    return kl_avr.avg.item(), lwlap, kl_noisy_avr.avg.item(), lwlap_noisy

```

```

def validate(val_loader, model):
    kl_avr = AverageMeter()
    lsoftmax = nn.LogSoftmax(dim=1).cuda()
    softmax = torch.nn.Softmax(dim=1).cuda()
    criterion_kl = nn.KLDivLoss().cuda()

    # switch to eval mode
    model.eval()

    # validate
    preds = np.zeros([0, NUM_CLASS], np.float32)
    y_true = np.zeros([0, NUM_CLASS], np.float32)
    for i, (input, target) in enumerate(val_loader):
        # get batches
        input = torch.autograd.Variable(input.cuda())
        target = torch.autograd.Variable(target.cuda())

        # compute output
        with torch.no_grad():
            output = model(input)
            kl = criterion_kl(lsoftmax(output), target)
            pred = softmax(output)
            pred = pred.data.cpu().numpy()

        # record log
        kl_avr.update(kl.data, input.size(0))
        preds = np.concatenate([preds, pred])
        y_true = np.concatenate([y_true, target.data.cpu().numpy()])

    # calc metric
    per_class_lwlap, weight_per_class = calculate_per_class_lwlap(y_true, preds)
    lwlap = np.sum(per_class_lwlap * weight_per_class)

    return kl_avr.avg.item(), lwlap

if __name__ == '__main__':
    main()
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import time, random, sys
import sklearn.metrics
from sklearn.model_selection import KFold

import torch
import torch.nn as nn

```

```

import torch.backends.cudnn as cudnn
import torch.optim as optim
from torch.utils.data import DataLoader

sys.path.append('.')
from utils import AverageMeter, cycle, CosineLR, WaveDataset, calculate_per_class_lwlrwrap
from models import EnvNetv2

# set parameters
NUM_FOLD = 5
NUM_CLASS = 80
SEED = 42
NUM_EPOCH = 4 #80*3
NUM_CYCLE = 2 #80
BATCH_SIZE = 64
LR = [1e-1, 1e-6]
FOLD_LIST = [1, 2, 3, 4, 5]
FOLD_LIST = [1, 2,]
CROP_LENGTH = 133300
LOAD_DIR = "../models/envnet_model4_0"
OUTPUT_DIR = "../models/envnet_model4"

cudnn.benchmark = True
starttime = time.time()

def main():
    # load table data
    df_train = pd.read_csv("../input/train_curated.csv")
    df_noisy = pd.read_csv("../input/train_noisy.csv")
    df_test = pd.read_csv("../input/sample_submission.csv")
    labels = df_test.columns[1:].tolist()
    for label in labels:
        df_train[label] = df_train['labels'].apply(lambda x: label in x)
        df_noisy[label] = df_noisy['labels'].apply(lambda x: label in x)

    df_train['path'] = "../input/train_curated/" + df_train['fname']
    df_test['path'] = "../input/test/" + df_train['fname']
    df_noisy['path'] = "../input/train_noisy/" + df_noisy['fname']

    # fold splitting
    folds = list(KFold(n_splits=NUM_FOLD, shuffle=True, random_state=SEED).split(np.arange(len(df_train))))

    # Training
    log_columns = ['epoch', 'bce', 'lwlrwrap', 'bce_noisy', 'lwlrwrap_noisy', 'val_bce', 'val_lwlrwrap', 'time']
    for fold, (ids_train_split, ids_valid_split) in enumerate(folds):
        if fold+1 not in FOLD_LIST: continue
        print("fold: {}".format(fold + 1))

```

```

train_log = pd.DataFrame(columns=log_columns)

# build model
model = EnvNetv2(NUM_CLASS).cuda()
# model.load_state_dict(torch.load("{}weight_fold_{}_epoch_400.pth".format(LoadDir, fold+1)))

# prepare data loaders
df_train_fold = df_train.iloc[ids_train_split].reset_index(drop=True)
dataset_train = WaveDataset(df_train_fold['path'], df_train_fold[labels].values,
                             crop=CROP_LENGTH, crop_mode='random', padding=CROP_LENGTH//2,
                             mixup=True, scaling=1.25, gaussian=6,
                             )
train_loader = DataLoader(dataset_train, batch_size=BATCH_SIZE,
                           shuffle=True, num_workers=1, pin_memory=True,
                           )

df_valid = df_train.iloc[ids_valid_split].reset_index(drop=True)
dataset_valid = WaveDataset(df_valid['path'], df_valid[labels].values, padding=CROP_LENGTH//2)
valid_loader = DataLoader(dataset_valid, batch_size=1,
                           shuffle=False, num_workers=1, pin_memory=True,
                           )

dataset_noisy = WaveDataset(df_noisy['path'], df_noisy[labels].values,
                             crop=CROP_LENGTH, crop_mode='random', padding=CROP_LENGTH//2,
                             mixup=True, scaling=1.25, gaussian=6,
                             )
noisy_loader = DataLoader(dataset_noisy, batch_size=BATCH_SIZE,
                           shuffle=True, num_workers=1, pin_memory=True,
                           )
noisy_itr = cycle(noisy_loader)

# set optimizer and loss
optimizer = optim.SGD(filter(lambda p: p.requires_grad, model.parameters()), lr=LR[0], momentum = 0.9, nesterov = True)
scheduler = CosineLR(optimizer, step_size_min=LR[1], t0=len(train_loader) * NUM_CYCLE, tmult=1)

# training
for epoch in range(NUM_EPOCH):
    # train for one epoch
    bce, lwlap, bce_noisy, lwlap_noisy = train((train_

```

```

loader, noisy_itr), model, optimizer, scheduler, epoch)

    # evaluate on validation set
    val_bce, val_lwlrp = validate(valid_loader, model)

    # print log
    endtime = time.time() - starttime
    print("Epoch: {}/{}".format(epoch + 1, NUM_EPOCH)
          + "CE: {:.4f}".format(bce)
          + "LwLRAP: {:.4f}".format(lwlrp)
          + "Noisy CE: {:.4f}".format(bce_noisy)
          + "Noisy LWLRAP: {:.4f}".format(lwlrp_noisy)
          + "Valid CE: {:.4f}".format(val_bce)
          + "Valid LWLRAP: {:.4f}".format(val_lwlrp)
          + "sec: {:.1f}".format(endtime)
        )

    # save log and weights
    train_log_epoch = pd.DataFrame(
        [[epoch+1, bce, lwlrp, bce_noisy, lwlrp_noisy,
         val_bce, val_lwlrp, endtime]],
        columns=log_columns)
    train_log = pd.concat([train_log, train_log_epoch])
    train_log.to_csv("{}train_log_fold{}.csv".format(OUTPUT_DIR,
        fold+1), index=False)
    if (epoch+1)%NUM_CYCLE==0:
        torch.save(model.state_dict(), "{}weight_fold_{}_epoch{}.pth".format(OUTPUT_DIR,
            fold+1, epoch+1))

def train(train_loaders, model, optimizer, scheduler, epoch):
    train_loader, noisy_itr = train_loaders
    bce_avr = AverageMeter()
    bce_noisy_avr = AverageMeter()
    criterion_bce = nn.BCEWithLogitsLoss().cuda()
    sigmoid = nn.Sigmoid().cuda()

    # switch to train mode
    model.train()

    # training
    preds = np.zeros([0, NUM_CLASS], np.float32)
    y_true = np.zeros([0, NUM_CLASS], np.float32)
    preds_noisy = np.zeros([0, NUM_CLASS], np.float32)
    y_true_noisy = np.zeros([0, NUM_CLASS], np.float32)
    for i, (input, target) in enumerate(train_loader):
        # get batches
        input = torch.autograd.Variable(input.cuda())
        target = torch.autograd.Variable(target.cuda())

        input_noisy, target_noisy = next(noisy_itr)
        input_noisy = torch.autograd.Variable(input_noisy.cuda())

        target_noisy = torch.autograd.Variable(target_noisy.cuda())

```



```

        # compute output
        output = model(input)
        bce = criterion_bce(sigmoid(output), target)
        output_noisy = model.noisy(input_noisy)
        bce_noisy = criterion_bce(sigmoid(output_noisy), target_
noisy)

        loss = bce + bce_noisy
        pred = sigmoid(output)
        pred = pred.data.cpu().numpy()
        pred_noisy = sigmoid(output_noisy)
        pred_noisy = pred_noisy.data.cpu().numpy()

        # backprop
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
        scheduler.step()

        # record log
        bce_avr.update(bce.data, input.size(0))
        bce_noisy_avr.update(bce_noisy.data, input.size(0))
        preds = np.concatenate([preds, pred])
        y_true = np.concatenate([y_true, target.data.cpu().numpy
    ()])

        preds_noisy = np.concatenate([preds_noisy, pred_noisy])
        y_true_noisy = np.concatenate([y_true_noisy, target_nois
y.data.cpu().numpy()])

        # calc metric
        per_class_lwlraps, weight_per_class = calculate_per_class_lwlraps(y_true, preds)
        lwlraps = np.sum(per_class_lwlraps * weight_per_class)
        per_class_lwlraps, weight_per_class = calculate_per_class_lwlraps(y_true_noisy, preds_noisy)
        lwlraps_noisy = np.sum(per_class_lwlraps * weight_per_class)

        return bce_avr.avg.item(), lwlraps, bce_noisy_avr.avg.item(), lwlraps_noisy

def validate(val_loader, model):
    bce_avr = AverageMeter()
    sigmoid = nn.Sigmoid().cuda()
    criterion_bce = nn.BCEWithLogitsLoss().cuda()

    # switch to eval mode
    model.eval()

    # validate
    preds = np.zeros([0, NUM_CLASS], np.float32)
    y_true = np.zeros([0, NUM_CLASS], np.float32)
    for i, (input, target) in enumerate(val_loader):
        # get batches
        input = torch.autograd.Variable(input.cuda())
        target = torch.autograd.Variable(target.cuda())

```

```

        # compute output
        with torch.no_grad():
            output = model(input)
            bce = criterion_bce(sigmoid(output), target)
            pred = sigmoid(output)
            pred = pred.data.cpu().numpy()

        # record log
        bce_avr.update(bce.data, input.size(0))
        preds = np.concatenate([preds, pred])
        y_true = np.concatenate([y_true, target.data.cpu().numpy()])

    ])

    # calc metric
    per_class_lwlap, weight_per_class = calculate_per_class_lwlap(y_true, preds)
    lwlap = np.sum(per_class_lwlap * weight_per_class)

    return bce_avr.avg.item(), lwlap

if __name__ == '__main__':
    main()
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import time, random, sys
import sklearn.metrics
from sklearn.model_selection import KFold

import torch
import torch.nn as nn
import torch.backends.cudnn as cudnn
import torch.optim as optim
from torch.utils.data import DataLoader

sys.path.append('.')
from utils import AverageMeter, cycle, CosineLR, WaveDataset, calculate_per_class_lwlap
from models import EnvNetv2

# set parameters
NUM_FOLD = 5
NUM_CLASS = 80
SEED = 42
NUM_EPOCH = 80 * 5
NUM_CYCLE = 80
BATCH_SIZE = 16
LR = [1e-1, 1e-6]
FOLD_LIST = [1, 2, 3, 4, 5]
CROP_LENGTH = 133300
LOAD_DIR = "../models/envnet_model4_0"
OUTPUT_DIR = "../models/envnet_model5"

cudnn.benchmark = True
starttime = time.time()

```

```

def main():
    # load table data
    df_train = pd.read_csv("../input/train_curated.csv")
    df_noisy = pd.read_csv("../input/train_noisy.csv")
    df_test = pd.read_csv("../input/sample_submission.csv")
    labels = df_test.columns[1:].tolist()
    for label in labels:
        df_train[label] = df_train['labels'].apply(lambda x: label in x)
        df_noisy[label] = df_noisy['labels'].apply(lambda x: label in x)

    df_train['path'] = "../input/train_curated/" + df_train['fname']
    df_test['path'] = "../input/test/" + df_train['fname']
    df_noisy['path'] = "../input/train_noisy/" + df_noisy['fname']

    # fold splitting
    folds = list(KFold(n_splits=NUM_FOLD, shuffle=True, random_state=SEED).split(np.arange(len(df_train))))

    # Training
    log_columns = ['epoch', 'kl', 'lwlrp', 'kl_noisy', 'lwlrp_noisy', 'val_kl', 'val_lwlrp', 'time']
    for fold, (ids_train_split, ids_valid_split) in enumerate(folds):
        if fold+1 not in FOLD_LIST: continue
        print("fold: {}".format(fold + 1))
        train_log = pd.DataFrame(columns=log_columns)

        # build model
        model = EnvNetv2(NUM_CLASS).cuda()
        model.load_state_dict(torch.load("{}weight_fold{}_epoch_400.pth".format(LoadDir, fold+1)))

        # prepare data loaders
        df_train_fold = df_train.iloc[ids_train_split].reset_index(drop=True)
        dataset_train = WaveDataset(df_train_fold['path'], df_train_fold[labels].values,
                                    crop=CROP_LENGTH, crop_mode='random', padding=CROP_LENGTH//2,
                                    mixup=True, scaling=1.25, gain=6,
                                    )
        train_loader = DataLoader(dataset_train, batch_size=BATCH_SIZE,
                                   shuffle=True, num_workers=1, pin_memory=True,
                                   )

        df_valid = df_train.iloc[ids_valid_split].reset_index(drop=True)
        dataset_valid = WaveDataset(df_valid['path'], df_valid[labels].values,
                                    crop=CROP_LENGTH, crop_mode='random', padding=CROP_LENGTH//2,
                                    mixup=True, scaling=1.25, gain=6,
                                    )

```

```

abels].values, padding=CROP_LENGTH//2)
    valid_loader = DataLoader(dataset_valid, batch_size=1,
                              shuffle=False, num_workers=1,
pin_memory=True,
                                )

    dataset_noisy = WaveDataset(df_noisy['path'], df_noisy[labels].values,
                                crop=CROP_LENGTH, crop_mode=
'random', padding=CROP_LENGTH//2,
                                mixup=True, scaling=1.25, ga
in=6,
                                )
    noisy_loader = DataLoader(dataset_noisy, batch_size=BATCH_SIZE,
                              shuffle=True, num_workers=1, pin
in_memory=True,
                                )
    noisy_itr = cycle(noisy_loader)

    # set optimizer and loss
    optimizer = optim.SGD(filter(lambda p: p.requires_grad,
model.parameters()), lr=LR[0], momentum = 0.9, nesterov = True)
    scheduler = CosineLR(optimizer, step_size_min=LR[1], t0=
len(train_loader) * NUM_CYCLE, tmult=1)

    # training
    for epoch in range(NUM_EPOCH):
        # train for one epoch
        kl, lwlap, kl_noisy, lwlap_noisy = train((train_loader, noisy_itr), model, optimizer, scheduler, epoch)

        # evaluate on validation set
        val_kl, val_lwlap = validate(valid_loader, model)

        # print log
        endtime = time.time() - starttime
        print("Epoch: {}/{}".format(epoch + 1, NUM_EPOCH)
              + "KL: {:.4f} ".format(kl)
              + "LwLRAP: {:.4f} ".format(lwlap)
              + "Noisy KL: {:.4f} ".format(kl_noisy)
              + "Noisy LwLRAP: {:.4f} ".format(lwlap_noisy)
              + "Valid KL: {:.4f} ".format(val_kl)
              + "Valid LwLRAP: {:.4f} ".format(val_lwlap)
              + "sec: {:.1f}".format(endtime)
              )

        # save log and weights
        train_log_epoch = pd.DataFrame(
            [[epoch+1, kl, lwlap, kl_noisy, lwlap_noisy, v
al_kl, val_lwlap, endtime]],
            columns=log_columns)
        train_log = pd.concat([train_log, train_log_epoch])
        train_log.to_csv("{}train_log_fold{}.csv".format(OUTPUT_DIR, fold+1), index=False)
        if (epoch+1)%NUM_CYCLE==0:

```

```

        torch.save(model.state_dict(), "{}_weight_fold_{}_epoch_{}.pth".format(OUTPUT_DIR, fold+1, epoch+1))

def train(train_loaders, model, optimizer, scheduler, epoch):
    train_loader, noisy_itr = train_loaders
    kl_avr = AverageMeter()
    kl_noisy_avr = AverageMeter()
    lsigmoid = nn.LogSigmoid().cuda()
    lsoftmax = nn.LogSoftmax(dim=1).cuda()
    softmax = nn.Softmax(dim=1).cuda()
    criterion_kl = nn.KLDivLoss().cuda()

    # switch to train mode
    model.train()

    # training
    preds = np.zeros([0, NUM_CLASS], np.float32)
    y_true = np.zeros([0, NUM_CLASS], np.float32)
    preds_noisy = np.zeros([0, NUM_CLASS], np.float32)
    y_true_noisy = np.zeros([0, NUM_CLASS], np.float32)
    for i, (input, target) in enumerate(train_loader):
        # get batches
        input = torch.autograd.Variable(input.cuda())
        target = torch.autograd.Variable(target.cuda())

        input_noisy, target_noisy = next(noisy_itr)
        input_noisy = torch.autograd.Variable(input_noisy.cuda())
        target_noisy = torch.autograd.Variable(target_noisy.cuda())

        # compute output
        output = model(input)
        kl = criterion_kl(lsoftmax(output), target)
        output_noisy = model(input_noisy)
        kl_noisy = criterion_kl(lsoftmax(output_noisy), target_noisy)

        loss = kl + kl_noisy
        pred = softmax(output)
        pred = pred.data.cpu().numpy()
        pred_noisy = softmax(output_noisy)
        pred_noisy = pred_noisy.data.cpu().numpy()

        # backprop
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
        scheduler.step()

        # record log
        kl_avr.update(kl.data, input.size(0))
        kl_noisy_avr.update(kl_noisy.data, input.size(0))
        preds = np.concatenate([preds, pred])
        y_true = np.concatenate([y_true, target.data.cpu().numpy()])
    ]]
```

```

        preds_noisy = np.concatenate([preds_noisy, pred_noisy])
        y_true_noisy = np.concatenate([y_true_noisy, target_noisy.data.cpu().numpy()])

        # calc metric
        per_class_lwlap, weight_per_class = calculate_per_class_lwlap(y_true, preds)
        lwlap = np.sum(per_class_lwlap * weight_per_class)
        per_class_lwlap, weight_per_class = calculate_per_class_lwlap(y_true_noisy, preds_noisy)
        lwlap_noisy = np.sum(per_class_lwlap * weight_per_class)

        return kl_avr.avg.item(), lwlap, kl_noisy_avr.avg.item(), lwlap_noisy

def validate(val_loader, model):
    kl_avr = AverageMeter()
    lsoftmax = nn.LogSoftmax(dim=1).cuda()
    softmax = torch.nn.Softmax(dim=1).cuda()
    criterion_kl = nn.KLDivLoss().cuda()

    # switch to eval mode
    model.eval()

    # validate
    preds = np.zeros([0, NUM_CLASS], np.float32)
    y_true = np.zeros([0, NUM_CLASS], np.float32)
    for i, (input, target) in enumerate(val_loader):
        # get batches
        input = torch.autograd.Variable(input.cuda())
        target = torch.autograd.Variable(target.cuda())

        # compute output
        with torch.no_grad():
            output = model(input)
            kl = criterion_kl(lsoftmax(output), target)
            pred = softmax(output)
            pred = pred.data.cpu().numpy()

        # record log
        kl_avr.update(kl.data, input.size(0))
        preds = np.concatenate([preds, pred])
        y_true = np.concatenate([y_true, target.data.cpu().numpy()])

    # calc metric
    per_class_lwlap, weight_per_class = calculate_per_class_lwlap(y_true, preds)
    lwlap = np.sum(per_class_lwlap * weight_per_class)

    return kl_avr.avg.item(), lwlap

if __name__ == '__main__':
    main()
import numpy as np

```

```

import pandas as pd
import matplotlib.pyplot as plt
import time, random, sys
import sklearn.metrics
from sklearn.model_selection import KFold

import torch
import torch.nn as nn
import torch.backends.cudnn as cudnn
import torch.optim as optim
from torch.utils.data import DataLoader

sys.path.append('.')
from utils import AverageMeter, cycle, CosineLR, WaveDataset, calculate_per_class_lwlrwrap
from models import EnvNetv2

# set parameters
NUM_FOLD = 5
NUM_CLASS = 80
SEED = 42
NUM_EPOCH = 80 * 5
NUM_CYCLE = 80
BATCH_SIZE = 16
LR = [1e-1, 1e-6]
FOLD_LIST = [1, 2, 3, 4, 5]
CROP_LENGTH = 200000
OUTPUT_DIR = "../models/envnet_model6_0"

cudnn.benchmark = True
starttime = time.time()

def main():
    # load table data
    df_train = pd.read_csv("../input/train_curated.csv")
    df_noisy = pd.read_csv("../input/train_noisy.csv")
    df_test = pd.read_csv("../input/sample_submission.csv")
    labels = df_test.columns[1:].tolist()
    for label in labels:
        df_train[label] = df_train['labels'].apply(lambda x: label in x)
        df_noisy[label] = df_noisy['labels'].apply(lambda x: label in x)

    df_train['path'] = "../input/train_curated/" + df_train['fname']
    df_test['path'] = "../input/test/" + df_train['fname']
    df_noisy['path'] = "../input/train_noisy/" + df_noisy['fname']

    # fold splitting
    folds = list(KFold(n_splits=NUM_FOLD, shuffle=True, random_state=SEED).split(np.arange(len(df_train))))

    # Training

```

```

log_columns = ['epoch', 'kl', 'lwlap', 'kl_noisy', 'lwlap_
noisy', 'val_kl', 'val_lwlap', 'time']
for fold, (ids_train_split, ids_valid_split) in enumerate(fo
lds):
    if fold+1 not in FOLD_LIST: continue
    print("fold: {}".format(fold + 1))
    train_log = pd.DataFrame(columns=log_columns)

    # build model
    model = EnvNetv2(NUM_CLASS).cuda()

    # prepare data loaders
    df_train_fold = df_train.iloc[ids_train_split].reset_ind
ex(drop=True)
    dataset_train = WaveDataset(df_train_fold['path'], df_tr
ain_fold[labels].values,
                                crop=CROP_LENGTH, crop_mode=
'random', padding=CROP_LENGTH//2,
                                mixup=True, scaling=1.25, ga
in=6,
                                )
    train_loader = DataLoader(dataset_train, batch_size=BATC
H_SIZE,
                                shuffle=True, num_workers=1, p
in_memory=True,
                                )

    df_valid = df_train.iloc[ids_valid_split].reset_index(dr
op=True)
    dataset_valid = WaveDataset(df_valid['path'], df_valid[l
abels].values, padding=CROP_LENGTH//2)
    valid_loader = DataLoader(dataset_valid, batch_size=1,
                                shuffle=False, num_workers=1,
pin_memory=True,
                                )

    dataset_noisy = WaveDataset(df_noisy['path'], df_noisy[l
abels].values,
                                crop=CROP_LENGTH, crop_mode=
'random', padding=CROP_LENGTH//2,
                                mixup=True, scaling=1.25, ga
in=6,
                                )
    noisy_loader = DataLoader(dataset_noisy, batch_size=BATC
H_SIZE,
                                shuffle=True, num_workers=1, p
in_memory=True,
                                )
    noisy_itr = cycle(noisy_loader)

    # set optimizer and loss
    optimizer = optim.SGD(filter(lambda p: p.requires_grad,
model.parameters()), lr=LR[0], momentum = 0.9, nesterov = True)
    scheduler = CosineLR(optimizer, step_size_min=LR[1], t0=
len(train_loader) * NUM_CYCLE, tmult=1)

```



```

    # training
    for epoch in range(NUM_EPOCH):
        # train for one epoch
        kl, lwlap, kl_noisy, lwlap_noisy = train((train_loader, noisy_itr), model, optimizer, scheduler, epoch)

        # evaluate on validation set
        val_kl, val_lwlap = validate(valid_loader, model)

        # print log
        endtime = time.time() - starttime
        print("Epoch: {}/{} ".format(epoch + 1, NUM_EPOCH)
              + "KL: {:.4f} ".format(kl)
              + "LwLRAP: {:.4f} ".format(lwlap)
              + "Noisy KL: {:.4f} ".format(kl_noisy)
              + "Noisy LWLRAP: {:.4f} ".format(lwlap_noisy)
              + "Valid KL: {:.4f} ".format(val_kl)
              + "Valid LWLRAP: {:.4f} ".format(val_lwlap)
              + "sec: {:.1f}".format(endtime)
              )

        # save log and weights
        train_log_epoch = pd.DataFrame(
            [[epoch+1, kl, lwlap, kl_noisy, lwlap_noisy, val_kl, val_lwlap, endtime]],
            columns=log_columns)
        train_log = pd.concat([train_log, train_log_epoch])
        train_log.to_csv("{}train_log_fold{}.csv".format(OUTPUT_DIR, fold+1), index=False)
        if (epoch+1)%NUM_CYCLE==0:
            torch.save(model.state_dict(), "{}weight_fold{}_epoch{}.pth".format(OUTPUT_DIR, fold+1, epoch+1))

def train(train_loaders, model, optimizer, scheduler, epoch):
    train_loader, noisy_itr = train_loaders
    kl_avr = AverageMeter()
    kl_noisy_avr = AverageMeter()
    lsigmoid = nn.LogSigmoid().cuda()
    lsoftmax = nn.LogSoftmax(dim=1).cuda()
    softmax = nn.Softmax(dim=1).cuda()
    criterion_kl = nn.KLDivLoss().cuda()

    # switch to train mode
    model.train()

    # training
    preds = np.zeros([0, NUM_CLASS], np.float32)
    y_true = np.zeros([0, NUM_CLASS], np.float32)
    preds_noisy = np.zeros([0, NUM_CLASS], np.float32)
    y_true_noisy = np.zeros([0, NUM_CLASS], np.float32)
    for i, (input, target) in enumerate(train_loader):
        # get batches
        input = torch.autograd.Variable(input.cuda())
        target = torch.autograd.Variable(target.cuda())

```

```

        input_noisy, target_noisy = next(noisy_itr)
        input_noisy = torch.autograd.Variable(input_noisy.cuda())
    )
    target_noisy = torch.autograd.Variable(target_noisy.cuda()
    ())

    # compute output
    output = model(input)
    kl = criterion_kl(lsoftmax(output), target)
    output_noisy = model.noisy(input_noisy)
    kl_noisy = criterion_kl(lsoftmax(output_noisy), target_n
    oisy)

    loss = kl + kl_noisy
    pred = softmax(output)
    pred = pred.data.cpu().numpy()
    pred_noisy = softmax(output_noisy)
    pred_noisy = pred_noisy.data.cpu().numpy()

    # backprop
    optimizer.zero_grad()
    loss.backward()
    optimizer.step()
    scheduler.step()

    # record log
    kl_avr.update(kl.data, input.size(0))
    kl_noisy_avr.update(kl_noisy.data, input.size(0))
    preds = np.concatenate([preds, pred])
    y_true = np.concatenate([y_true, target.data.cpu().numpy
    ()])

    preds_noisy = np.concatenate([preds_noisy, pred_noisy])
    y_true_noisy = np.concatenate([y_true_noisy, target_nois
    y.data.cpu().numpy()])

    # calc metric
    per_class_lwlrp, weight_per_class = calculate_per_class_lwl
    rap(y_true, preds)
    lwlrp = np.sum(per_class_lwlrp * weight_per_class)
    per_class_lwlrp, weight_per_class = calculate_per_class_lwl
    rap(y_true_noisy, preds_noisy)
    lwlrp_noisy = np.sum(per_class_lwlrp * weight_per_class)

    return kl_avr.avg.item(), lwlrp, kl_noisy_avr.avg.item(), l
    wlrp_noisy

def validate(val_loader, model):
    kl_avr = AverageMeter()
    lsoftmax = nn.LogSoftmax(dim=1).cuda()
    softmax = torch.nn.Softmax(dim=1).cuda()
    criterion_kl = nn.KLDivLoss().cuda()

    # switch to eval mode
    model.eval()

    # validate

```

```

preds = np.zeros([0, NUM_CLASS], np.float32)
y_true = np.zeros([0, NUM_CLASS], np.float32)
for i, (input, target) in enumerate(val_loader):
    # get batches
    input = torch.autograd.Variable(input.cuda())
    target = torch.autograd.Variable(target.cuda())

    # compute output
    with torch.no_grad():
        output = model(input)
        kl = criterion_kl(softmax(output), target)
        pred = softmax(output)
        pred = pred.data.cpu().numpy()

    # record log
    kl_avr.update(kl.data, input.size(0))
    preds = np.concatenate([preds, pred])
    y_true = np.concatenate([y_true, target.data.cpu().numpy()])

# calc metric
per_class_lwlap, weight_per_class = calculate_per_class_lwlap(y_true, preds)
lwlap = np.sum(per_class_lwlap * weight_per_class)

return kl_avr.avg.item(), lwlap

if __name__ == '__main__':
    main()
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import time, random, sys
import sklearn.metrics
from sklearn.model_selection import KFold

import torch
import torch.nn as nn
import torch.backends.cudnn as cudnn
import torch.optim as optim
from torch.utils.data import DataLoader

sys.path.append('.')
from utils import AverageMeter, cycle, CosineLR, WaveDataset, calculate_per_class_lwlap
from models import EnvNetv2

# set parameters
NUM_FOLD = 5
NUM_CLASS = 80
SEED = 42
NUM_EPOCH = 80 * 5
NUM_CYCLE = 80
BATCH_SIZE = 16
LR = [1e-1, 1e-6]
FOLD_LIST = [1, 2, 3, 4, 5]

```

```

CROP_LENGTH = 200000
LOAD_DIR = "../models/envnet_model6_0"
OUTPUT_DIR = "../models/envnet_model6"

cudnn.benchmark = True
starttime = time.time()

def main():
    # load table data
    df_train = pd.read_csv("../input/train_curated.csv")
    df_noisy = pd.read_csv("../input/train_noisy.csv")
    df_test = pd.read_csv("../input/sample_submission.csv")
    labels = df_test.columns[1:].tolist()
    for label in labels:
        df_train[label] = df_train['labels'].apply(lambda x: label in x)
        df_noisy[label] = df_noisy['labels'].apply(lambda x: label in x)

    df_train['path'] = "../input/train_curated/" + df_train['fname']
    df_test['path'] = "../input/test/" + df_train['fname']
    df_noisy['path'] = "../input/train_noisy/" + df_noisy['fname']

    # fold splitting
    folds = list(KFold(n_splits=NUM_FOLD, shuffle=True, random_state=SEED).split(np.arange(len(df_train))))

    # Training
    log_columns = ['epoch', 'kl', 'lwlrwrap', 'kl_noisy', 'lwlrwrap_noisy', 'val_kl', 'val_lwlrwrap', 'time']
    for fold, (ids_train_split, ids_valid_split) in enumerate(folds):
        if fold+1 not in FOLD_LIST: continue
        print("fold: {}".format(fold + 1))
        train_log = pd.DataFrame(columns=log_columns)

        # build model
        model = EnvNetv2(NUM_CLASS).cuda()
        model.load_state_dict(torch.load("{}weight_fold{}_epoch_400.pth".format(LOAD_DIR, fold+1)))

        # prepare data loaders
        df_train_fold = df_train.iloc[ids_train_split].reset_index(drop=True)
        dataset_train = WaveDataset(df_train_fold['path'], df_train_fold[labels].values,
                                    crop=CROP_LENGTH, crop_mode='random', padding=CROP_LENGTH//2,
                                    mixup=True, scaling=1.25, gain=6,
                                    )
        train_loader = DataLoader(dataset_train, batch_size=BATCH_SIZE,

```

```

shuffle=True, num_workers=1, p
in_memory=True,
    )

    df_valid = df_train.iloc[ids_valid_split].reset_index(drop=True)
    dataset_valid = WaveDataset(df_valid['path'], df_valid[labels].values, padding=CROP_LENGTH//2)
    valid_loader = DataLoader(dataset_valid, batch_size=1, shuffle=False, num_workers=1,
pin_memory=True,
    )

    dataset_noisy = WaveDataset(df_noisy['path'], df_noisy[labels].values,
                                crop=CROP_LENGTH, crop_mode='random', padding=CROP_LENGTH//2,
                                mixup=True, scaling=1.25, gain=6,
    )
    noisy_loader = DataLoader(dataset_noisy, batch_size=BATCH_SIZE,
                                shuffle=True, num_workers=1, pin_memory=True,
    )
    noisy_itr = cycle(noisy_loader)

    # set optimizer and loss
    optimizer = optim.SGD(filter(lambda p: p.requires_grad, model.parameters()), lr=LR[0], momentum = 0.9, nesterov = True)
    scheduler = CosineLR(optimizer, step_size_min=LR[1], t0=len(train_loader) * NUM_CYCLE, tmult=1)

    # training
    for epoch in range(NUM_EPOCH):
        # train for one epoch
        kl, lwlap, kl_noisy, lwlap_noisy = train((train_loader, noisy_itr), model, optimizer, scheduler, epoch)

        # evaluate on validation set
        val_kl, val_lwlap = validate(valid_loader, model)

        # print log
        endtime = time.time() - starttime
        print("Epoch: {}/{}".format(epoch + 1, NUM_EPOCH)
              + "KL: {:.4f}".format(kl)
              + "LwLRAP: {:.4f}".format(lwlap)
              + "Noisy KL: {:.4f}".format(kl_noisy)
              + "Noisy LWLRAP: {:.4f}".format(lwlap_noisy)
              + "Valid KL: {:.4f}".format(val_kl)
              + "Valid LWLRAP: {:.4f}".format(val_lwlap)
              + "sec: {:.1f}".format(endtime)
        )

        # save log and weights
        train_log_epoch = pd.DataFrame(

```

```

        [[epoch+1, kl, lwlap, kl_noisy, lwlap_noisy, v
al_kl, val_lwlap, endtime]],
        columns=log_columns)
        train_log = pd.concat([train_log, train_log_epoch])
        train_log.to_csv("{}train_log_fold{}.csv".format(OU
TPUT_DIR, fold+1), index=False)
        if (epoch+1)%NUM_CYCLE==0:
            torch.save(model.state_dict(), "{}weight_fold_
{}_epoch{}.pth".format(OUTPUT_DIR, fold+1, epoch+1))

def train(train_loaders, model, optimizer, scheduler, epoch):
    train_loader, noisy_itr = train_loaders
    kl_avr = AverageMeter()
    kl_noisy_avr = AverageMeter()
    lsigmoid = nn.LogSigmoid().cuda()
    lsoftmax = nn.LogSoftmax(dim=1).cuda()
    softmax = nn.Softmax(dim=1).cuda()
    criterion_kl = nn.KLDivLoss().cuda()

    # switch to train mode
    model.train()

    # training
    preds = np.zeros([0, NUM_CLASS], np.float32)
    y_true = np.zeros([0, NUM_CLASS], np.float32)
    preds_noisy = np.zeros([0, NUM_CLASS], np.float32)
    y_true_noisy = np.zeros([0, NUM_CLASS], np.float32)
    for i, (input, target) in enumerate(train_loader):
        # get batches
        input = torch.autograd.Variable(input.cuda())
        target = torch.autograd.Variable(target.cuda())

        input_noisy, target_noisy = next(noisy_itr)
        input_noisy = torch.autograd.Variable(input_noisy.cuda()
)
        target_noisy = torch.autograd.Variable(target_noisy.cuda
())

        # compute output
        output = model(input)
        kl = criterion_kl(lsoftmax(output), target)
        output_noisy = model.noisy(input_noisy)
        kl_noisy = criterion_kl(lsoftmax(output_noisy), target_n
oisy)
        loss = kl + kl_noisy
        pred = softmax(output)
        pred = pred.data.cpu().numpy()
        pred_noisy = softmax(output_noisy)
        pred_noisy = pred_noisy.data.cpu().numpy()

        # backprop
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
        scheduler.step()

```

```

        # record log
        kl_avr.update(kl.data, input.size(0))
        kl_noisy_avr.update(kl_noisy.data, input.size(0))
        preds = np.concatenate([preds, pred])
        y_true = np.concatenate([y_true, target.data.cpu().numpy
    (]])

        preds_noisy = np.concatenate([preds_noisy, pred_noisy])
        y_true_noisy = np.concatenate([y_true_noisy, target_nois
y.data.cpu().numpy()])

        # calc metric
        per_class_lwlraps, weight_per_class = calculate_per_class_lwlraps(y_true, preds)
        lwlraps = np.sum(per_class_lwlraps * weight_per_class)
        per_class_lwlraps, weight_per_class = calculate_per_class_lwlraps(y_true_noisy, preds_noisy)
        lwlraps_noisy = np.sum(per_class_lwlraps * weight_per_class)

        return kl_avr.avg.item(), lwlraps, kl_noisy_avr.avg.item(), lwlraps_noisy

def validate(val_loader, model):
    kl_avr = AverageMeter()
    lsoftmax = nn.LogSoftmax(dim=1).cuda()
    softmax = torch.nn.Softmax(dim=1).cuda()
    criterion_kl = nn.KLDivLoss().cuda()

    # switch to eval mode
    model.eval()

    # validate
    preds = np.zeros([0, NUM_CLASS], np.float32)
    y_true = np.zeros([0, NUM_CLASS], np.float32)
    for i, (input, target) in enumerate(val_loader):
        # get batches
        input = torch.autograd.Variable(input.cuda())
        target = torch.autograd.Variable(target.cuda())

        # compute output
        with torch.no_grad():
            output = model(input)
            kl = criterion_kl(lsoftmax(output), target)
            pred = softmax(output)
            pred = pred.data.cpu().numpy()

        # record log
        kl_avr.update(kl.data, input.size(0))
        preds = np.concatenate([preds, pred])
        y_true = np.concatenate([y_true, target.data.cpu().numpy
    (]])

    # calc metric
    per_class_lwlraps, weight_per_class = calculate_per_class_lwlraps(y_true, preds)

```

```

    lwlrap = np.sum(per_class_lwlrap * weight_per_class)

    return kl_avr.avg.item(), lwlrap

if __name__ == '__main__':
    main()
import numpy as np
from torch.optim.lr_scheduler import _LRScheduler
from torch.utils.data.dataset import Dataset
from math import cos, pi
import librosa
from scipy.io import wavfile
import random

class AverageMeter(object):
    """Computes and stores the average and current value"""

    def __init__(self):
        self.reset()

    def reset(self):
        self.val = 0
        self.avg = 0
        self.sum = 0
        self.count = 0

    def update(self, val, n=1):
        self.val = val
        self.sum += val * n
        self.count += n
        self.avg = self.sum / self.count

def cycle(iterable):
    """
    convert dataloader to iterator
    :param iterable:
    :return:
    """
    while True:
        for x in iterable:
            yield x

class CosineLR(_LRScheduler):
    """cosine annealing.
    """
    def __init__(self, optimizer, step_size_min=1e-5, t0=100, tmult=2, curr_epoch=-1, last_epoch=-1):
        self.step_size_min = step_size_min
        self.t0 = t0
        self.tmult = tmult
        self.epochs_since_restart = curr_epoch
        super(CosineLR, self).__init__(optimizer, last_epoch)

    def get_lr(self):

```



```

        self.epochs_since_restart += 1

        if self.epochs_since_restart > self.t0:
            self.t0 *= self.tmult
            self.epochs_since_restart = 0

        lrs = [self.step_size_min + (
            0.5 * (base_lr - self.step_size_min) * (1 + cos(
self.epochs_since_restart * pi / self.t0)))
            for base_lr in self.base_lrs]

        return lrs

class MelDataset(Dataset):
    def __init__(self, X, y, crop=-1,
        mixup=False, freqmask=False, gain=False,
        crop_mode='original', crop_rate=0.25
    ):
        self.X= X
        self.y= y
        self.crop = crop
        self.mixup = mixup
        self.freqmask = freqmask
        self.gain = gain
        self.crop_mode = crop_mode
        self.crop_rate = crop_rate

    def do_additional_crop(self, img):
        len_img = img.shape[1]
        img_new = np.zeros([img.shape[0], self.crop], np.float32
)
        rate = np.random.random() * (1 - self.crop_rate) + self.
crop_rate
        if np.random.random() < 0.5: rate = 1

        if img.shape[1] <= self.crop:
            len_crop = int(img.shape[1] * rate)
            if img.shape[1] - len_crop == 0:
                shift_crop = 0
            else:
                shift_crop = np.random.randint(0, img.shape[1] -
len_crop)
            img = img[:, shift_crop:shift_crop + len_crop]
            if self.crop - len_crop == 0:
                shift = 0
            else:
                shift = np.random.randint(0, self.crop - len_cro
p)
            img_new[:, shift:shift + len_crop] = img
        else:
            shift = np.random.randint(0, img.shape[1] - self.cro
p)

            img_new = img[:, shift:shift + self.crop]
            len_crop = int(self.crop * rate)
            if self.crop - len_crop == 0:

```

```

        shift_crop = 0
    else:
        shift_crop = np.random.randint(0, self.crop - len
n_crop)

        img_new[:shift_crop] = 0
        img_new[shift_crop + len_crop:] = 0
    return img_new

def do_random_crop(self, img):
    img_new = np.zeros([img.shape[0], self.crop], np.float32
)

    if img.shape[1] < self.crop:
        shift = np.random.randint(0, self.crop - img.shape[1
])

        img_new[:, shift:shift + img.shape[1]] = img
    elif img.shape[1] == self.crop:
        img_new = img
    else:
        shift = np.random.randint(0, img.shape[1] - self.cro
p)

        img_new = img[:, shift:shift + self.crop]
    return img_new

def do_crop(self, img):
    if self.crop_mode == 'random':
        return self.do_random_crop(img)
    elif self.crop_mode == 'additional':
        return self.do_additional_crop(img)
    elif self.crop_mode == 'original':
        return img

def do_mixup(self, img, label, alpha=1.):
    idx = np.random.randint(0, len(self.X))
    img2 = np.load("{}_np.npy".format(self.X[idx][:4]))
    img2 = self.do_crop(img2)

    label2 = self.y[idx].astype(np.float32)

    rate = np.random.beta(alpha, alpha)
    img = img * rate + img2 * (1 - rate)
    label = label * rate + label2 * (1 - rate)
    return img, label

def do_freqmask(self, img, max=32):
    coord = np.random.randint(0, img.shape[0])
    width = np.random.randint(8, max)
    cut = np.array([coord - width, coord + width])
    cut = np.clip(cut, 0, img.shape[0])
    img[cut[0]:cut[1]] = 0
    return img

def do_gain(self, img, max=0.1):
    rate = 1 - max + np.random.random() * max * 2
    return img * rate

```

```

def __getitem__(self, index):
    img = np.load("{}_npy".format(self.X[index][:4]))
    img = self.do_crop(img)
    label = self.y[index].astype(np.float32)

    if self.mixup and np.random.random() < 0.5:
        img, label = self.do_mixup(img, label)
    if self.gain and np.random.random() < 0.5:
        img = self.do_gain(img)
    if self.freqmask and np.random.random() < 0.5:
        img = self.do_freqmask(img)

    img = librosa.power_to_db(img)
    img = (img - img.mean()) / (img.std() + 1e-7)
    img = img.reshape([1, img.shape[0], img.shape[1]])

    return img, label

def __len__(self):
    return len(self.X)

def compute_gain(sound, fs, min_db=-80.0, mode='RMSE'):
    if fs == 16000:
        n_fft = 2048
    elif fs == 44100:
        n_fft = 4096
    else:
        raise Exception('Invalid fs {}'.format(fs))
    stride = n_fft // 2

    gain = []
    for i in range(0, len(sound) - n_fft + 1, stride):
        if mode == 'RMSE':
            g = np.mean(sound[i: i + n_fft] ** 2)
        elif mode == 'A_weighting':
            spec = np.fft.rfft(np.hanning(n_fft + 1)[:n_fft] * sound[i: i + n_fft])
            power_spec = np.abs(spec) ** 2
            a_weighted_spec = power_spec * np.power(10, a_weighting(fs, n_fft) / 10)
            g = np.sum(a_weighted_spec)
        else:
            raise Exception('Invalid mode {}'.format(mode))
        gain.append(g)

    gain = np.array(gain)
    gain = np.maximum(gain, np.power(10, min_db / 10))
    gain_db = 10 * np.log10(gain)

    return gain_db

def mix(sound1, sound2, r, fs):
    gain1 = np.max(compute_gain(sound1, fs)) # Decibel
    gain2 = np.max(compute_gain(sound2, fs))

```

```

    t = 1.0 / (1 + np.power(10, (gain1 - gain2) / 20.) * (1 - r)
    / r)
    sound = ((sound1 * t + sound2 * (1 - t)) / np.sqrt(t ** 2 +
(1 - t) ** 2))
    sound = sound.astype(np.float32)

    return sound

```

```

class WaveDataset(Dataset):
    def __init__(self, X, y,
                 crop=-1, crop_mode='original', padding=0,
                 mixup=False, scaling=-1, gain=-1,
                 fs=44100,
                 ):
        self.X = X
        self.y = y
        self.crop = crop
        self.crop_mode = crop_mode
        self.padding = padding
        self.mixup = mixup
        self.scaling = scaling
        self.gain = gain
        self.fs = fs

    def preprocess(self, sound):
        for f in self.preprocess_funcs:
            sound = f(sound)

        return sound

    def do_padding(self, snd):
        snd_new = np.pad(snd, self.padding, 'constant')
        return snd_new

    def do_crop(self, snd):
        if self.crop_mode=='random':
            shift = np.random.randint(0, snd.shape[0] - self.cro
p)

            snd_new = snd[shift:shift + self.crop]
        else:
            snd_new = snd
        return snd_new

    def do_gain(self, snd):
        snd_new = snd * np.power(10, random.uniform(-self.gain,
self.gain) / 20.0)
        return snd_new

    def do_scaling(self, snd, interpolate='Nearest'):
        scale = np.power(self.scaling, random.uniform(-1, 1))
        output_size = int(len(snd) * scale)
        ref = np.arange(output_size) / scale
        if interpolate == 'Linear':
            ref1 = ref.astype(np.int32)
            ref2 = np.minimum(ref1+1, len(snd)-1)

```

```

        r = ref - ref1
        snd_new = snd[ref1] * (1-r) + snd[ref2] * r
    elif interpolate == 'Nearest':
        snd_new = snd[ref.astype(np.int32)]
    else:
        raise Exception('Invalid interpolation mode {}'.format(interpolate))

    return snd_new

def do_mixup(self, snd, label, alpha=1):
    idx2 = np.random.randint(0, len(self.X))
    _, snd2 = wavfile.read("{}".format(self.X[idx2]))
    label2 = self.y[idx2].astype(np.float32)
    if self.scaling != -1:
        snd2 = self.do_scaling(snd2)
    snd2 = self.do_padding(snd2)
    snd2 = self.do_crop(snd2)

    rate = np.random.beta(alpha, alpha)
    snd_new = mix(snd, snd2, rate, self.fs)
    label_new = label * rate + label2 * (1 - rate)
    return snd_new, label_new

def __getitem__(self, index):
    _, snd = wavfile.read("{}".format(self.X[index]))
    label = self.y[index].astype(np.float32)
    if self.scaling != -1:
        snd = self.do_scaling(snd)
    snd = self.do_padding(snd)
    snd = self.do_crop(snd)
    if self.mixup:
        snd, label = self.do_mixup(snd, label)
    if self.gain != -1:
        snd = self.do_gain(snd)
    snd = snd.reshape([1, 1, -1]).astype(np.float32) / 32768
    .0
    return snd, label

def __len__(self):
    return len(self.X)

def _one_sample_positive_class_precisions(scores, truth):
    """Calculate precisions for each true class for a single sample.

    Args:
        scores: np.array of (num_classes,) giving the individual classifier scores.
        truth: np.array of (num_classes,) bools indicating which classes are true.

    Returns:
        pos_class_indices: np.array of indices of the true classes for this sample.

```

```

    pos_class_precisions: np.array of precisions corresponding
    to each of those
    classes.
    """
    num_classes = scores.shape[0]
    pos_class_indices = np.flatnonzero(truth > 0)
    # Only calculate precisions if there are some true classes.
    if not len(pos_class_indices):
        return pos_class_indices, np.zeros(0)
    # Retrieval list of classes for this sample.
    retrieved_classes = np.argsort(scores)[::-1]
    # class_rankings[top_scoring_class_index] == 0 etc.
    class_rankings = np.zeros(num_classes, dtype=np.int)
    class_rankings[retrieved_classes] = range(num_classes)
    # Which of these is a true label?
    retrieved_class_true = np.zeros(num_classes, dtype=np.bool)
    retrieved_class_true[class_rankings[pos_class_indices]] = True
    # Num hits for every truncated retrieval list.
    retrieved_cumulative_hits = np.cumsum(retrieved_class_true)
    # Precision of retrieval list truncated at each hit, in order
    # of pos_labels.
    precision_at_hits = (
        retrieved_cumulative_hits[class_rankings[pos_class_i
indices]] /
        (1 + class_rankings[pos_class_indices].astype(np.flo
at)))
    return pos_class_indices, precision_at_hits

# All-in-one calculation of per-class lwlap.

def calculate_per_class_lwlap(truth, scores):
    """Calculate label-weighted label-ranking average precision.

    Arguments:
        truth: np.array of (num_samples, num_classes) giving boole
an ground-truth
        of presence of that class in that sample.
        scores: np.array of (num_samples, num_classes) giving the
classifier-under-
        test's real-valued score for each class for each sample.

    Returns:
        per_class_lwlap: np.array of (num_classes,) giving the lw
lap for each
        class.
        weight_per_class: np.array of (num_classes,) giving the pr
ior of each
        class within the truth labels. Then the overall unbalan
ced lwlap is
        simply np.sum(per_class_lwlap * weight_per_class)
    """
    assert truth.shape == scores.shape
    num_samples, num_classes = scores.shape
    # Space to store a distinct precision value for each class o

```

```

n each sample.
    # Only the classes that are true for each sample will be filled in.
    precisions_for_samples_by_classes = np.zeros((num_samples, num_classes))
    for sample_num in range(num_samples):
        pos_class_indices, precision_at_hits = (
            _one_sample_positive_class_precisions(scores[sample_num, :],
                                                    truth[sample_num, :]))
        precisions_for_samples_by_classes[sample_num, pos_class_indices] = (
            precision_at_hits)
        labels_per_class = np.sum(truth > 0, axis=0)
        weight_per_class = labels_per_class / float(np.sum(labels_per_class))
        # Form average of each column, i.e. all the precisions assigned to labels in
        # a particular class.
        per_class_lwlap = (np.sum(precisions_for_samples_by_classes, axis=0) /
                           np.maximum(1, labels_per_class))
        # overall_lwlap = simple average of all the actual per-class, per-sample precisions
        # = np.sum(precisions_for_samples_by_classes) / np.sum(precisions_for_samples_by_classes > 0)
        # also = weighted mean of per-class lwlaps, weighted by class label prior across samples
        # = np.sum(per_class_lwlap * weight_per_classes)
    return per_class_lwlap, weight_per_class

```