

9th-place-modeling-kernel

January 11, 2020

This kernel is used for training our models. In the end, we went with an ensemble of 15 models similar to the one trained here.

Their CV are all above 0.86 and 5-folds LB varies between 0.71 and 0.72

For our solution writeup, please check : <https://www.kaggle.com/c/freesound-audio-tagging-2019/discussion/95409#latest-551352>

1 Initialization

1.1 Imports

```
[1]: import re
import gc
import os
import time
import pywt
import pickle
import random
import operator
import warnings
import numpy as np
import pandas as pd
import seaborn as sns
import tensorflow as tf
import IPython.display as ipd
import matplotlib.pyplot as plt

from PIL import Image
from sklearn.metrics import *
from collections import Counter
from sklearn.model_selection import *
from tqdm import tqdm_notebook as tqdm
from keras.utils import to_categorical
from multiprocessing import Pool, cpu_count
from collections import OrderedDict
```

```

import torch
import torch.nn as nn
import torch.utils.data
import torch.nn.functional as F
import torchvision.models as models
from torch.autograd import Variable
from torch.optim.lr_scheduler import *
from torchvision.transforms import transforms
from torch.utils.data import Dataset, DataLoader

import librosa
import librosa.display

begin = time.time()
sns.set_style('white')
warnings.simplefilter(action='ignore', category=FutureWarning)
warnings.simplefilter(action='ignore', category=UserWarning)
warnings.simplefilter(action='ignore', category=RuntimeWarning)

print("Number of available cpu cores: {}".format(cpu_count()))

```

Using TensorFlow backend.

Number of available cpu cores: 2

1.2 Seeding

```

[2]: def seed_everything(seed):
    random.seed(seed)
    os.environ['PYTHONHASHSEED'] = str(seed)
    np.random.seed(seed)
    torch.manual_seed(seed)
    torch.cuda.manual_seed(seed)
    torch.backends.cudnn.deterministic = True
    tf.set_random_seed(seed)

```

```

[3]: seed = 2019
    seed_everything(seed)

```

```

[4]: N_JOBS = cpu_count()
    os.environ['MKL_NUM_THREADS'] = str(N_JOBS)
    os.environ['OMP_NUM_THREADS'] = str(N_JOBS)
    DataLoader = partial(DataLoader, num_workers=N_JOBS)

```

1.3 Load Data

```
[5]: DATA_PATH = '../input/freesound-audio-tagging-2019/'  
TRAIN_CURATED_PATH = DATA_PATH + 'train_curated/'  
TRAIN_NOISY_PATH = DATA_PATH + 'train_noisy/'  
TEST_PATH = DATA_PATH + 'test/'
```

```
[6]: df_train = pd.read_csv(DATA_PATH + 'train_curated.csv')  
df_noisy = pd.read_csv(DATA_PATH + 'train_noisy.csv')  
df_test = pd.read_csv(DATA_PATH + 'sample_submission.csv')  
  
df_train.head()
```

```
[6]:      fname      labels  
0  0006ae4e.wav      Bark  
1  0019ef41.wav  Raindrop  
2  001ec0ad.wav  Finger_snapping  
3  0026c7cb.wav      Run  
4  0026f116.wav  Finger_snapping
```

```
[7]: labels = list(df_test.columns[1:])  
num_classes = len(labels)  
  
print("Number of classes :", num_classes)
```

Number of classes : 80

Removing “corrupted” files

```
[8]: to_remove = ["f76181c4.wav", "77b925c2.wav", "6a1f682a.wav", "c7db12aa.wav",  
↪ "7752cc8a.wav", "1d44b0bd.wav"]  
  
for i in df_train.copy().index:  
    if df_train['fname'][i] in to_remove:  
        df_train.drop(i, inplace=True)  
  
df_train = df_train.reset_index(drop=True)
```

Restraining noisy data to samples with only one label

```
[9]: df_noisy["nb_labels"] = df_noisy["labels"].apply(lambda x: len(x.split(',')))  
df_noisy = df_noisy[df_noisy["nb_labels"] == 1].copy().reset_index(drop=True)
```

2 Signal Processing

2.1 Config

```
[10]: class config:
    sampling_rate = 44100 # 44.1 kHz
    duration = 4 #2 # Minimum length for short samples (seconds)
    samples = sampling_rate * duration # Minimum sample size

    top_db = 60 # Noise filtering, default = 60

    # Frequencies kept in spectrograms
    fmin = 20
    fmax = sampling_rate // 2 # Shannon theorem

    # Spectrogram parameters
    n_mels = 64 # = spec_height
    n_fft = n_mels * 30 # Size of fft window - smooths the spectrogram
    spec_min_width = 256 #128
    x_mean, x_std = -35.7, 21.6
    hop_length = duration * sampling_rate // spec_min_width + 1 # Number of
    ↪ samples between each frame - impacts y size of spectrogram
```

2.2 Read Audio

```
[11]: def read_audio(pathname, conf, trim_long_data):
    y, sr = librosa.load(pathname, sr=conf.sampling_rate)
    # trim silence
    if len(y) > 0: # workaround: 0 length causes error
        y, _ = librosa.effects.trim(y) # trim, top_db=default(60)
    # make it unified length to conf.samples
    if len(y) > conf.samples: # long enough
        if trim_long_data: y = y[0:0+conf.samples]
    else: # pad blank
        padding = conf.samples - len(y) # add padding at both ends
        offset = padding // 2
        y = np.pad(y, (offset, conf.samples - len(y) - offset), 'constant')
    return y
```

2.3 MEL Spectrogram

The three channel mode adds first and second or delta to the data. It did not improve the results.

```
[12]: def audio_to_melspectrogram(audio, config, three_channels=False):
    spectrogram = librosa.feature.melspectrogram(audio,
                                                  sr=config.sampling_rate,
                                                  n_mels=config.n_mels,
                                                  hop_length=config.hop_length,
                                                  n_fft=config.n_fft,
                                                  fmin=config.fmin,
                                                  fmax=config.fmax)

    logmel = librosa.power_to_db(spectrogram).astype(np.float32)

    if three_channels:
        return np.array([logmel, librosa.feature.delta(logmel), librosa.feature.
→delta(logmel, order=2)])
    else:
        return logmel
```

2.4 Normalize

We can either use individual image statistics or statistics computed on train data.

```
[13]: def normalize(X, mean=None, std=None):
    mean = mean or X.mean()
    std = std or (X-X.mean()).std()
    return ((X - mean)/std).astype(np.float16)
```

3 Labels

```
[14]: def get_labels(df, labels):
    mapping = {label:i for i, label in enumerate(labels)}
    y = np.zeros((len(df), len(labels)))
    all_labels = df['labels'].apply(lambda x: x.split(','))
    for i, label in enumerate(all_labels):
        for l in label:
            y[i, mapping[l]] = 1
    return y.astype(int)
```

```
[15]: y_train = get_labels(df_train, labels)
```

3.1 Select best noisy samples

We load predictions that were computed using a model that was trained only on curated data. The idea is to keep the ones that were the most correctly predicted, using the log loss as metric. As

some class were better predicted than others, we restrict the number of added samples to 50 per class.

```
[16]: scores_noisy = np.load("../input/fat-cp/scores.npy")
      y_noisy = get_labels(df_noisy, labels)

[17]: def sort_by_loss(y_noisy, scores_noisy):
      losses_dic = {i : log_loss(y_noisy[i], scores_noisy[i]) for i in
      ↪range(scores_noisy.shape[0])}
      sorted_dic = sorted(losses_dic.items(), key=operator.itemgetter(1))
      return sorted_dic

[18]: sorted_by_loss = sort_by_loss(y_noisy, scores_noisy)

[19]: def select_noisy(sorted_by_loss, nb_noisy, max_per_class=50):
      selected = np.array(sorted_by_loss[:nb_noisy, 0]).astype(int)
      to_keep = []
      counts = {}
      for s in selected:
          l = df_noisy["labels"][s]
          try:
              counts[l] += 1
          except:
              counts[l] = 0
          if counts[l] < max_per_class:
              to_keep.append(s)

      return df_noisy.iloc[to_keep].reset_index(drop=True)

      df_noisy_selected

[20]: nb_noisy = 5000

[21]: df_noisy_selected = select_noisy(sorted_by_loss, nb_noisy, max_per_class=50)
      y_noisy = get_labels(df_noisy_selected, labels)
```

3.2 Input data

```
[22]: three_channels = False
      crop = False

[23]: def process(df, path, config, crop=False, three_channels=False):
      X = []
      for i in df.index:
          signal = read_audio(path + df['fname'][i], config, crop)
```

```

        X.append(normalize(audio_to_melspectrogram(signal, config), config.
↪x_mean, config.x_std)) #normalize based on global statistics
        #X.append(normalize(audio_to_melspectrogram(signal, config)))↪
↪#normalize based on individual statistics
    return X

```

```

[24]: # df_train = df_train.head(100)
      # df_test = df_test.head(100)
      # df_noisy_selected = df_noisy_selected.head(100)

```

```

[25]: %%time
      X_train = process(df_train, TRAIN_CURATED_PATH, config, crop=crop,↪
↪three_channels=three_channels)
      X_noisy = process(df_noisy_selected, TRAIN_NOISY_PATH, config, crop=crop,↪
↪three_channels=three_channels)
      X_test = process(df_test, TEST_PATH, config, three_channels=three_channels)


```

CPU times: user 10min 3s, sys: 8min 7s, total: 18min 11s
 Wall time: 9min 48s

```

[26]: plt.matshow(X_train[0].astype(np.float32))
      plt.show()

```



3.3 Datasets

3.3.1 Augmentation

We augment spectrograms by hiding random time & frequency intervals

```

[27]: def spec_augment(spec: np.ndarray, num_mask=2, freq_masking_max_percentage=0.
↪15, time_masking_max_percentage=0.3):
      spec = spec.copy()
      for i in range(num_mask):
          num_freqs, num_frames = spec.shape
          freq_percentage = random.uniform(0.0, freq_masking_max_percentage)
          time_percentage = random.uniform(0.0, time_masking_max_percentage)

```

```

num_freqs_to_mask = int(freq_percentage * num_freqs)
num_frames_to_mask = int(time_percentage * num_frames)

t0 = int(np.random.uniform(low=0.0, high=num_frames -
↪ num_frames_to_mask))
f0 = int(np.random.uniform(low=0.0, high=num_freqs - num_freqs_to_mask))

spec[:, t0:t0 + num_frames_to_mask] = 0
spec[f0:f0 + num_freqs_to_mask, :] = 0


return spec

```

```

[28]: plt.matshow(spec_augment(X_train[0].astype(np.float32)))
plt.show()

```



3.3.2 Usual transforms

Only the conversion to tensor is kept, we also tried flipping and resized cropping.

```

[29]: transforms_dict = {
    'train': transforms.Compose([
        transforms.ToTensor(),
    ]),
    'test': transforms.Compose([
        transforms.ToTensor(),
    ]),
    'viz': transforms.Compose([
    ]),
}

```

3.3.3 Train

```

[30]: class FATDatasetTrain(Dataset):
    def __init__(self, mels, transforms, y, apply_spec_aug=False):
        super().__init__()
        self.mels = mels

```



```

self.labels = y
self.transforms = transforms
self.apply_spec_aug = apply_spec_aug

def __len__(self):
    return len(self.mels)

def __getitem__(self, idx):
    data = self.mels[idx].astype(np.float32)
    base_dim, time_dim = data.shape
    crop = random.randint(0, max(time_dim - config.spec_min_width, 0))
    data = data[0:base_dim, crop:crop + config.spec_min_width]

    if self.apply_spec_aug:
        data = spec_augment(data)

    data = np.expand_dims(data, axis=2)
    data = self.transforms(data)

    label = self.labels[idx]
    label = torch.from_numpy(label).float()
    return data, label

```

3.3.4 Test + TTA

TTA helped to fight instability, it worked well for us.

```

[31]: class FATDatasetTest(Dataset):
    def __init__(self, mels, transforms, y=None, nb_tta=5):
        super().__init__()
        self.mels = mels
        self.transforms = transforms
        self.tta = nb_tta
        if y is None:
            self.y = np.zeros(len(self.mels))
        else:
            self.y = y

    def __len__(self):
        return len(self.mels) * self.tta

    def __getitem__(self, idx):
        new_idx = idx % len(self.mels)

        data = self.mels[new_idx].astype(np.float32) #/255.0
        base_dim, time_dim = data.shape

```

```

        crop = random.randint(0, max(time_dim - config.spec_min_width, 0))
        data = data[0:base_dim, crop:crop + config.spec_min_width]
        data = np.expand_dims(data, axis=2)
        data = self.transforms(data)

        return data, self.y[new_idx]


```

```

[32]: dataset = FATDatasetTrain(X_train[:10], transforms_dict['viz'], y_train[:10],
    ↪ apply_spec_aug=True)
print('Input shape : ', dataset[0][0].squeeze().shape)
plt.matshow(dataset[0][0].squeeze())
plt.show()

```

Input shape : (64, 256)



4 Modeling

4.1 Tools

4.1.1 Adaptive Pooling

```

[33]: def adaptive_concat_pool2d(x, sz=(1,1)):
        out1 = F.adaptive_avg_pool2d(x, sz).view(x.size(0), -1)
        out2 = F.adaptive_max_pool2d(x, sz).view(x.size(0), -1)
        return torch.cat([out1, out2], 1)

```

4.1.2 ConvBlock

Adapted from <https://www.kaggle.com/mhiro2/simple-2d-cnn-classifier-with-pytorch>

```

[34]: class ConvBlock(nn.Module):
        def __init__(self, in_channels, out_channels, kernel_size=3, pool=True):

```

```

        super().__init__()

        padding = kernel_size // 2
        self.pool = pool

        self.conv1 = nn.Sequential(
            nn.Conv2d(in_channels, out_channels, kernel_size=kernel_size,
→stride=1, padding=padding),
            nn.BatchNorm2d(out_channels),
            nn.ReLU(),
        )
        self.conv2 = nn.Sequential(
            nn.Conv2d(out_channels + in_channels, out_channels,
→kernel_size=kernel_size, stride=1, padding=padding),
            nn.BatchNorm2d(out_channels),
            nn.ReLU(),
        )

        self._init_weights()

    def _init_weights(self):
        for m in self.modules():
            if isinstance(m, nn.Conv2d):
                nn.init.kaiming_normal_(m.weight)
                if m.bias is not None:
                    nn.init.zeros_(m.bias)
            elif isinstance(m, nn.BatchNorm2d):
                nn.init.constant_(m.weight, 1)
                nn.init.zeros_(m.bias)

    def forward(self, x): # x.shape = [batch_size, in_channels, a, b]
        x1 = self.conv1(x)
        x = self.conv2(torch.cat([x, x1], 1))
        if(self.pool): x = F.avg_pool2d(x, 2)
        return x # x.shape = [batch_size, out_channels, a//2, b//2]

```

4.2 Models

```

[35]: class Classifier_M1(nn.Module):
    def __init__(self, num_classes=num_classes):
        super().__init__()
        self.conv = nn.Sequential(
            ConvBlock(in_channels=1, out_channels=64),
            ConvBlock(in_channels=64, out_channels=128),
            ConvBlock(in_channels=128, out_channels=256),
            ConvBlock(in_channels=256, out_channels=512),

```

```

    )

    self.fc = nn.Sequential(
        nn.BatchNorm1d(512*2),
        nn.Linear(512*2, 128),
        nn.PReLU(),
        nn.BatchNorm1d(128),
        nn.Linear(128, num_classes),
    )

    def forward(self, x): # batch_size, 3, a, b
        x = self.conv(x) # batch_size, 512, a//16, b//16
        x = self.fc(adaptive_concat_pool2d(x))
        return x

```

```

[36]: class Classifier_M2(nn.Module):
    def __init__(self, num_classes=num_classes):
        super().__init__()
        self.conv1 = ConvBlock(1,64)
        self.conv2 = ConvBlock(64,128)
        self.conv3 = ConvBlock(128,256)
        self.conv4 = ConvBlock(256,512,pool=False)

        self.fc = nn.Sequential(
            nn.BatchNorm1d(1792),
            nn.Linear(1792, 256),
            nn.PReLU(),
            nn.BatchNorm1d(256),
            nn.Linear(256, num_classes),
        )

    def forward(self, x): # batch_size, 3, a, b
        x1 = self.conv1(x)
        x2 = self.conv2(x1)
        x3 = self.conv3(x2)
        x4 = self.conv4(x3)
        #pyramid pooling
        x = torch.cat([adaptive_concat_pool2d(x2), adaptive_concat_pool2d(x3),
                       adaptive_concat_pool2d(x4)], 1)
        x = self.fc(x)
        return x

```

```

[37]: class Classifier_M3(nn.Module):
    def __init__(self, num_classes=num_classes):
        super().__init__()
        self.conv1 = ConvBlock(1,64)
        self.conv2 = ConvBlock(64,128)

```

```

self.conv3 = ConvBlock(128,256)
self.conv4 = ConvBlock(256,512)
self.conv5 = ConvBlock(512,1024,pool=False)

self.fc = nn.Sequential(
    nn.BatchNorm1d(3840),
    nn.Linear(3840, 256),
    nn.PReLU(),
    nn.BatchNorm1d(256),
    nn.Linear(256, num_classes),
)

def forward(self, x): # batch_size, 3, a, b
    x1 = self.conv1(x)
    x2 = self.conv2(x1)
    x3 = self.conv3(x2)
    x4 = self.conv4(x3)
    x5 = self.conv5(x4)
    #pyramid pooling
    x = torch.cat([adaptive_concat_pool2d(x2), adaptive_concat_pool2d(x3),
                  adaptive_concat_pool2d(x4),adaptive_concat_pool2d(x5)],  

↪1)

    x = self.fc(x)
    return x

```

5 Tools

5.1 LwLRAP

```

[38]: def _one_sample_positive_class_precisions(scores, truth):
    num_classes = scores.shape[0]
    pos_class_indices = np.flatnonzero(truth > 0)

    if not len(pos_class_indices):
        return pos_class_indices, np.zeros(0)

    retrieved_classes = np.argsort(scores)[::-1]

    class_rankings = np.zeros(num_classes, dtype=np.int)
    class_rankings[retrieved_classes] = range(num_classes)

    retrieved_class_true = np.zeros(num_classes, dtype=np.bool)
    retrieved_class_true[class_rankings[pos_class_indices]] = True

    retrieved_cumulative_hits = np.cumsum(retrieved_class_true)

```

```
precision_at_hits = (
    retrieved_cumulative_hits[class_rankings[pos_class_indices]] /
    (1 + class_rankings[pos_class_indices].astype(np.float)))
return pos_class_indices, precision_at_hits
```

```
[39]: def lwlrp(truth, scores):
    assert truth.shape == scores.shape
    num_samples, num_classes = scores.shape
    precisions_for_samples_by_classes = np.zeros((num_samples, num_classes))
    for sample_num in range(num_samples):
        pos_class_indices, precision_at_hits =
        ↪_one_sample_positive_class_precisions(scores[sample_num, :],
        ↪truth[sample_num, :])
        precisions_for_samples_by_classes[sample_num, pos_class_indices] =
        ↪precision_at_hits


    labels_per_class = np.sum(truth > 0, axis=0)
    weight_per_class = labels_per_class / float(np.sum(labels_per_class))

    per_class_lwlrp = (np.sum(precisions_for_samples_by_classes, axis=0) /
        np.maximum(1, labels_per_class))
    return per_class_lwlrp, weight_per_class
```

5.2 Mixup

Facebook's implementation is used for Mixup. With its default parameters.

```
[40]: alpha_ = 0.4
sns.distplot(np.random.beta(alpha_, alpha_, 100000))
plt.xlim(0, 1)
plt.title(f'Distribution used for Mixup (alpha = {alpha_})', size=15)
plt.show()
```



```
[41]: def mixup_data(x, y, alpha=alpha_, use_cuda=True):
    if alpha > 0:
        lam = np.random.beta(alpha, alpha)
    else:
        lam = 1

    batch_size = x.size()[0]
    if use_cuda:
        index = torch.randperm(batch_size).cuda()
    else:
        index = torch.randperm(batch_size)

    mixed_x = lam * x + (1 - lam) * x[index, :]
    y_a, y_b = y, y[index]
    return mixed_x, y_a, y_b, lam
```

```
[42]: def mixup_criterion(criterion, pred, y_a, y_b, lam):
    return lam * criterion(pred.float().cuda(), y_a.float().cuda()) + (1 - lam) *
    ↪ criterion(pred.float().cuda(), y_b.float().cuda())
```

5.3 Label Smoothing

We tried this as well, without any success. I do believe that clever use of label smoothing & noisy data can improve results

```
[43]: def smooth(y, eps=0.4):
        a = 1 - eps * (1 + 1/y.shape[1])
        b = eps / y.shape[1]
        return a * y + b
```

5.4 Learning Rate


I reimplemented a custom Scheduler, that does Cosine Annealing and warmup. We found out that warmup played a quite important role

```
[44]: class WarmupCosineAnnealingLR(CosineAnnealingLR):
        def __init__(self, optimizer, warmup_prop, T_max, eta_min=0, last_epoch=-1):
            self.warmup_prop = warmup_prop
            super(CosineAnnealingLR, self).__init__(optimizer, T_max, eta_min,
            ↪last_epoch)

        def get_lr(self):
            if self.last_epoch < T_max * warmup_prop:
                return [self.base_lrs * self.last_epoch/T_max
                        for base_lr, group in zip(self.base_lrs, self.optimizer.
            ↪param_groups)]
            else:
                return super(CosineAnnealingLR, self).step()
```

```
[45]: lr_max = 0.001
        lr_min = 1e-5
        warmup_prop = 0.1
        nb_epochs = 80
        epochs = [i for i in range(nb_epochs)]
        lrs_warmup = [lr_max * i / (warmup_prop * nb_epochs) for i in range(1,
            ↪int(warmup_prop * nb_epochs) + 1)]
        lrs_scheduler = [lr_min + 0.5 * (lr_max - lr_min) * (1 + np.cos(np.pi * i / ((1
            ↪- warmup_prop) * nb_epochs))) for i in range(int((1 - warmup_prop) *
            ↪nb_epochs))]

        plt.figure(figsize=(15, 5))
        plt.axhline(lr_min, label='lr_min', c='skyblue', linestyle='--')
        plt.axhline(lr_max, label='lr_max', c='orangered', linestyle='--')
        plt.plot(epochs, lrs_warmup + lrs_scheduler[:len(epochs)])
        plt.grid(True)
        plt.legend()
        plt.title('Learning Rate Scheduling Policy', size=15)
        plt.show()
```

5.5 Sigmoid

```
[46]: def sigmoid(x):
       return 1 / (1 + np.exp(-x))
```

5.6 Save & Load

```
[47]: def save_model_weights(model, filename, verbose=1):
       if verbose:
           print(f'Saving weights to {filename}\n')
       torch.save(model.state_dict(), filename)
```

```
[48]: def load_model_weights(model, filename, verbose=1):
       if verbose:
           print(f'Loading weights from {filename}\n')
       model.load_state_dict(torch.load(filename))
       return model
```

5.7 Predict

```
[49]: def predict(dataset, model, batch_size=64):
       model.eval()
       y = np.array([[]] * num_classes).T
       loader = torch.utils.data.DataLoader(dataset, batch_size=batch_size,
       ↪shuffle=False)
       with torch.no_grad():
           for x, target in loader:
               y = np.concatenate((y, sigmoid(model(x.cuda())).view(-1,
               ↪num_classes).detach().cpu().numpy())))
```

```
return y
```

5.8 Fit

Code is a bit ugly, you might not want to dive into it. It's PyTorch. Here are the main points :

- Weight decay (0.001)
- Gradient clipping (1)
- Binary Cross Entropy loss
- Adam is used here, Adabound was also tried
- Custom lr scheduler (see above), with a warmup proportion 0.1 usually.
- Apply mixup with a probability *mixup* (chosen at 1)
- Checkpointing is used

```
[50]: def fit(model, train_dataset, val_dataset, epochs=50, batch_size=128, nb_tta=5,
    ↪mixup=False, warmup_prop=0.1,
    ↪verbose=1, cp=False, model_name='model', lr=0.001):
    avg_val_loss = 1000
    best_score = 0
    clip_value = 1.0

    model.cuda()

    opt_params = [
        {'params': [p for n, p in list(model.named_parameters()) if not any(nd
    ↪in n for nd in ['bias', 'LayerNorm.bias', 'LayerNorm.weight'])],
    ↪'weight_decay': 0.01},
        {'params': [p for n, p in list(model.named_parameters()) if any(nd in n
    ↪for nd in ['bias', 'LayerNorm.bias', 'LayerNorm.weight'])], 'weight_decay':
    ↪0.0}
    ]

    for p in model.parameters():
        p.register_hook(lambda grad: torch.clamp(grad, -clip_value, clip_value))

    optimizer = torch.optim.Adam(opt_params, lr=lr)
    scheduler = CosineAnnealingLR(optimizer, T_max=epochs, eta_min=1e-7)
    loss_fn = nn.BCEWithLogitsLoss(reduction='sum')

    train_loader = torch.utils.data.DataLoader(train_dataset,
    ↪batch_size=batch_size, shuffle=True)
    val_loader = torch.utils.data.DataLoader(val_dataset,
    ↪batch_size=batch_size, shuffle=False)

    if warmup_prop >= 1:
```

```

warmup_prop = warmup_prop / epochs

for epoch in range(epochs):
    model.train()
    avg_loss = 0
    start_time = time.time()

    if epoch < epochs * warmup_prop:
        lr = 5e-5 + 0.001 * epoch / (epochs * warmup_prop)
        for param_group in optimizer.param_groups:
            param_group['lr'] = lr
    else:
        scheduler.step()
        lr = np.mean([param_group['lr'] for param_group in optimizer.
→param_groups])

    for x, y_batch in train_loader:
        optimizer.zero_grad()

        dice = np.random.random() < mixup
        if dice:
            x, y_a, y_b, lam = mixup_data(x, y_batch.float())
            x, y_a, y_b = map(Variable, (x, y_a, y_b))

        y_pred = model(x.cuda()).view(-1, num_classes)
        if dice:
            loss = mixup_criterion(loss_fn, y_pred, y_a, y_b, lam)
        else:
            loss = loss_fn(y_pred.float(), y_batch.float().cuda())

        loss.backward()
        avg_loss += loss.item() / len(train_dataset)
        optimizer.step()

    model.eval()
    avg_val_loss = 0.
    pred_val = np.array([[]]*num_classes).T

    for x, y_batch in val_loader:
        y_pred = model(x.cuda()).view(-1, num_classes).detach()
        avg_val_loss += loss_fn(y_pred.float(), y_batch.float().cuda()).
→item() / len(val_dataset)
        pred_val = np.concatenate((pred_val, sigmoid(y_pred.cpu().numpy())))

    pred_val = np.mean(pred_val.reshape((nb_tta, -1, num_classes,)), axis=0)
    score_class, weight = lwrap(val_dataset.y, pred_val)

```

```

score = (score_class * weight).sum()

elapsed_time = time.time() - start_time

if cp:
    if score > best_score:
        save_model_weights(model, f"{model_name}.pt", verbose=0)
        best_score = score

    if (epoch + 1) % verbose == 0:
        elapsed_time = elapsed_time * verbose
        print(f'Epoch {epoch+1}/{epochs}      lr={lr:.1e}      lwlap={score:.
↪5f}      ', end='')
        print(f'loss={avg_loss:.2f}      val_loss={avg_val_loss:.2f}      ↪
↪t={elapsed_time:.0f}s')

return best_score

```

5.9 Refit

- Same thing as the previous function, but:
- We change the scheduler to reduce_lr_on_plateau, starting from 0.001
- Used with curated data only

```

[51]: def refit(model, train_dataset, val_dataset, epochs=10, batch_size=128, ↪
↪nb_tta=5, mixup=False,
        verbose=1, cp=False, model_name='model', best_score=0, lr=0.001):
    avg_val_loss = 1000
    clip_value = 1.0

    model.cuda()

    opt_params = [
        {'params': [p for n, p in list(model.named_parameters()) if not any(nd ↪
↪in n for nd in ['bias', 'LayerNorm.bias', 'LayerNorm.weight'])], ↪
↪'weight_decay': 0.01},
        {'params': [p for n, p in list(model.named_parameters()) if any(nd in n ↪
↪for nd in ['bias', 'LayerNorm.bias', 'LayerNorm.weight'])], 'weight_decay': ↪
↪0.0}
    ]

    for p in model.parameters():
        p.register_hook(lambda grad: torch.clamp(grad, -clip_value, clip_value))

    optimizer = torch.optim.Adam(opt_params, lr)

```

```

    scheduler = ReduceLROnPlateau(optimizer, mode='min', patience=3, factor=0.
↪1, verbose=0)
    loss_fn = nn.BCEWithLogitsLoss(reduction='sum')

    train_loader = torch.utils.data.DataLoader(train_dataset, ↪
↪batch_size=batch_size, shuffle=True)
    val_loader = torch.utils.data.DataLoader(val_dataset, ↪
↪batch_size=batch_size, shuffle=False)

    for epoch in range(epochs):
        model.train()
        avg_loss = 0
        start_time = time.time()

        scheduler.step(avg_val_loss)
        lr = np.mean([param_group['lr'] for param_group in optimizer.
↪param_groups])

        for x, y_batch in train_loader:
            optimizer.zero_grad()

            dice = np.random.random() < mixup
            if dice:
                x, y_a, y_b, lam = mixup_data(x, y_batch.float())
                x, y_a, y_b = map(Variable, (x, y_a, y_b))

            y_pred = model(x.cuda()).view(-1, num_classes)
            if dice:
                loss = mixup_criterion(loss_fn, y_pred, y_a, y_b, lam)
            else:
                loss = loss_fn(y_pred.float(), y_batch.float().cuda())

            loss.backward()
            avg_loss += loss.item() / len(train_dataset)
            optimizer.step()

        model.eval()
        avg_val_loss = 0.
        pred_val = np.array([[]]*num_classes).T

        for x, y_batch in val_loader:
            y_pred = model(x.cuda()).view(-1, num_classes).detach()
            avg_val_loss += loss_fn(y_pred.float(), y_batch.float().cuda()).
↪item() / len(val_dataset)
            pred_val = np.concatenate((pred_val, sigmoid(y_pred.cpu().numpy())))

```

```

pred_val = np.mean(pred_val.reshape((nb_tta, -1, num_classes,)), axis=0)
score_class, weight = lwlap(val_dataset.y, pred_val)
score = (score_class * weight).sum()

elapsed_time = time.time() - start_time

if cp:
    if score > best_score:
        save_model_weights(model, f"{model_name}.pt", verbose=0)
        best_score = score

if (epoch + 1) % verbose == 0:
    elapsed_time = elapsed_time * verbose
    print(f'Epoch {epoch+1}/{epochs}      lr={lr:.1e}      lwlap={score:.
↪5f}      ', end='')
    print(f'loss={avg_loss:.2f}      val_loss={avg_val_loss:.2f}      □
↪t={elapsed_time:.0f}s')

    if lr <= 1e-6:
        break
return best_score

```

5.10 *k*-folds

- 5 folds Stratified, using one of the labels if there are more. Only curated data is kept.
- Fit for 65 (approx.) epochs on curated + selected noisy (optional)
- Refit twice for a few epochs on curated only

```

[52]: def k_fold(model_class, X, y, X_test, transform_dic, X_noisy=[], y_noisy=np.
↪array([]),
        k=5, batch_size=32, epochs=5, seed=2019, nb_tta=5, tta_eval=25,
↪mixup=False,
        verbose=1, save=True, cp=False, warmup_prop=0.1, model_name="model",
↪pretrained_path=''):

    splits = list(StratifiedKFold(n_splits=k, shuffle=True, random_state=seed).
↪split(X, np.argmax(y, axis=1)))

    pred_test = np.zeros((len(X_test), num_classes))
    pred_oof = np.zeros((y.shape[0], num_classes))

    for i, (train_idx, val_idx) in enumerate(splits):
        print(f"----- Fold {i+1} -----")
        seed_everything(seed + i)
        start_time = time.time()

```

```

model = model_class(num_classes=num_classes)

if len(pretrained_path):
    load_model_weights(model, pretrained_path)

train_dataset = FATDatasetTrain([X[i] for i in train_idx] + X_noisy,
↳transforms=transform_dic['train'],
                                y=np.concatenate((y[train_idx],
↳y_noisy))),apply_spec_aug=True)
val_dataset = FATDatasetTest([X[i] for i in val_idx],
↳transforms=transform_dic['test'], y=y[val_idx], nb_tta=nb_tta)
test_dataset = FATDatasetTest(X_test, transforms=transform_dic['test'],
↳nb_tta=tta_eval)

print('\n - Fitting \n')

best_score = fit(model, train_dataset, val_dataset, epochs=epochs,
↳batch_size=batch_size, nb_tta=nb_tta, mixup=mixup,
                    warmup_prop=warmup_prop, verbose=verbose, cp=cp,
↳model_name=f'{model_name}_{i+1}', lr=1e-3)

print('\n - Re-fitting with curated data only (1/2)\n')

retrain_dataset = FATDatasetTrain([X[i] for i in train_idx],
↳transforms=transform_dic['train'], y=y[train_idx])

best_score = refit(model, retrain_dataset, val_dataset, epochs=epochs,
↳batch_size=batch_size, nb_tta=nb_tta, mixup=mixup,
                    verbose=verbose, cp=cp,
↳model_name=f'{model_name}_{i+1}', best_score=best_score, lr=1e-4)

if cp:
    load_model_weights(model, f"{model_name}_{i+1}.pt", verbose=1)
elif save:
    save_model_weights(model, f"{model_name}_{i+1}.pt", verbose=1)

print('\n - Re-fitting with curated data only (2/2)\n')

best_score = refit(model, retrain_dataset, val_dataset, epochs=epochs,
↳batch_size=batch_size, nb_tta=nb_tta, mixup=mixup,
                    verbose=verbose, cp=cp,
↳model_name=f'{model_name}_{i+1}', best_score=best_score, lr=1e-4)

if cp:
    load_model_weights(model, f"{model_name}_{i+1}.pt", verbose=1)
elif save:

```

```

        save_model_weights(model, f"{model_name}_{i+1}.pt", verbose=1)

    print(f'\n - Predicting with {tta_eval} TTA \n')

    val_dataset = FATDatasetTest([X[i] for i in val_idx],
    ↪ transforms=transform_dic['test'], y=y[val_idx], nb_tta=tta_eval)
    pred_val = np.mean(predict(val_dataset, model).reshape((tta_eval, -1,
    ↪ num_classes,)), axis=0)
    pred_oof[val_idx, :] = pred_val

    score_class, weight = lwrap(y[val_idx], pred_val)
    score = (score_class * weight).sum()

    pred_test += np.mean(predict(test_dataset, model).reshape((tta_eval,
    ↪ -1, num_classes,)), axis=0)

    print(f"\n lwrap : Scored {score :.4f} on validation data")
    print(f"\n     Done in {(time.time() - start_time) / 60 :.1f} minutes
    ↪ \n")

    return pred_test, pred_oof

```

6 Training

6.0.1 Parameters

```

[53]: model_name = "model"
      k = 5
      epochs = 65
      batch_size = 64
      nb_tta = 5
      mixup = 1
      cp = True
      warmup_prop = 0.1

```

```

[54]: model = Classifier_M3
      transforms_ = transforms_dict

```

6.0.2 Weights pretrained on noisy data

In case we want to use weights pretrained on noisy, we specify weights to load.


```
[55]: # cp_path = "../input/fat-cp/new_default_65.pt"
      cp_path = ''
```

6.0.3 Fitting

```
[56]: pred_test, pred_oof = k_fold(model, X_train, y_train, X_test, transforms_,
                                   X_noisy=X_noisy, y_noisy=y_noisy,
                                   k=k, batch_size=batch_size, epochs=epochs,
                                   nb_tta=nb_tta, cp=cp, warmup_prop=warmup_prop,
                                   mixup=mixup, seed=seed,
                                   model_name=model_name, pretrained_path=cp_path)
```

----- Fold 1 -----

- Fitting

Epoch 1/65 t=51s	lr=5.0e-05	lwlr=0.26018	loss=55.27	val_loss=55.02
Epoch 2/65 t=51s	lr=2.0e-04	lwlr=0.32222	loss=52.35	val_loss=47.70
Epoch 3/65 t=50s	lr=3.6e-04	lwlr=0.26679	loss=34.57	val_loss=20.81
Epoch 4/65 t=50s	lr=5.1e-04	lwlr=0.31297	loss=10.67	val_loss=5.44
Epoch 5/65 t=50s	lr=6.7e-04	lwlr=0.36285	loss=5.08	val_loss=5.00
Epoch 6/65 t=50s	lr=8.2e-04	lwlr=0.37723	loss=4.69	val_loss=4.69
Epoch 7/65 t=50s	lr=9.7e-04	lwlr=0.43856	loss=4.48	val_loss=4.38
Epoch 8/65 t=50s	lr=1.0e-03	lwlr=0.54052	loss=4.26	val_loss=3.80
Epoch 9/65 t=50s	lr=1.0e-03	lwlr=0.60203	loss=4.01	val_loss=3.60
Epoch 10/65 t=50s	lr=1.0e-03	lwlr=0.50895	loss=3.93	val_loss=4.25
Epoch 11/65 t=50s	lr=9.9e-04	lwlr=0.61065	loss=3.86	val_loss=3.43
Epoch 12/65 t=51s	lr=9.9e-04	lwlr=0.68352	loss=3.68	val_loss=3.00
Epoch 13/65 t=51s	lr=9.9e-04	lwlr=0.66243	loss=3.74	val_loss=3.14
Epoch 14/65 t=50s	lr=9.8e-04	lwlr=0.66513	loss=3.65	val_loss=3.08
Epoch 15/65 t=50s	lr=9.7e-04	lwlr=0.69114	loss=3.68	val_loss=2.97

Epoch 16/65 t=51s	lr=9.6e-04	lwlrwrap=0.67038	loss=3.59	val_loss=3.05
Epoch 17/65 t=50s	lr=9.5e-04	lwlrwrap=0.74167	loss=3.55	val_loss=2.65
Epoch 18/65 t=51s	lr=9.4e-04	lwlrwrap=0.73690	loss=3.45	val_loss=2.61
Epoch 19/65 t=51s	lr=9.3e-04	lwlrwrap=0.73952	loss=3.38	val_loss=2.56
Epoch 20/65 t=50s	lr=9.2e-04	lwlrwrap=0.76187	loss=3.31	val_loss=2.55
Epoch 21/65 t=51s	lr=9.0e-04	lwlrwrap=0.73961	loss=3.37	val_loss=2.65
Epoch 22/65 t=50s	lr=8.9e-04	lwlrwrap=0.75324	loss=3.26	val_loss=2.45
Epoch 23/65 t=51s	lr=8.7e-04	lwlrwrap=0.75069	loss=3.30	val_loss=2.48
Epoch 24/65 t=50s	lr=8.6e-04	lwlrwrap=0.69725	loss=3.04	val_loss=2.94
Epoch 25/65 t=51s	lr=8.4e-04	lwlrwrap=0.76349	loss=3.24	val_loss=2.56
Epoch 26/65 t=51s	lr=8.2e-04	lwlrwrap=0.78099	loss=3.13	val_loss=2.32
Epoch 27/65 t=51s	lr=8.0e-04	lwlrwrap=0.76385	loss=3.30	val_loss=2.36
Epoch 28/65 t=51s	lr=7.8e-04	lwlrwrap=0.77796	loss=3.14	val_loss=2.37
Epoch 29/65 t=50s	lr=7.6e-04	lwlrwrap=0.78999	loss=3.08	val_loss=2.23
Epoch 30/65 t=50s	lr=7.4e-04	lwlrwrap=0.77614	loss=3.11	val_loss=2.27
Epoch 31/65 t=50s	lr=7.2e-04	lwlrwrap=0.79385	loss=2.93	val_loss=2.14
Epoch 32/65 t=51s	lr=7.0e-04	lwlrwrap=0.80254	loss=2.91	val_loss=2.08
Epoch 33/65 t=51s	lr=6.8e-04	lwlrwrap=0.82031	loss=2.95	val_loss=2.03
Epoch 34/65 t=51s	lr=6.5e-04	lwlrwrap=0.81929	loss=2.93	val_loss=2.02
Epoch 35/65 t=50s	lr=6.3e-04	lwlrwrap=0.82256	loss=2.82	val_loss=2.01
Epoch 36/65 t=50s	lr=6.1e-04	lwlrwrap=0.82071	loss=2.86	val_loss=1.95
Epoch 37/65 t=51s	lr=5.8e-04	lwlrwrap=0.82696	loss=2.78	val_loss=1.86
Epoch 38/65 t=51s	lr=5.6e-04	lwlrwrap=0.83216	loss=2.59	val_loss=1.93
Epoch 39/65 t=50s	lr=5.4e-04	lwlrwrap=0.83329	loss=2.78	val_loss=1.89

Epoch 40/65 t=51s	lr=5.1e-04	lwlrwrap=0.82690	loss=2.74	val_loss=1.91
Epoch 41/65 t=50s	lr=4.9e-04	lwlrwrap=0.82497	loss=2.55	val_loss=1.96
Epoch 42/65 t=50s	lr=4.6e-04	lwlrwrap=0.81964	loss=2.60	val_loss=1.95
Epoch 43/65 t=50s	lr=4.4e-04	lwlrwrap=0.82881	loss=2.66	val_loss=1.94
Epoch 44/65 t=51s	lr=4.2e-04	lwlrwrap=0.84151	loss=2.56	val_loss=1.82
Epoch 45/65 t=50s	lr=3.9e-04	lwlrwrap=0.83742	loss=2.42	val_loss=1.82
Epoch 46/65 t=50s	lr=3.7e-04	lwlrwrap=0.85139	loss=2.27	val_loss=1.72
Epoch 47/65 t=50s	lr=3.5e-04	lwlrwrap=0.83361	loss=2.42	val_loss=1.89
Epoch 48/65 t=50s	lr=3.2e-04	lwlrwrap=0.84908	loss=2.34	val_loss=1.76
Epoch 49/65 t=50s	lr=3.0e-04	lwlrwrap=0.85275	loss=2.42	val_loss=1.72
Epoch 50/65 t=51s	lr=2.8e-04	lwlrwrap=0.84569	loss=2.29	val_loss=1.75
Epoch 51/65 t=50s	lr=2.6e-04	lwlrwrap=0.85379	loss=2.30	val_loss=1.72
Epoch 52/65 t=50s	lr=2.4e-04	lwlrwrap=0.85505	loss=2.17	val_loss=1.66
Epoch 53/65 t=50s	lr=2.2e-04	lwlrwrap=0.85766	loss=2.19	val_loss=1.72
Epoch 54/65 t=50s	lr=2.0e-04	lwlrwrap=0.85389	loss=2.11	val_loss=1.71
Epoch 55/65 t=50s	lr=1.8e-04	lwlrwrap=0.86408	loss=2.08	val_loss=1.64
Epoch 56/65 t=50s	lr=1.6e-04	lwlrwrap=0.85847	loss=2.11	val_loss=1.68
Epoch 57/65 t=50s	lr=1.4e-04	lwlrwrap=0.86186	loss=2.16	val_loss=1.64
Epoch 58/65 t=50s	lr=1.3e-04	lwlrwrap=0.85690	loss=2.12	val_loss=1.71
Epoch 59/65 t=50s	lr=1.1e-04	lwlrwrap=0.86328	loss=2.12	val_loss=1.62
Epoch 60/65 t=50s	lr=9.6e-05	lwlrwrap=0.86328	loss=2.15	val_loss=1.64
Epoch 61/65 t=50s	lr=8.2e-05	lwlrwrap=0.85629	loss=2.10	val_loss=1.72
Epoch 62/65 t=50s	lr=6.9e-05	lwlrwrap=0.86327	loss=2.09	val_loss=1.63
Epoch 63/65 t=50s	lr=5.7e-05	lwlrwrap=0.86837	loss=1.95	val_loss=1.61

Epoch 64/65 t=50s	lr=4.7e-05	lwlrwrap=0.86044	loss=1.99	val_loss=1.64
Epoch 65/65 t=50s	lr=3.7e-05	lwlrwrap=0.86022	loss=2.02	val_loss=1.63

- Re-fitting with curated data only (1/2)

Epoch 1/65 t=34s	lr=1.0e-04	lwlrwrap=0.86119	loss=1.81	val_loss=1.66
Epoch 2/65 t=34s	lr=1.0e-04	lwlrwrap=0.87080	loss=1.71	val_loss=1.63
Epoch 3/65 t=34s	lr=1.0e-04	lwlrwrap=0.86493	loss=1.70	val_loss=1.60
Epoch 4/65 t=34s	lr=1.0e-04	lwlrwrap=0.87125	loss=1.57	val_loss=1.61
Epoch 5/65 t=34s	lr=1.0e-04	lwlrwrap=0.86143	loss=1.68	val_loss=1.66
Epoch 6/65 t=34s	lr=1.0e-04	lwlrwrap=0.86141	loss=1.56	val_loss=1.67
Epoch 7/65 t=34s	lr=1.0e-04	lwlrwrap=0.86181	loss=1.64	val_loss=1.68
Epoch 8/65 t=34s	lr=1.0e-05	lwlrwrap=0.86455	loss=1.71	val_loss=1.66
Epoch 9/65 t=34s	lr=1.0e-05	lwlrwrap=0.86781	loss=1.69	val_loss=1.65
Epoch 10/65 t=34s	lr=1.0e-05	lwlrwrap=0.86688	loss=1.64	val_loss=1.66
Epoch 11/65 t=34s	lr=1.0e-05	lwlrwrap=0.86353	loss=1.54	val_loss=1.67
Epoch 12/65 t=34s	lr=1.0e-06	lwlrwrap=0.86722	loss=1.54	val_loss=1.65
Epoch 13/65 t=34s	lr=1.0e-06	lwlrwrap=0.86256	loss=1.68	val_loss=1.66
Epoch 14/65 t=34s	lr=1.0e-06	lwlrwrap=0.86433	loss=1.69	val_loss=1.64
Epoch 15/65 t=34s	lr=1.0e-06	lwlrwrap=0.87051	loss=1.69	val_loss=1.60
Epoch 16/65 t=34s	lr=1.0e-07	lwlrwrap=0.86571	loss=1.89	val_loss=1.64

Loading weights from model_1.pt

- Re-fitting with curated data only (2/2)

Epoch 1/65 t=34s	lr=1.0e-04	lwlrwrap=0.86580	loss=1.79	val_loss=1.63
Epoch 2/65 t=34s	lr=1.0e-04	lwlrwrap=0.85720	loss=1.90	val_loss=1.68

Epoch 3/65 t=34s	lr=1.0e-04	lwlap=0.86975	loss=1.65	val_loss=1.63
Epoch 4/65 t=34s	lr=1.0e-04	lwlap=0.86020	loss=1.85	val_loss=1.69
Epoch 5/65 t=34s	lr=1.0e-04	lwlap=0.85750	loss=1.50	val_loss=1.66
Epoch 6/65 t=34s	lr=1.0e-04	lwlap=0.86362	loss=1.67	val_loss=1.64
Epoch 7/65 t=34s	lr=1.0e-04	lwlap=0.87108	loss=1.37	val_loss=1.58
Epoch 8/65 t=34s	lr=1.0e-04	lwlap=0.87157	loss=1.63	val_loss=1.63
Epoch 9/65 t=34s	lr=1.0e-04	lwlap=0.86306	loss=1.86	val_loss=1.70
Epoch 10/65 t=34s	lr=1.0e-04	lwlap=0.86813	loss=1.57	val_loss=1.68
Epoch 11/65 t=34s	lr=1.0e-04	lwlap=0.85985	loss=1.70	val_loss=1.66
Epoch 12/65 t=34s	lr=1.0e-05	lwlap=0.86882	loss=1.59	val_loss=1.63
Epoch 13/65 t=34s	lr=1.0e-05	lwlap=0.86365	loss=1.61	val_loss=1.64
Epoch 14/65 t=34s	lr=1.0e-05	lwlap=0.86310	loss=1.55	val_loss=1.66
Epoch 15/65 t=34s	lr=1.0e-05	lwlap=0.86480	loss=1.73	val_loss=1.66
Epoch 16/65 t=34s	lr=1.0e-06	lwlap=0.85833	loss=1.82	val_loss=1.69
Epoch 17/65 t=34s	lr=1.0e-06	lwlap=0.87328	loss=1.50	val_loss=1.62
Epoch 18/65 t=34s	lr=1.0e-06	lwlap=0.86849	loss=1.57	val_loss=1.64
Epoch 19/65 t=34s	lr=1.0e-06	lwlap=0.86565	loss=1.51	val_loss=1.66
Epoch 20/65 t=34s	lr=1.0e-07	lwlap=0.86396	loss=1.62	val_loss=1.67

Loading weights from model_1.pt

- Predicting with 25 TTA

lwlap : Scored 0.8720 on validation data

Done in 79.2 minutes

----- Fold 2 -----

- Fitting

Epoch 1/65 t=50s	lr=5.0e-05	lwlrwrap=0.25902	loss=55.37	val_loss=55.17
Epoch 2/65 t=50s	lr=2.0e-04	lwlrwrap=0.32307	loss=52.38	val_loss=48.82
Epoch 3/65 t=50s	lr=3.6e-04	lwlrwrap=0.25209	loss=34.42	val_loss=18.35
Epoch 4/65 t=50s	lr=5.1e-04	lwlrwrap=0.33309	loss=10.60	val_loss=6.06
Epoch 5/65 t=50s	lr=6.7e-04	lwlrwrap=0.40691	loss=5.06	val_loss=4.67
Epoch 6/65 t=50s	lr=8.2e-04	lwlrwrap=0.38906	loss=4.66	val_loss=4.56
Epoch 7/65 t=50s	lr=9.7e-04	lwlrwrap=0.48700	loss=4.38	val_loss=4.36
Epoch 8/65 t=50s	lr=1.0e-03	lwlrwrap=0.52742	loss=4.40	val_loss=3.99
Epoch 9/65 t=50s	lr=1.0e-03	lwlrwrap=0.62530	loss=4.10	val_loss=3.44
Epoch 10/65 t=50s	lr=1.0e-03	lwlrwrap=0.64406	loss=4.03	val_loss=3.24
Epoch 11/65 t=50s	lr=9.9e-04	lwlrwrap=0.65728	loss=3.90	val_loss=3.18
Epoch 12/65 t=50s	lr=9.9e-04	lwlrwrap=0.64846	loss=3.93	val_loss=3.13
Epoch 13/65 t=50s	lr=9.9e-04	lwlrwrap=0.68736	loss=3.70	val_loss=2.98
Epoch 14/65 t=50s	lr=9.8e-04	lwlrwrap=0.66574	loss=3.65	val_loss=3.13
Epoch 15/65 t=50s	lr=9.7e-04	lwlrwrap=0.71426	loss=3.61	val_loss=2.76
Epoch 16/65 t=50s	lr=9.6e-04	lwlrwrap=0.68025	loss=3.63	val_loss=2.97
Epoch 17/65 t=50s	lr=9.5e-04	lwlrwrap=0.71947	loss=3.52	val_loss=2.78
Epoch 18/65 t=50s	lr=9.4e-04	lwlrwrap=0.74868	loss=3.42	val_loss=2.53
Epoch 19/65 t=50s	lr=9.3e-04	lwlrwrap=0.71231	loss=3.41	val_loss=2.79
Epoch 20/65 t=50s	lr=9.2e-04	lwlrwrap=0.73408	loss=3.55	val_loss=2.73
Epoch 21/65 t=50s	lr=9.0e-04	lwlrwrap=0.73509	loss=3.31	val_loss=2.62
Epoch 22/65 t=50s	lr=8.9e-04	lwlrwrap=0.75465	loss=3.25	val_loss=2.52
Epoch 23/65 t=50s	lr=8.7e-04	lwlrwrap=0.74000	loss=3.39	val_loss=2.49

Epoch 24/65 t=50s	lr=8.6e-04	lwlrwrap=0.75720	loss=3.28	val_loss=2.49
Epoch 25/65 t=50s	lr=8.4e-04	lwlrwrap=0.80292	loss=3.23	val_loss=2.22
Epoch 26/65 t=50s	lr=8.2e-04	lwlrwrap=0.77969	loss=3.08	val_loss=2.29
Epoch 27/65 t=50s	lr=8.0e-04	lwlrwrap=0.80895	loss=3.03	val_loss=2.17
Epoch 28/65 t=50s	lr=7.8e-04	lwlrwrap=0.76786	loss=3.14	val_loss=2.39
Epoch 29/65 t=50s	lr=7.6e-04	lwlrwrap=0.79010	loss=3.04	val_loss=2.28
Epoch 30/65 t=50s	lr=7.4e-04	lwlrwrap=0.79507	loss=3.06	val_loss=2.26
Epoch 31/65 t=50s	lr=7.2e-04	lwlrwrap=0.81043	loss=2.89	val_loss=2.05
Epoch 32/65 t=50s	lr=7.0e-04	lwlrwrap=0.79879	loss=3.04	val_loss=2.18
Epoch 33/65 t=50s	lr=6.8e-04	lwlrwrap=0.82446	loss=2.90	val_loss=1.98
Epoch 34/65 t=50s	lr=6.5e-04	lwlrwrap=0.81679	loss=2.73	val_loss=1.97
Epoch 35/65 t=50s	lr=6.3e-04	lwlrwrap=0.81098	loss=2.80	val_loss=2.14
Epoch 36/65 t=50s	lr=6.1e-04	lwlrwrap=0.83682	loss=2.79	val_loss=1.92
Epoch 37/65 t=50s	lr=5.8e-04	lwlrwrap=0.83089	loss=2.77	val_loss=1.87
Epoch 38/65 t=50s	lr=5.6e-04	lwlrwrap=0.82411	loss=2.64	val_loss=1.92
Epoch 39/65 t=50s	lr=5.4e-04	lwlrwrap=0.82888	loss=2.62	val_loss=1.90
Epoch 40/65 t=50s	lr=5.1e-04	lwlrwrap=0.83106	loss=2.65	val_loss=1.89
Epoch 41/65 t=50s	lr=4.9e-04	lwlrwrap=0.82978	loss=2.68	val_loss=1.87
Epoch 42/65 t=50s	lr=4.6e-04	lwlrwrap=0.83332	loss=2.75	val_loss=1.85
Epoch 43/65 t=50s	lr=4.4e-04	lwlrwrap=0.82566	loss=2.55	val_loss=1.94
Epoch 44/65 t=50s	lr=4.2e-04	lwlrwrap=0.85103	loss=2.45	val_loss=1.74
Epoch 45/65 t=50s	lr=3.9e-04	lwlrwrap=0.85199	loss=2.39	val_loss=1.71
Epoch 46/65 t=50s	lr=3.7e-04	lwlrwrap=0.83055	loss=2.42	val_loss=1.83
Epoch 47/65 t=50s	lr=3.5e-04	lwlrwrap=0.84993	loss=2.41	val_loss=1.72

Epoch 48/65 t=50s	lr=3.2e-04	lwlrwrap=0.85309	loss=2.16	val_loss=1.68
Epoch 49/65 t=50s	lr=3.0e-04	lwlrwrap=0.83150	loss=2.35	val_loss=1.90
Epoch 50/65 t=50s	lr=2.8e-04	lwlrwrap=0.84976	loss=2.37	val_loss=1.74
Epoch 51/65 t=50s	lr=2.6e-04	lwlrwrap=0.85226	loss=2.23	val_loss=1.69
Epoch 52/65 t=50s	lr=2.4e-04	lwlrwrap=0.85748	loss=2.32	val_loss=1.64
Epoch 53/65 t=50s	lr=2.2e-04	lwlrwrap=0.86178	loss=2.19	val_loss=1.63
Epoch 54/65 t=50s	lr=2.0e-04	lwlrwrap=0.85680	loss=2.24	val_loss=1.66
Epoch 55/65 t=50s	lr=1.8e-04	lwlrwrap=0.86760	loss=2.22	val_loss=1.58
Epoch 56/65 t=50s	lr=1.6e-04	lwlrwrap=0.86286	loss=2.23	val_loss=1.59
Epoch 57/65 t=50s	lr=1.4e-04	lwlrwrap=0.86606	loss=1.92	val_loss=1.60
Epoch 58/65 t=50s	lr=1.3e-04	lwlrwrap=0.86971	loss=2.10	val_loss=1.57
Epoch 59/65 t=50s	lr=1.1e-04	lwlrwrap=0.86235	loss=2.11	val_loss=1.60
Epoch 60/65 t=50s	lr=9.6e-05	lwlrwrap=0.86598	loss=2.01	val_loss=1.57
Epoch 61/65 t=50s	lr=8.2e-05	lwlrwrap=0.86078	loss=1.91	val_loss=1.59
Epoch 62/65 t=50s	lr=6.9e-05	lwlrwrap=0.86622	loss=2.07	val_loss=1.59
Epoch 63/65 t=50s	lr=5.7e-05	lwlrwrap=0.86156	loss=2.20	val_loss=1.58
Epoch 64/65 t=50s	lr=4.7e-05	lwlrwrap=0.86164	loss=2.06	val_loss=1.59
Epoch 65/65 t=50s	lr=3.7e-05	lwlrwrap=0.87069	loss=2.04	val_loss=1.55

- Re-fitting with curated data only (1/2)

Epoch 1/65 t=34s	lr=1.0e-04	lwlrwrap=0.86663	loss=1.89	val_loss=1.66
Epoch 2/65 t=34s	lr=1.0e-04	lwlrwrap=0.86828	loss=1.80	val_loss=1.63
Epoch 3/65 t=34s	lr=1.0e-04	lwlrwrap=0.86489	loss=1.68	val_loss=1.62
Epoch 4/65 t=34s	lr=1.0e-04	lwlrwrap=0.86348	loss=1.61	val_loss=1.62
Epoch 5/65	lr=1.0e-04	lwlrwrap=0.86184	loss=1.66	val_loss=1.66

t=34s					
Epoch 6/65	lr=1.0e-04	lwlraps=0.86053	loss=1.73	val_loss=1.63	
t=34s					
Epoch 7/65	lr=1.0e-04	lwlraps=0.86825	loss=1.82	val_loss=1.64	
t=34s					
Epoch 8/65	lr=1.0e-04	lwlraps=0.86474	loss=1.60	val_loss=1.58	
t=34s					
Epoch 9/65	lr=1.0e-04	lwlraps=0.85932	loss=2.03	val_loss=1.69	
t=34s					
Epoch 10/65	lr=1.0e-04	lwlraps=0.86314	loss=1.70	val_loss=1.66	
t=34s					
Epoch 11/65	lr=1.0e-04	lwlraps=0.86507	loss=1.48	val_loss=1.63	
t=34s					
Epoch 12/65	lr=1.0e-04	lwlraps=0.85898	loss=1.68	val_loss=1.60	
t=34s					
Epoch 13/65	lr=1.0e-05	lwlraps=0.86231	loss=1.74	val_loss=1.63	
t=34s					
Epoch 14/65	lr=1.0e-05	lwlraps=0.86461	loss=1.56	val_loss=1.60	
t=34s					
Epoch 15/65	lr=1.0e-05	lwlraps=0.86098	loss=1.62	val_loss=1.64	
t=34s					
Epoch 16/65	lr=1.0e-05	lwlraps=0.86423	loss=1.71	val_loss=1.62	
t=34s					
Epoch 17/65	lr=1.0e-06	lwlraps=0.86734	loss=1.71	val_loss=1.60	
t=34s					
Epoch 18/65	lr=1.0e-06	lwlraps=0.86453	loss=1.60	val_loss=1.63	
t=34s					
Epoch 19/65	lr=1.0e-06	lwlraps=0.86644	loss=1.60	val_loss=1.60	
t=34s					
Epoch 20/65	lr=1.0e-06	lwlraps=0.86256	loss=1.87	val_loss=1.62	
t=34s					
Epoch 21/65	lr=1.0e-07	lwlraps=0.86764	loss=1.49	val_loss=1.60	
t=34s					

Loading weights from model_2.pt

- Re-fitting with curated data only (2/2)

Epoch 1/65	lr=1.0e-04	lwlraps=0.86222	loss=1.44	val_loss=1.60	
t=34s					
Epoch 2/65	lr=1.0e-04	lwlraps=0.85909	loss=1.77	val_loss=1.65	
t=34s					
Epoch 3/65	lr=1.0e-04	lwlraps=0.85946	loss=1.73	val_loss=1.70	
t=34s					
Epoch 4/65	lr=1.0e-04	lwlraps=0.86801	loss=1.85	val_loss=1.60	
t=34s					
Epoch 5/65	lr=1.0e-04	lwlraps=0.86079	loss=1.68	val_loss=1.63	
t=34s					

Epoch 6/65 t=34s	lr=1.0e-05	lwlrwrap=0.85941	loss=1.58	val_loss=1.65
Epoch 7/65 t=34s	lr=1.0e-05	lwlrwrap=0.86306	loss=1.48	val_loss=1.59
Epoch 8/65 t=34s	lr=1.0e-05	lwlrwrap=0.86403	loss=1.71	val_loss=1.62
Epoch 9/65 t=34s	lr=1.0e-05	lwlrwrap=0.86150	loss=1.60	val_loss=1.62
Epoch 10/65 t=34s	lr=1.0e-05	lwlrwrap=0.86144	loss=1.48	val_loss=1.59
Epoch 11/65 t=34s	lr=1.0e-05	lwlrwrap=0.86447	loss=1.66	val_loss=1.63
Epoch 12/65 t=34s	lr=1.0e-05	lwlrwrap=0.86403	loss=1.72	val_loss=1.63
Epoch 13/65 t=34s	lr=1.0e-05	lwlrwrap=0.86341	loss=1.71	val_loss=1.64
Epoch 14/65 t=34s	lr=1.0e-05	lwlrwrap=0.86714	loss=1.67	val_loss=1.60
Epoch 15/65 t=34s	lr=1.0e-06	lwlrwrap=0.86530	loss=1.47	val_loss=1.60
Epoch 16/65 t=34s	lr=1.0e-06	lwlrwrap=0.86637	loss=1.69	val_loss=1.62
Epoch 17/65 t=34s	lr=1.0e-06	lwlrwrap=0.86267	loss=1.68	val_loss=1.65
Epoch 18/65 t=34s	lr=1.0e-06	lwlrwrap=0.86653	loss=1.80	val_loss=1.62
Epoch 19/65 t=34s	lr=1.0e-07	lwlrwrap=0.86404	loss=1.66	val_loss=1.62

Loading weights from model_2.pt

- Predicting with 25 TTA

lwlrwrap : Scored 0.8733 on validation data

Done in 81.1 minutes

----- Fold 3 -----

- Fitting

Epoch 1/65 t=50s	lr=5.0e-05	lwlrwrap=0.24414	loss=55.44	val_loss=55.32
Epoch 2/65 t=50s	lr=2.0e-04	lwlrwrap=0.30122	loss=52.45	val_loss=47.40
Epoch 3/65 t=50s	lr=3.6e-04	lwlrwrap=0.27746	loss=34.35	val_loss=21.69

Epoch 4/65 t=50s	lr=5.1e-04	lwlraps=0.35608	loss=10.50	val_loss=5.62
Epoch 5/65 t=50s	lr=6.7e-04	lwlraps=0.36117	loss=5.11	val_loss=5.15
Epoch 6/65 t=50s	lr=8.2e-04	lwlraps=0.42873	loss=4.66	val_loss=4.71
Epoch 7/65 t=50s	lr=9.7e-04	lwlraps=0.53984	loss=4.51	val_loss=3.84
Epoch 8/65 t=50s	lr=1.0e-03	lwlraps=0.46399	loss=4.30	val_loss=4.12
Epoch 9/65 t=50s	lr=1.0e-03	lwlraps=0.57711	loss=4.20	val_loss=3.63
Epoch 10/65 t=50s	lr=1.0e-03	lwlraps=0.53918	loss=3.95	val_loss=3.81
Epoch 11/65 t=50s	lr=9.9e-04	lwlraps=0.61374	loss=4.00	val_loss=3.40
Epoch 12/65 t=50s	lr=9.9e-04	lwlraps=0.56248	loss=3.86	val_loss=3.63
Epoch 13/65 t=50s	lr=9.9e-04	lwlraps=0.69307	loss=3.69	val_loss=2.90
Epoch 14/65 t=50s	lr=9.8e-04	lwlraps=0.72925	loss=3.73	val_loss=2.77
Epoch 15/65 t=50s	lr=9.7e-04	lwlraps=0.69479	loss=3.73	val_loss=2.92
Epoch 16/65 t=50s	lr=9.6e-04	lwlraps=0.69144	loss=3.63	val_loss=2.95
Epoch 17/65 t=50s	lr=9.5e-04	lwlraps=0.66020	loss=3.49	val_loss=3.14
Epoch 18/65 t=50s	lr=9.4e-04	lwlraps=0.71340	loss=3.61	val_loss=2.81
Epoch 19/65 t=50s	lr=9.3e-04	lwlraps=0.72814	loss=3.39	val_loss=2.66
Epoch 20/65 t=50s	lr=9.2e-04	lwlraps=0.71464	loss=3.40	val_loss=2.79
Epoch 21/65 t=50s	lr=9.0e-04	lwlraps=0.73708	loss=3.38	val_loss=2.64
Epoch 22/65 t=50s	lr=8.9e-04	lwlraps=0.78120	loss=3.20	val_loss=2.41
Epoch 23/65 t=50s	lr=8.7e-04	lwlraps=0.74806	loss=3.23	val_loss=2.52
Epoch 24/65 t=50s	lr=8.6e-04	lwlraps=0.75222	loss=3.29	val_loss=2.50
Epoch 25/65 t=50s	lr=8.4e-04	lwlraps=0.73320	loss=3.18	val_loss=2.63
Epoch 26/65 t=50s	lr=8.2e-04	lwlraps=0.77417	loss=3.21	val_loss=2.33
Epoch 27/65 t=50s	lr=8.0e-04	lwlraps=0.79568	loss=3.09	val_loss=2.23

Epoch 28/65 t=50s	lr=7.8e-04	lwlrwrap=0.74394	loss=3.08	val_loss=2.56
Epoch 29/65 t=50s	lr=7.6e-04	lwlrwrap=0.80027	loss=3.17	val_loss=2.15
Epoch 30/65 t=50s	lr=7.4e-04	lwlrwrap=0.79424	loss=2.97	val_loss=2.22
Epoch 31/65 t=50s	lr=7.2e-04	lwlrwrap=0.80583	loss=2.95	val_loss=2.13
Epoch 32/65 t=50s	lr=7.0e-04	lwlrwrap=0.79787	loss=3.06	val_loss=2.25
Epoch 33/65 t=50s	lr=6.8e-04	lwlrwrap=0.80842	loss=2.92	val_loss=2.12
Epoch 34/65 t=50s	lr=6.5e-04	lwlrwrap=0.79978	loss=2.72	val_loss=2.19
Epoch 35/65 t=50s	lr=6.3e-04	lwlrwrap=0.80208	loss=2.93	val_loss=2.15
Epoch 36/65 t=50s	lr=6.1e-04	lwlrwrap=0.80972	loss=2.73	val_loss=2.07
Epoch 37/65 t=50s	lr=5.8e-04	lwlrwrap=0.82559	loss=2.95	val_loss=1.99
Epoch 38/65 t=50s	lr=5.6e-04	lwlrwrap=0.82783	loss=2.75	val_loss=1.90
Epoch 39/65 t=50s	lr=5.4e-04	lwlrwrap=0.82657	loss=2.79	val_loss=1.94
Epoch 40/65 t=50s	lr=5.1e-04	lwlrwrap=0.83205	loss=2.69	val_loss=1.91
Epoch 41/65 t=50s	lr=4.9e-04	lwlrwrap=0.83138	loss=2.64	val_loss=1.92
Epoch 42/65 t=50s	lr=4.6e-04	lwlrwrap=0.83177	loss=2.48	val_loss=1.95
Epoch 43/65 t=50s	lr=4.4e-04	lwlrwrap=0.83187	loss=2.53	val_loss=1.90
Epoch 44/65 t=50s	lr=4.2e-04	lwlrwrap=0.84055	loss=2.56	val_loss=1.81
Epoch 45/65 t=50s	lr=3.9e-04	lwlrwrap=0.83422	loss=2.58	val_loss=1.88
Epoch 46/65 t=50s	lr=3.7e-04	lwlrwrap=0.84742	loss=2.46	val_loss=1.77
Epoch 47/65 t=50s	lr=3.5e-04	lwlrwrap=0.84145	loss=2.29	val_loss=1.80
Epoch 48/65 t=50s	lr=3.2e-04	lwlrwrap=0.85495	loss=2.43	val_loss=1.70
Epoch 49/65 t=50s	lr=3.0e-04	lwlrwrap=0.84674	loss=2.39	val_loss=1.75
Epoch 50/65 t=50s	lr=2.8e-04	lwlrwrap=0.85250	loss=2.14	val_loss=1.67
Epoch 51/65 t=50s	lr=2.6e-04	lwlrwrap=0.84839	loss=2.42	val_loss=1.70

Epoch 52/65 t=50s	lr=2.4e-04	lwlrwrap=0.84246	loss=2.26	val_loss=1.72
Epoch 53/65 t=50s	lr=2.2e-04	lwlrwrap=0.84758	loss=2.17	val_loss=1.73
Epoch 54/65 t=50s	lr=2.0e-04	lwlrwrap=0.85836	loss=2.10	val_loss=1.68
Epoch 55/65 t=50s	lr=1.8e-04	lwlrwrap=0.85988	loss=2.25	val_loss=1.62
Epoch 56/65 t=50s	lr=1.6e-04	lwlrwrap=0.86222	loss=2.20	val_loss=1.66
Epoch 57/65 t=50s	lr=1.4e-04	lwlrwrap=0.86285	loss=2.17	val_loss=1.64
Epoch 58/65 t=50s	lr=1.3e-04	lwlrwrap=0.85746	loss=2.18	val_loss=1.63
Epoch 59/65 t=50s	lr=1.1e-04	lwlrwrap=0.86874	loss=1.87	val_loss=1.59
Epoch 60/65 t=50s	lr=9.6e-05	lwlrwrap=0.86401	loss=2.15	val_loss=1.59
Epoch 61/65 t=50s	lr=8.2e-05	lwlrwrap=0.86070	loss=2.03	val_loss=1.61
Epoch 62/65 t=50s	lr=6.9e-05	lwlrwrap=0.85960	loss=2.10	val_loss=1.64
Epoch 63/65 t=50s	lr=5.7e-05	lwlrwrap=0.86222	loss=2.00	val_loss=1.60
Epoch 64/65 t=50s	lr=4.7e-05	lwlrwrap=0.86746	loss=1.96	val_loss=1.56
Epoch 65/65 t=50s	lr=3.7e-05	lwlrwrap=0.86102	loss=2.04	val_loss=1.65

- Re-fitting with curated data only (1/2)

Epoch 1/65 t=34s	lr=1.0e-04	lwlrwrap=0.85773	loss=1.65	val_loss=1.70
Epoch 2/65 t=34s	lr=1.0e-04	lwlrwrap=0.86105	loss=1.70	val_loss=1.66
Epoch 3/65 t=34s	lr=1.0e-04	lwlrwrap=0.86759	loss=1.65	val_loss=1.62
Epoch 4/65 t=34s	lr=1.0e-04	lwlrwrap=0.86566	loss=1.71	val_loss=1.61
Epoch 5/65 t=34s	lr=1.0e-04	lwlrwrap=0.86776	loss=1.70	val_loss=1.61
Epoch 6/65 t=34s	lr=1.0e-04	lwlrwrap=0.86389	loss=1.62	val_loss=1.66
Epoch 7/65 t=34s	lr=1.0e-04	lwlrwrap=0.86537	loss=1.59	val_loss=1.63
Epoch 8/65 t=34s	lr=1.0e-04	lwlrwrap=0.85872	loss=2.01	val_loss=1.62
Epoch 9/65	lr=1.0e-04	lwlrwrap=0.86514	loss=1.93	val_loss=1.67

```

t=34s
Epoch 10/65      lr=1.0e-05      lwlraps=0.86215      loss=1.88      val_loss=1.70
t=34s
Epoch 11/65      lr=1.0e-05      lwlraps=0.86664      loss=1.60      val_loss=1.61
t=34s
Epoch 12/65      lr=1.0e-05      lwlraps=0.86899      loss=1.58      val_loss=1.62
t=34s
Epoch 13/65      lr=1.0e-05      lwlraps=0.86687      loss=1.68      val_loss=1.63
t=34s
Epoch 14/65      lr=1.0e-06      lwlraps=0.86990      loss=1.61      val_loss=1.59
t=34s
Epoch 15/65      lr=1.0e-06      lwlraps=0.86878      loss=1.60      val_loss=1.61
t=34s
Epoch 16/65      lr=1.0e-06      lwlraps=0.86478      loss=1.40      val_loss=1.60
t=34s
Epoch 17/65      lr=1.0e-06      lwlraps=0.86378      loss=1.35      val_loss=1.59
t=34s
Epoch 18/65      lr=1.0e-06      lwlraps=0.86469      loss=1.50      val_loss=1.63
t=34s
Epoch 19/65      lr=1.0e-07      lwlraps=0.87557      loss=1.50      val_loss=1.58
t=34s
Loading weights from model_3.pt

```

- Re-fitting with curated data only (2/2)

```

Epoch 1/65      lr=1.0e-04      lwlraps=0.86746      loss=1.58      val_loss=1.58
t=34s
Epoch 2/65      lr=1.0e-04      lwlraps=0.86260      loss=1.87      val_loss=1.64
t=34s
Epoch 3/65      lr=1.0e-04      lwlraps=0.86753      loss=1.76      val_loss=1.62
t=34s
Epoch 4/65      lr=1.0e-04      lwlraps=0.85922      loss=1.76      val_loss=1.72
t=34s
Epoch 5/65      lr=1.0e-04      lwlraps=0.86343      loss=1.67      val_loss=1.63
t=34s
Epoch 6/65      lr=1.0e-05      lwlraps=0.86529      loss=1.55      val_loss=1.63
t=34s
Epoch 7/65      lr=1.0e-05      lwlraps=0.87115      loss=1.62      val_loss=1.56
t=34s
Epoch 8/65      lr=1.0e-05      lwlraps=0.86554      loss=1.43      val_loss=1.64
t=34s
Epoch 9/65      lr=1.0e-05      lwlraps=0.86417      loss=1.53      val_loss=1.61
t=34s
Epoch 10/65     lr=1.0e-05      lwlraps=0.86943      loss=1.85      val_loss=1.64
t=34s
Epoch 11/65     lr=1.0e-05      lwlraps=0.87077      loss=1.56      val_loss=1.60
t=34s

```

Epoch 12/65 t=34s	lr=1.0e-06	lwlrwrap=0.86846	loss=1.49	val_loss=1.61
Epoch 13/65 t=34s	lr=1.0e-06	lwlrwrap=0.86334	loss=1.58	val_loss=1.64
Epoch 14/65 t=34s	lr=1.0e-06	lwlrwrap=0.86422	loss=1.68	val_loss=1.62
Epoch 15/65 t=34s	lr=1.0e-06	lwlrwrap=0.86728	loss=1.60	val_loss=1.63
Epoch 16/65 t=34s	lr=1.0e-07	lwlrwrap=0.86321	loss=1.45	val_loss=1.64

Loading weights from model_3.pt

- Predicting with 25 TTA

lwlrwrap : Scored 0.8722 on validation data

Done in 78.0 minutes

----- Fold 4 -----

- Fitting

Epoch 1/65 t=50s	lr=5.0e-05	lwlrwrap=0.29778	loss=55.13	val_loss=54.73
Epoch 2/65 t=50s	lr=2.0e-04	lwlrwrap=0.31540	loss=52.07	val_loss=47.28
Epoch 3/65 t=50s	lr=3.6e-04	lwlrwrap=0.34737	loss=33.77	val_loss=19.99
Epoch 4/65 t=50s	lr=5.1e-04	lwlrwrap=0.35495	loss=10.24	val_loss=5.97
Epoch 5/65 t=50s	lr=6.7e-04	lwlrwrap=0.42731	loss=5.06	val_loss=4.36
Epoch 6/65 t=50s	lr=8.2e-04	lwlrwrap=0.40044	loss=4.62	val_loss=4.49
Epoch 7/65 t=50s	lr=9.7e-04	lwlrwrap=0.52729	loss=4.48	val_loss=3.93
Epoch 8/65 t=50s	lr=1.0e-03	lwlrwrap=0.55094	loss=4.28	val_loss=3.75
Epoch 9/65 t=50s	lr=1.0e-03	lwlrwrap=0.50160	loss=4.12	val_loss=4.12
Epoch 10/65 t=50s	lr=1.0e-03	lwlrwrap=0.64308	loss=3.95	val_loss=3.25
Epoch 11/65 t=50s	lr=9.9e-04	lwlrwrap=0.61630	loss=3.76	val_loss=3.34
Epoch 12/65 t=50s	lr=9.9e-04	lwlrwrap=0.64073	loss=3.78	val_loss=3.21

Epoch 13/65 t=50s	lr=9.9e-04	lwlrwrap=0.67801	loss=3.68	val_loss=2.93
Epoch 14/65 t=50s	lr=9.8e-04	lwlrwrap=0.69722	loss=3.62	val_loss=2.85
Epoch 15/65 t=50s	lr=9.7e-04	lwlrwrap=0.73045	loss=3.52	val_loss=2.69
Epoch 16/65 t=50s	lr=9.6e-04	lwlrwrap=0.70490	loss=3.63	val_loss=2.77
Epoch 17/65 t=50s	lr=9.5e-04	lwlrwrap=0.75872	loss=3.37	val_loss=2.47
Epoch 18/65 t=50s	lr=9.4e-04	lwlrwrap=0.72439	loss=3.49	val_loss=2.72
Epoch 19/65 t=50s	lr=9.3e-04	lwlrwrap=0.76779	loss=3.50	val_loss=2.47
Epoch 20/65 t=50s	lr=9.2e-04	lwlrwrap=0.74829	loss=3.34	val_loss=2.57
Epoch 21/65 t=50s	lr=9.0e-04	lwlrwrap=0.73985	loss=3.43	val_loss=2.63
Epoch 22/65 t=50s	lr=8.9e-04	lwlrwrap=0.77689	loss=3.30	val_loss=2.33
Epoch 23/65 t=50s	lr=8.7e-04	lwlrwrap=0.78717	loss=3.30	val_loss=2.31
Epoch 24/65 t=50s	lr=8.6e-04	lwlrwrap=0.75625	loss=3.15	val_loss=2.49
Epoch 25/65 t=50s	lr=8.4e-04	lwlrwrap=0.76155	loss=3.25	val_loss=2.47
Epoch 26/65 t=50s	lr=8.2e-04	lwlrwrap=0.76212	loss=3.10	val_loss=2.34
Epoch 27/65 t=50s	lr=8.0e-04	lwlrwrap=0.80013	loss=3.08	val_loss=2.14
Epoch 28/65 t=50s	lr=7.8e-04	lwlrwrap=0.78997	loss=3.10	val_loss=2.26
Epoch 29/65 t=50s	lr=7.6e-04	lwlrwrap=0.80464	loss=3.08	val_loss=2.15
Epoch 30/65 t=50s	lr=7.4e-04	lwlrwrap=0.77194	loss=3.07	val_loss=2.36
Epoch 31/65 t=50s	lr=7.2e-04	lwlrwrap=0.81336	loss=3.07	val_loss=2.11
Epoch 32/65 t=50s	lr=7.0e-04	lwlrwrap=0.80988	loss=2.96	val_loss=2.05
Epoch 33/65 t=50s	lr=6.8e-04	lwlrwrap=0.82385	loss=2.99	val_loss=2.00
Epoch 34/65 t=50s	lr=6.5e-04	lwlrwrap=0.80126	loss=2.86	val_loss=2.12
Epoch 35/65 t=50s	lr=6.3e-04	lwlrwrap=0.80687	loss=2.74	val_loss=2.12
Epoch 36/65 t=50s	lr=6.1e-04	lwlrwrap=0.81401	loss=2.74	val_loss=2.10

Epoch 37/65 t=50s	lr=5.8e-04	lwlrwrap=0.81070	loss=2.61	val_loss=2.06
Epoch 38/65 t=50s	lr=5.6e-04	lwlrwrap=0.82920	loss=2.83	val_loss=1.90
Epoch 39/65 t=50s	lr=5.4e-04	lwlrwrap=0.82754	loss=2.71	val_loss=1.91
Epoch 40/65 t=50s	lr=5.1e-04	lwlrwrap=0.84746	loss=2.67	val_loss=1.73
Epoch 41/65 t=50s	lr=4.9e-04	lwlrwrap=0.83863	loss=2.54	val_loss=1.81
Epoch 42/65 t=50s	lr=4.6e-04	lwlrwrap=0.84833	loss=2.69	val_loss=1.79
Epoch 43/65 t=50s	lr=4.4e-04	lwlrwrap=0.84437	loss=2.55	val_loss=1.77
Epoch 44/65 t=50s	lr=4.2e-04	lwlrwrap=0.84532	loss=2.72	val_loss=1.76
Epoch 45/65 t=50s	lr=3.9e-04	lwlrwrap=0.82494	loss=2.59	val_loss=1.96
Epoch 46/65 t=50s	lr=3.7e-04	lwlrwrap=0.85755	loss=2.41	val_loss=1.74
Epoch 47/65 t=50s	lr=3.5e-04	lwlrwrap=0.85326	loss=2.56	val_loss=1.71
Epoch 48/65 t=50s	lr=3.2e-04	lwlrwrap=0.85765	loss=2.30	val_loss=1.67
Epoch 49/65 t=50s	lr=3.0e-04	lwlrwrap=0.85991	loss=2.40	val_loss=1.67
Epoch 50/65 t=50s	lr=2.8e-04	lwlrwrap=0.85265	loss=2.56	val_loss=1.67
Epoch 51/65 t=50s	lr=2.6e-04	lwlrwrap=0.85622	loss=2.28	val_loss=1.68
Epoch 52/65 t=50s	lr=2.4e-04	lwlrwrap=0.85404	loss=2.36	val_loss=1.70
Epoch 53/65 t=50s	lr=2.2e-04	lwlrwrap=0.87021	loss=2.22	val_loss=1.60
Epoch 54/65 t=50s	lr=2.0e-04	lwlrwrap=0.86594	loss=2.19	val_loss=1.62
Epoch 55/65 t=50s	lr=1.8e-04	lwlrwrap=0.86740	loss=2.14	val_loss=1.56
Epoch 56/65 t=50s	lr=1.6e-04	lwlrwrap=0.86882	loss=2.16	val_loss=1.56
Epoch 57/65 t=50s	lr=1.4e-04	lwlrwrap=0.86947	loss=2.00	val_loss=1.55
Epoch 58/65 t=50s	lr=1.3e-04	lwlrwrap=0.86948	loss=2.19	val_loss=1.56
Epoch 59/65 t=50s	lr=1.1e-04	lwlrwrap=0.87085	loss=2.00	val_loss=1.55
Epoch 60/65 t=50s	lr=9.6e-05	lwlrwrap=0.87223	loss=2.21	val_loss=1.56

Epoch 61/65 t=50s	lr=8.2e-05	lwlrwrap=0.87082	loss=2.07	val_loss=1.55
Epoch 62/65 t=50s	lr=6.9e-05	lwlrwrap=0.86583	loss=2.05	val_loss=1.54
Epoch 63/65 t=50s	lr=5.7e-05	lwlrwrap=0.86745	loss=2.05	val_loss=1.59
Epoch 64/65 t=50s	lr=4.7e-05	lwlrwrap=0.86983	loss=2.00	val_loss=1.54
Epoch 65/65 t=50s	lr=3.7e-05	lwlrwrap=0.86802	loss=2.05	val_loss=1.60

- Re-fitting with curated data only (1/2)

Epoch 1/65 t=34s	lr=1.0e-04	lwlrwrap=0.87000	loss=1.69	val_loss=1.58
Epoch 2/65 t=34s	lr=1.0e-04	lwlrwrap=0.86720	loss=1.75	val_loss=1.60
Epoch 3/65 t=34s	lr=1.0e-04	lwlrwrap=0.86865	loss=1.63	val_loss=1.58
Epoch 4/65 t=34s	lr=1.0e-04	lwlrwrap=0.86356	loss=1.75	val_loss=1.64
Epoch 5/65 t=34s	lr=1.0e-04	lwlrwrap=0.86371	loss=1.85	val_loss=1.61
Epoch 6/65 t=34s	lr=1.0e-05	lwlrwrap=0.86565	loss=1.59	val_loss=1.60
Epoch 7/65 t=34s	lr=1.0e-05	lwlrwrap=0.86302	loss=1.68	val_loss=1.61
Epoch 8/65 t=34s	lr=1.0e-05	lwlrwrap=0.87480	loss=1.79	val_loss=1.56
Epoch 9/65 t=34s	lr=1.0e-05	lwlrwrap=0.86982	loss=1.48	val_loss=1.57
Epoch 10/65 t=34s	lr=1.0e-05	lwlrwrap=0.86446	loss=1.49	val_loss=1.60
Epoch 11/65 t=34s	lr=1.0e-05	lwlrwrap=0.86673	loss=1.79	val_loss=1.64
Epoch 12/65 t=34s	lr=1.0e-05	lwlrwrap=0.86619	loss=1.62	val_loss=1.60
Epoch 13/65 t=34s	lr=1.0e-06	lwlrwrap=0.86983	loss=1.47	val_loss=1.56
Epoch 14/65 t=34s	lr=1.0e-06	lwlrwrap=0.86619	loss=1.64	val_loss=1.60
Epoch 15/65 t=34s	lr=1.0e-06	lwlrwrap=0.86529	loss=1.96	val_loss=1.65
Epoch 16/65 t=34s	lr=1.0e-06	lwlrwrap=0.86539	loss=1.48	val_loss=1.58
Epoch 17/65 t=34s	lr=1.0e-06	lwlrwrap=0.85853	loss=1.97	val_loss=1.66
Epoch 18/65	lr=1.0e-07	lwlrwrap=0.86285	loss=1.57	val_loss=1.62

t=34s

Loading weights from model_4.pt

- Re-fitting with curated data only (2/2)

Epoch 1/65 t=34s	lr=1.0e-04	lwlraps=0.86174	loss=1.67	val_loss=1.63
Epoch 2/65 t=34s	lr=1.0e-04	lwlraps=0.87102	loss=1.48	val_loss=1.59
Epoch 3/65 t=34s	lr=1.0e-04	lwlraps=0.87003	loss=1.57	val_loss=1.60
Epoch 4/65 t=34s	lr=1.0e-04	lwlraps=0.86408	loss=1.66	val_loss=1.66
Epoch 5/65 t=34s	lr=1.0e-04	lwlraps=0.86615	loss=1.71	val_loss=1.65
Epoch 6/65 t=34s	lr=1.0e-04	lwlraps=0.87117	loss=1.73	val_loss=1.63
Epoch 7/65 t=34s	lr=1.0e-05	lwlraps=0.87320	loss=1.60	val_loss=1.60
Epoch 8/65 t=34s	lr=1.0e-05	lwlraps=0.86905	loss=1.50	val_loss=1.65
Epoch 9/65 t=34s	lr=1.0e-05	lwlraps=0.86685	loss=1.75	val_loss=1.66
Epoch 10/65 t=34s	lr=1.0e-05	lwlraps=0.87011	loss=1.52	val_loss=1.61
Epoch 11/65 t=34s	lr=1.0e-06	lwlraps=0.87083	loss=1.52	val_loss=1.62
Epoch 12/65 t=34s	lr=1.0e-06	lwlraps=0.87009	loss=1.67	val_loss=1.64
Epoch 13/65 t=34s	lr=1.0e-06	lwlraps=0.87265	loss=1.39	val_loss=1.57
Epoch 14/65 t=34s	lr=1.0e-06	lwlraps=0.86978	loss=1.36	val_loss=1.58
Epoch 15/65 t=34s	lr=1.0e-06	lwlraps=0.86828	loss=1.52	val_loss=1.64
Epoch 16/65 t=34s	lr=1.0e-06	lwlraps=0.87776	loss=1.54	val_loss=1.58
Epoch 17/65 t=34s	lr=1.0e-06	lwlraps=0.86879	loss=1.66	val_loss=1.62
Epoch 18/65 t=34s	lr=1.0e-07	lwlraps=0.87019	loss=1.40	val_loss=1.60

Loading weights from model_4.pt

- Predicting with 25 TTA

lwlrap : Scored 0.8784 on validation data

Done in 78.8 minutes

----- Fold 5 -----

- Fitting

Epoch 1/65 t=50s	lr=5.0e-05	lwlrap=0.25448	loss=55.14	val_loss=54.92
Epoch 2/65 t=50s	lr=2.0e-04	lwlrap=0.30980	loss=52.22	val_loss=47.82
Epoch 3/65 t=50s	lr=3.6e-04	lwlrap=0.27796	loss=34.42	val_loss=21.34
Epoch 4/65 t=50s	lr=5.1e-04	lwlrap=0.32000	loss=10.62	val_loss=5.56
Epoch 5/65 t=50s	lr=6.7e-04	lwlrap=0.33746	loss=5.05	val_loss=4.93
Epoch 6/65 t=50s	lr=8.2e-04	lwlrap=0.47756	loss=4.68	val_loss=4.30
Epoch 7/65 t=50s	lr=9.7e-04	lwlrap=0.51861	loss=4.43	val_loss=4.03
Epoch 8/65 t=50s	lr=1.0e-03	lwlrap=0.44956	loss=4.23	val_loss=4.27
Epoch 9/65 t=50s	lr=1.0e-03	lwlrap=0.45506	loss=4.16	val_loss=4.38
Epoch 10/65 t=50s	lr=1.0e-03	lwlrap=0.56731	loss=4.08	val_loss=3.59
Epoch 11/65 t=50s	lr=9.9e-04	lwlrap=0.60904	loss=3.86	val_loss=3.40
Epoch 12/65 t=50s	lr=9.9e-04	lwlrap=0.65139	loss=3.75	val_loss=3.13
Epoch 13/65 t=50s	lr=9.9e-04	lwlrap=0.64167	loss=3.78	val_loss=3.19
Epoch 14/65 t=50s	lr=9.8e-04	lwlrap=0.65755	loss=3.65	val_loss=3.10
Epoch 15/65 t=50s	lr=9.7e-04	lwlrap=0.68737	loss=3.68	val_loss=2.89
Epoch 16/65 t=50s	lr=9.6e-04	lwlrap=0.64564	loss=3.56	val_loss=3.13
Epoch 17/65 t=50s	lr=9.5e-04	lwlrap=0.70171	loss=3.48	val_loss=2.76
Epoch 18/65 t=50s	lr=9.4e-04	lwlrap=0.71702	loss=3.48	val_loss=2.74
Epoch 19/65 t=50s	lr=9.3e-04	lwlrap=0.67062	loss=3.49	val_loss=2.92
Epoch 20/65 t=50s	lr=9.2e-04	lwlrap=0.71658	loss=3.40	val_loss=2.69

Epoch 21/65 t=50s	lr=9.0e-04	lwlrwrap=0.73064	loss=3.35	val_loss=2.58
Epoch 22/65 t=50s	lr=8.9e-04	lwlrwrap=0.73799	loss=3.29	val_loss=2.51
Epoch 23/65 t=50s	lr=8.7e-04	lwlrwrap=0.75789	loss=3.28	val_loss=2.38
Epoch 24/65 t=50s	lr=8.6e-04	lwlrwrap=0.75757	loss=3.21	val_loss=2.42
Epoch 25/65 t=50s	lr=8.4e-04	lwlrwrap=0.73907	loss=3.42	val_loss=2.47
Epoch 26/65 t=50s	lr=8.2e-04	lwlrwrap=0.73268	loss=3.13	val_loss=2.57
Epoch 27/65 t=50s	lr=8.0e-04	lwlrwrap=0.77122	loss=3.05	val_loss=2.35
Epoch 28/65 t=50s	lr=7.8e-04	lwlrwrap=0.72806	loss=3.04	val_loss=2.52
Epoch 29/65 t=50s	lr=7.6e-04	lwlrwrap=0.77886	loss=3.11	val_loss=2.30
Epoch 30/65 t=50s	lr=7.4e-04	lwlrwrap=0.79369	loss=2.93	val_loss=2.17
Epoch 31/65 t=50s	lr=7.2e-04	lwlrwrap=0.78362	loss=2.85	val_loss=2.18
Epoch 32/65 t=50s	lr=7.0e-04	lwlrwrap=0.77815	loss=2.88	val_loss=2.25
Epoch 33/65 t=50s	lr=6.8e-04	lwlrwrap=0.79891	loss=2.84	val_loss=2.10
Epoch 34/65 t=50s	lr=6.5e-04	lwlrwrap=0.79023	loss=2.80	val_loss=2.20
Epoch 35/65 t=50s	lr=6.3e-04	lwlrwrap=0.78654	loss=3.01	val_loss=2.22
Epoch 36/65 t=50s	lr=6.1e-04	lwlrwrap=0.79897	loss=2.84	val_loss=2.11
Epoch 37/65 t=50s	lr=5.8e-04	lwlrwrap=0.78560	loss=2.74	val_loss=2.18
Epoch 38/65 t=50s	lr=5.6e-04	lwlrwrap=0.80790	loss=2.81	val_loss=2.01
Epoch 39/65 t=50s	lr=5.4e-04	lwlrwrap=0.81328	loss=2.71	val_loss=2.01
Epoch 40/65 t=50s	lr=5.1e-04	lwlrwrap=0.82248	loss=2.67	val_loss=1.92
Epoch 41/65 t=50s	lr=4.9e-04	lwlrwrap=0.81284	loss=2.42	val_loss=1.93
Epoch 42/65 t=50s	lr=4.6e-04	lwlrwrap=0.81495	loss=2.60	val_loss=2.05
Epoch 43/65 t=50s	lr=4.4e-04	lwlrwrap=0.82048	loss=2.53	val_loss=1.89
Epoch 44/65 t=50s	lr=4.2e-04	lwlrwrap=0.80833	loss=2.43	val_loss=2.02

Epoch 45/65 t=50s	lr=3.9e-04	lwlrwrap=0.81206	loss=2.49	val_loss=1.96
Epoch 46/65 t=50s	lr=3.7e-04	lwlrwrap=0.82150	loss=2.51	val_loss=1.90
Epoch 47/65 t=50s	lr=3.5e-04	lwlrwrap=0.81518	loss=2.43	val_loss=1.90
Epoch 48/65 t=50s	lr=3.2e-04	lwlrwrap=0.83123	loss=2.30	val_loss=1.84
Epoch 49/65 t=50s	lr=3.0e-04	lwlrwrap=0.82540	loss=2.54	val_loss=1.85
Epoch 50/65 t=50s	lr=2.8e-04	lwlrwrap=0.83077	loss=2.26	val_loss=1.77
Epoch 51/65 t=50s	lr=2.6e-04	lwlrwrap=0.84408	loss=2.06	val_loss=1.71
Epoch 52/65 t=50s	lr=2.4e-04	lwlrwrap=0.83360	loss=2.19	val_loss=1.78
Epoch 53/65 t=50s	lr=2.2e-04	lwlrwrap=0.83359	loss=2.28	val_loss=1.79
Epoch 54/65 t=50s	lr=2.0e-04	lwlrwrap=0.84448	loss=2.12	val_loss=1.71
Epoch 55/65 t=50s	lr=1.8e-04	lwlrwrap=0.84334	loss=2.15	val_loss=1.70
Epoch 56/65 t=50s	lr=1.6e-04	lwlrwrap=0.83961	loss=2.07	val_loss=1.74
Epoch 57/65 t=50s	lr=1.4e-04	lwlrwrap=0.85070	loss=2.09	val_loss=1.70
Epoch 58/65 t=50s	lr=1.3e-04	lwlrwrap=0.84274	loss=2.07	val_loss=1.71
Epoch 59/65 t=50s	lr=1.1e-04	lwlrwrap=0.84327	loss=1.93	val_loss=1.71
Epoch 60/65 t=50s	lr=9.6e-05	lwlrwrap=0.85504	loss=1.73	val_loss=1.65
Epoch 61/65 t=50s	lr=8.2e-05	lwlrwrap=0.84435	loss=2.03	val_loss=1.71
Epoch 62/65 t=50s	lr=6.9e-05	lwlrwrap=0.84088	loss=1.98	val_loss=1.74
Epoch 63/65 t=50s	lr=5.7e-05	lwlrwrap=0.84461	loss=1.99	val_loss=1.71
Epoch 64/65 t=50s	lr=4.7e-05	lwlrwrap=0.84903	loss=2.02	val_loss=1.67
Epoch 65/65 t=50s	lr=3.7e-05	lwlrwrap=0.84859	loss=2.06	val_loss=1.68

- Re-fitting with curated data only (1/2)

Epoch 1/65 t=34s	lr=1.0e-04	lwlrwrap=0.84777	loss=1.78	val_loss=1.70
Epoch 2/65	lr=1.0e-04	lwlrwrap=0.85251	loss=1.99	val_loss=1.73

t=34s					
Epoch 3/65	lr=1.0e-04	lwlraps=0.85258	loss=1.79	val_loss=1.69	
t=34s					
Epoch 4/65	lr=1.0e-04	lwlraps=0.85211	loss=2.04	val_loss=1.74	
t=34s					
Epoch 5/65	lr=1.0e-04	lwlraps=0.83904	loss=1.69	val_loss=1.77	
t=34s					
Epoch 6/65	lr=1.0e-04	lwlraps=0.84374	loss=1.56	val_loss=1.74	
t=34s					
Epoch 7/65	lr=1.0e-04	lwlraps=0.84986	loss=1.58	val_loss=1.71	
t=34s					
Epoch 8/65	lr=1.0e-05	lwlraps=0.84990	loss=1.73	val_loss=1.72	
t=34s					
Epoch 9/65	lr=1.0e-05	lwlraps=0.85124	loss=1.52	val_loss=1.70	
t=34s					
Epoch 10/65	lr=1.0e-05	lwlraps=0.84645	loss=1.55	val_loss=1.74	
t=34s					
Epoch 11/65	lr=1.0e-05	lwlraps=0.84482	loss=1.75	val_loss=1.75	
t=34s					
Epoch 12/65	lr=1.0e-06	lwlraps=0.85043	loss=1.58	val_loss=1.72	
t=34s					
Epoch 13/65	lr=1.0e-06	lwlraps=0.85628	loss=1.75	val_loss=1.69	
t=34s					
Epoch 14/65	lr=1.0e-06	lwlraps=0.84924	loss=1.58	val_loss=1.74	
t=34s					
Epoch 15/65	lr=1.0e-06	lwlraps=0.84901	loss=1.70	val_loss=1.72	
t=34s					
Epoch 16/65	lr=1.0e-06	lwlraps=0.84862	loss=1.74	val_loss=1.76	
t=34s					
Epoch 17/65	lr=1.0e-06	lwlraps=0.85307	loss=1.77	val_loss=1.71	
t=34s					
Epoch 18/65	lr=1.0e-07	lwlraps=0.85125	loss=1.55	val_loss=1.71	
t=34s					

Loading weights from model_5.pt

- Re-fitting with curated data only (2/2)

Epoch 1/65	lr=1.0e-04	lwlraps=0.85309	loss=1.68	val_loss=1.71	
t=34s					
Epoch 2/65	lr=1.0e-04	lwlraps=0.85129	loss=1.68	val_loss=1.72	
t=34s					
Epoch 3/65	lr=1.0e-04	lwlraps=0.84475	loss=1.66	val_loss=1.73	
t=34s					
Epoch 4/65	lr=1.0e-04	lwlraps=0.84503	loss=1.63	val_loss=1.73	
t=34s					
Epoch 5/65	lr=1.0e-04	lwlraps=0.84878	loss=1.65	val_loss=1.75	
t=34s					

Epoch	lr	lwlap	loss	val_loss
Epoch 6/65 t=34s	1r=1.0e-05	lwlap=0.85322	loss=1.42	val_loss=1.68
Epoch 7/65 t=34s	1r=1.0e-05	lwlap=0.85235	loss=1.69	val_loss=1.73
Epoch 8/65 t=34s	1r=1.0e-05	lwlap=0.85302	loss=1.27	val_loss=1.69
Epoch 9/65 t=34s	1r=1.0e-05	lwlap=0.85209	loss=1.89	val_loss=1.72
Epoch 10/65 t=34s	1r=1.0e-05	lwlap=0.85469	loss=1.89	val_loss=1.72
Epoch 11/65 t=34s	1r=1.0e-06	lwlap=0.84940	loss=1.72	val_loss=1.73
Epoch 12/65 t=34s	1r=1.0e-06	lwlap=0.84937	loss=1.53	val_loss=1.74
Epoch 13/65 t=34s	1r=1.0e-06	lwlap=0.85159	loss=1.41	val_loss=1.70
Epoch 14/65 t=34s	1r=1.0e-06	lwlap=0.84585	loss=1.57	val_loss=1.73
Epoch 15/65 t=34s	1r=1.0e-07	lwlap=0.85115	loss=1.66	val_loss=1.69

Loading weights from model_5.pt

- Predicting with 25 TTA

lwlap : Scored 0.8551 on validation data

Done in 76.9 minutes

```
[57]: score_class, weight = lwlap(y_train, pred_oof)
      score = (score_class * weight).sum()
```

```
[58]: print(f'Local CV score is : {score:.4f}')
```

Local CV score is : 0.8703

```
[59]: score_df = pd.DataFrame(np.array([labels, np.round(score_class, 5), np.
      ↳round(weight, 5)]).T, columns=['Label', 'Score', 'Weight']).
      ↳set_index('Label').sort_values('Score')
      score_df
```

Label	Score	Weight
Walk_and_footsteps	0.61816	0.01305
Squeak	0.66539	0.01305

Fill_(with_liquid)	0.685	0.0087
Buzz	0.7367	0.00975
Yell	0.73886	0.01305
Sink_(filling_or_washing)	0.74065	0.01305
Traffic_noise_and_roadway_noise	0.74117	0.01305
Water_tap_and_faucet	0.74604	0.01305
Mechanical_fan	0.76589	0.00853
Hiss	0.76768	0.01305
Tap	0.77258	0.01305
Cutlery_and_silverware	0.77369	0.01305
Microwave_oven	0.77416	0.01305
Bus	0.78029	0.01305
Motorcycle	0.7861	0.01305
Trickle_and_dribble	0.79233	0.00922
Gasp	0.79252	0.00835
Chink_and_clink	0.80138	0.01305
Dishes_and_pots_and_pans	0.81386	0.01305
Bathtub_(filling_or_washing)	0.81556	0.01305
Drip	0.81887	0.01305
Cupboard_open_or_close	0.82087	0.01305
Frying_(food)	0.82089	0.01096
Accelerating_and_revving_and_vroom	0.82391	0.01305
Run	0.82789	0.01305
Slam	0.83424	0.01305
Stream	0.83956	0.01288
Chewing_and_mastication	0.84337	0.01305
Crowd	0.85428	0.01305
Printer	0.86122	0.01305
...
Raindrop	0.90862	0.01305
Toilet_flush	0.90946	0.01305
Keys_jangling	0.91089	0.01305
Child_speech_and_kid_speaking	0.9148	0.01305
Shatter	0.91578	0.01305
Race_car_and_auto_racing	0.91667	0.00975
Writing	0.91689	0.01305
Gong	0.91689	0.01305
Zipper_(clothing)	0.91949	0.01305
Meow	0.92131	0.01305
Church_bell	0.92703	0.01305
Screaming	0.92722	0.01305
Sigh	0.93567	0.00992
Acoustic_guitar	0.94529	0.01305
Bark	0.94714	0.01305
Bass_drum	0.94768	0.01305
Fart	0.952	0.01305
Whispering	0.95833	0.01288

Harmonica	0.96003	0.01305
Accordion	0.96099	0.00818
Glockenspiel	0.96429	0.00975
Burping_and_eructation	0.96468	0.01305
Bass_guitar	0.96487	0.01305
Purr	0.97692	0.01131
Bicycle_bell	0.97811	0.01166
Finger_snapping	0.98444	0.01305
Hi-hat	0.98444	0.01305
Marimba_and_xylophone	0.98444	0.01305
Skateboard	0.98444	0.01305
Strum	0.99556	0.01305

[80 rows x 2 columns]

7 Submission

```
[60]: df_test[df_test.columns[1:]] = pred_test
```

```
[61]: df_test.to_csv('submission.csv', index=False)
```